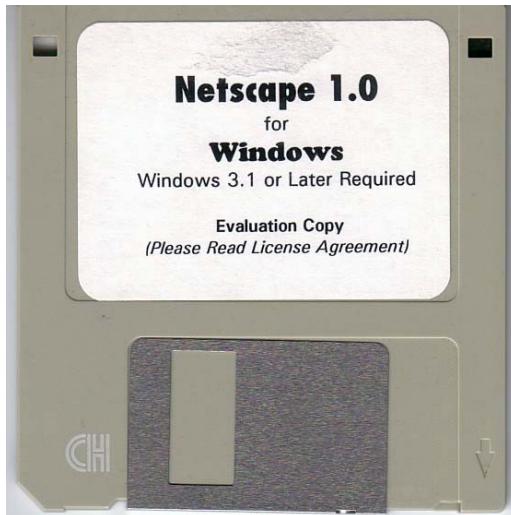


ARPANET 1974
<http://bpastudio.csudh.edu/fac/lpress/history/arpamaps/>

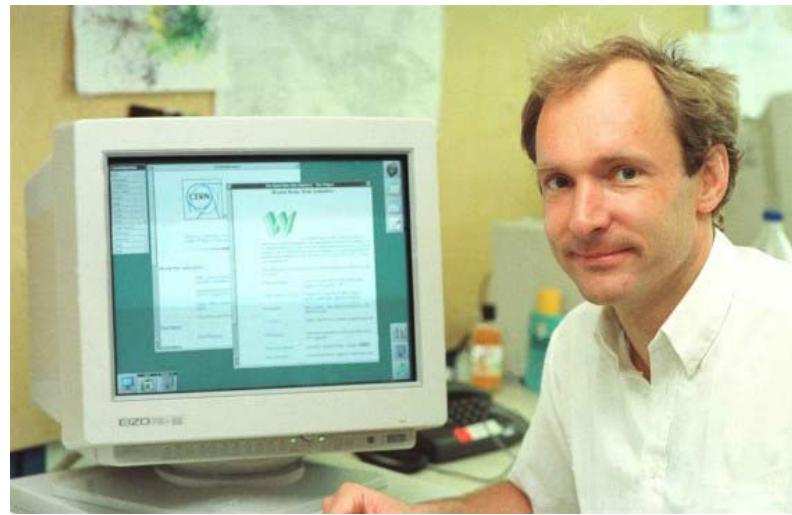
CSCE 560

Introduction to Computer Networking

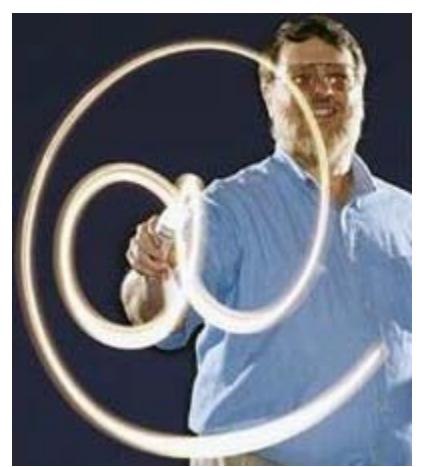
Dr. Barry Mullins
 AFIT/ENG
 Bldg 642, Room 209
 255-3636 x7979



<http://www.computerhistory.org/>



Sir Tim Berners-Lee - "inventor" of the WWW
 The world's first website, at CERN,
 went online August 6, 1991



Ray Tomlinson
 inventor of email 1971

Chapter 2: Application Layer

Our goals:

- Conceptual, implementation aspects of network **application protocols**
 - ❖ Transport-layer services
 - ❖ Client-server paradigm
 - ❖ Peer-to-peer paradigm
- Examine popular application-level protocols
 - ❖ HTTP
 - HyperText Transfer Protocol
 - ❖ FTP
 - File Transfer Protocol
 - ❖ SMTP / POP3 / IMAP
 - Simple Mail Transfer Protocol
 - Post Office Protocol
 - Internet Message Access Protocol
 - ❖ DNS
 - Domain Name System
- Create network applications
 - ❖ Socket API - application program interface

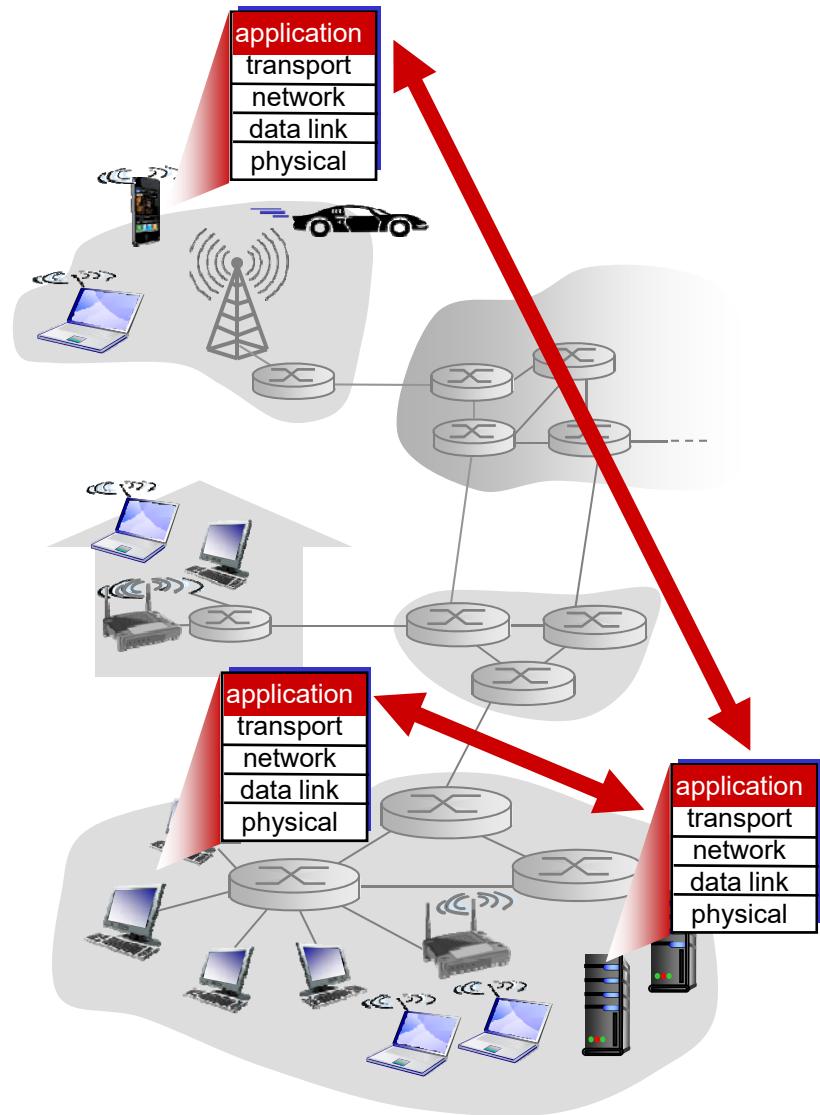
Creating a Network App

Write programs that

- Run on different end systems
- Communicate over a network
 - ❖ Web server software communicates with browser software

Very little software written for devices in network core

- Network core devices do not run user applications



Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications

Communicating Processes

Process: program (e.g., Chrome browser) running within a host

- Within same host, two processes communicate using **interprocess communication (IPC)** (defined by OS)
- Processes in different hosts communicate by exchanging **messages**

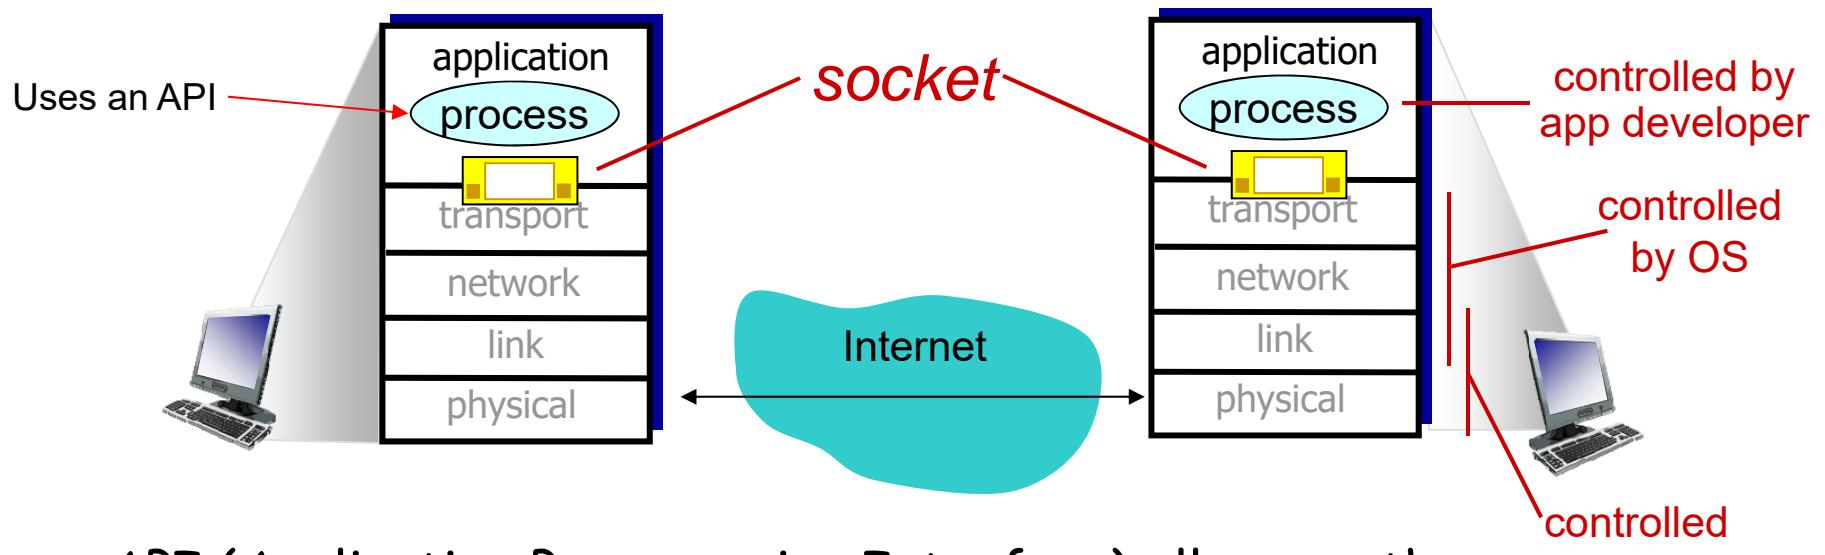
Client process: process that initiates communication

Server process: process that waits to be contacted

Fun fact: Applications with P2P architectures have client processes AND server processes

Sockets

- Process sends/receives messages to/from its socket
- Socket analogous to a door
 - ❖ Sending process shoves message out door
 - ❖ Sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



- API (Application Programming Interface) allows us the:
 1. choice of transport protocol
 2. ability to select a few parameters

Addressing Processes

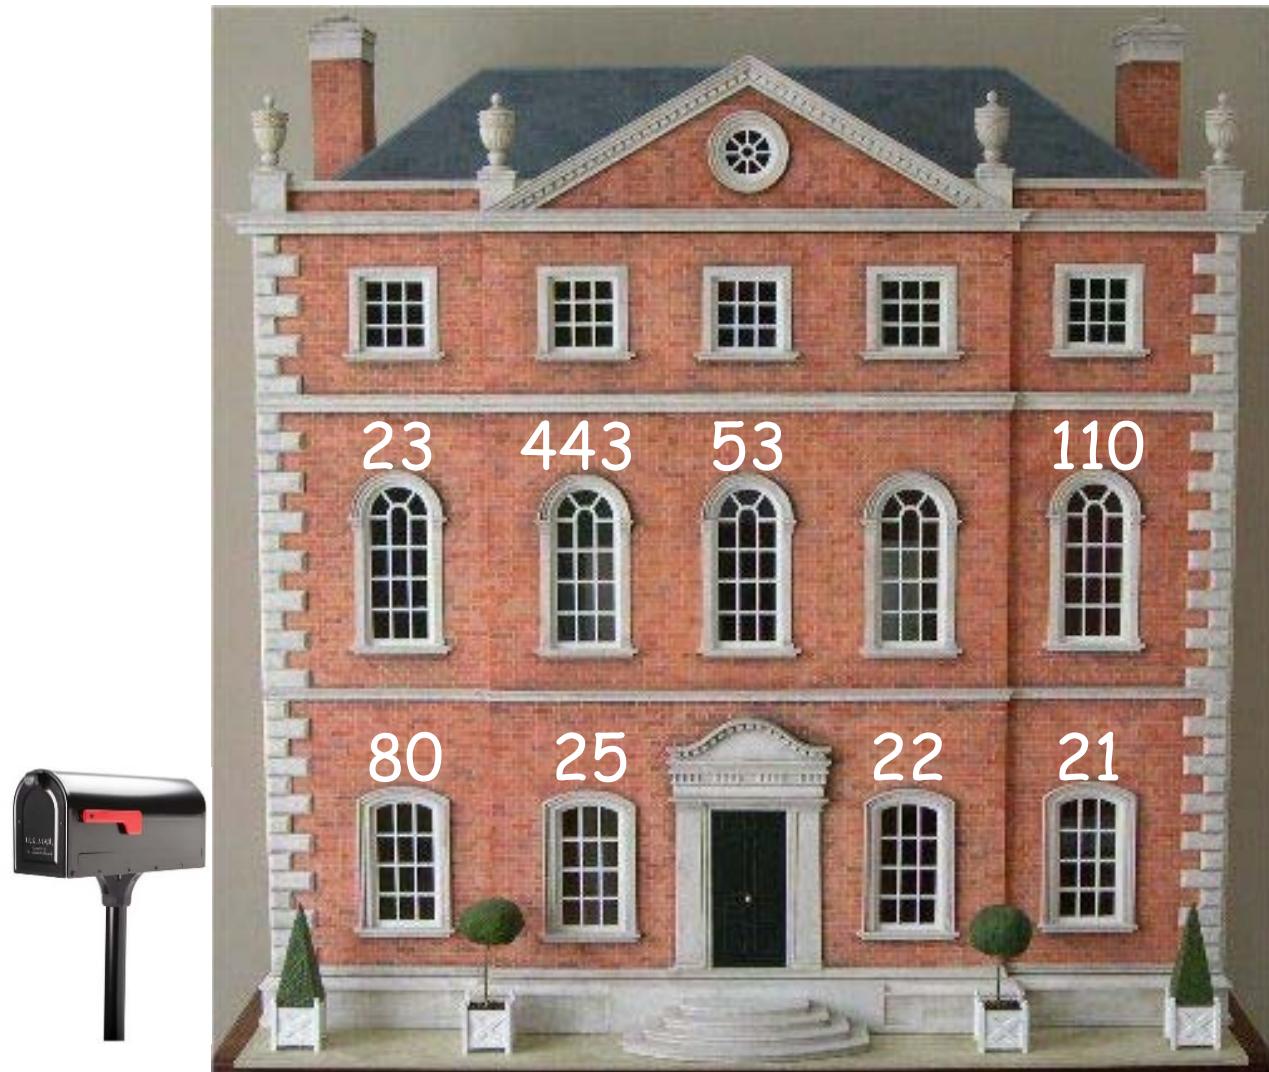
- ❑ For a process to receive messages, it must have an identifier
 - ❑ A host has a unique 32-bit IP address
- Q:** Does the IP address of the host on which the process runs suffice for identifying the process?
- A:** No, many processes can be running on same host
- ❑ Identifier includes both the **IP address** and **port numbers** associated with the process on the host
 - ❑ Example port numbers:
 - ❖ HTTP server: 80
 - ❖ Mail server: 25
 - ❑ To send HTTP message to `gaia.cs.umass.edu` web server:
 - ❖ **IP address:** 128.119.245.12
 - ❖ **Port number:** 80
 - ❑ More on this later

Preview - IP and Port Numbers

IP address → server

: Postal address → house

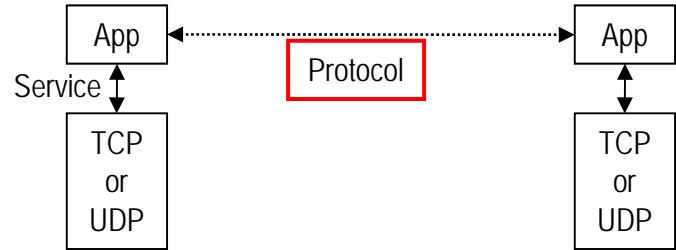
Port number → application (web, email) : Window number → person



Application Layer

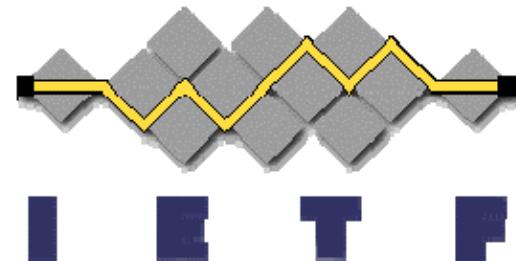
Protocol Defines:

- Types of messages exchanged
 - ❖ Request & response messages
 - ❖ Control messages
- Message syntax
 - ❖ What fields are in messages & how fields are delineated
- Message semantics
 - ❖ Meaning of information in fields
- Rules for when and how processes send & respond to messages



Public-domain protocols:

- Defined in RFCs
- Facilitates interoperability
- e.g., HTTP, SMTP
- Internet Engineering Task Force
 - ❖ www.ietf.org/rfc/
 - ❖ "The goal of the IETF is to make the Internet work better."



Proprietary protocols:

- e.g., Skype

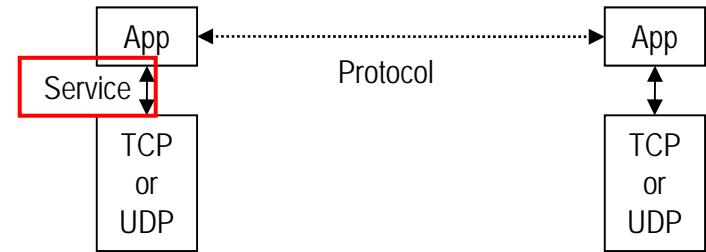
What Transport Service Does An App Need?

Data Integrity

- Some apps can tolerate some loss
 - ❖ Audio/Video
 - ❖ Interactive games
- Other apps require 100% reliable data transfer
 - ❖ File transfer
 - ❖ Web transactions

Timing

- Some apps require low delay to be "effective"
 - ❖ Internet telephony
 - ❖ Interactive games



Throughput

- Some apps require a minimum amount of throughput to be "effective"
 - ❖ Multimedia
- Other apps make use of whatever bandwidth they get → "elastic apps"
 - ❖ File transfer
 - ❖ Web traffic
 - ❖ Email

Internet Transport Protocol Services

TCP service:

- *Connection-oriented*: setup required between client and server processes
- *Reliable transport* between sending and receiving process
- *Flow control*: sender won't overwhelm receiver
- *Congestion control*: throttle sender when network overloaded
- *Does not provide*:
 - ❖ Timing
 - ❖ Minimum bandwidth guarantees

UDP service:

- Unreliable data transfer between sending and receiving process
- Does not provide:
 - ❖ Connection setup
 - ❖ Reliability
 - ❖ Flow control
 - ❖ Congestion control
 - ❖ Timing
 - ❖ Bandwidth guarantee

Q: Why bother?
Why is there a UDP?

Internet Apps: Application, Transport Protocols

| | Application | Application layer protocol | Underlying transport protocol |
|------------------------|--------------------|--|--|
| | Email | SMTP [RFC 2821] | TCP |
| Remote terminal access | | Telnet [RFC 854] | TCP |
| | Web | HTTP [RFC 2616] | TCP |
| | File transfer | FTP [RFC 959] | TCP |
| Streaming multimedia | | HTTP (e.g., YouTube) RTP [RFC 1889] | TCP or UDP |
| | Internet telephony | SIP, RTP, Proprietary (e.g., Skype) | TCP or UDP |
| Domain Name System | | DNS [RFC 1035 + many more] | UDP primarily |

RTP – Real-time Transport Protocol

SIP – Session Initiation Protocol

Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications



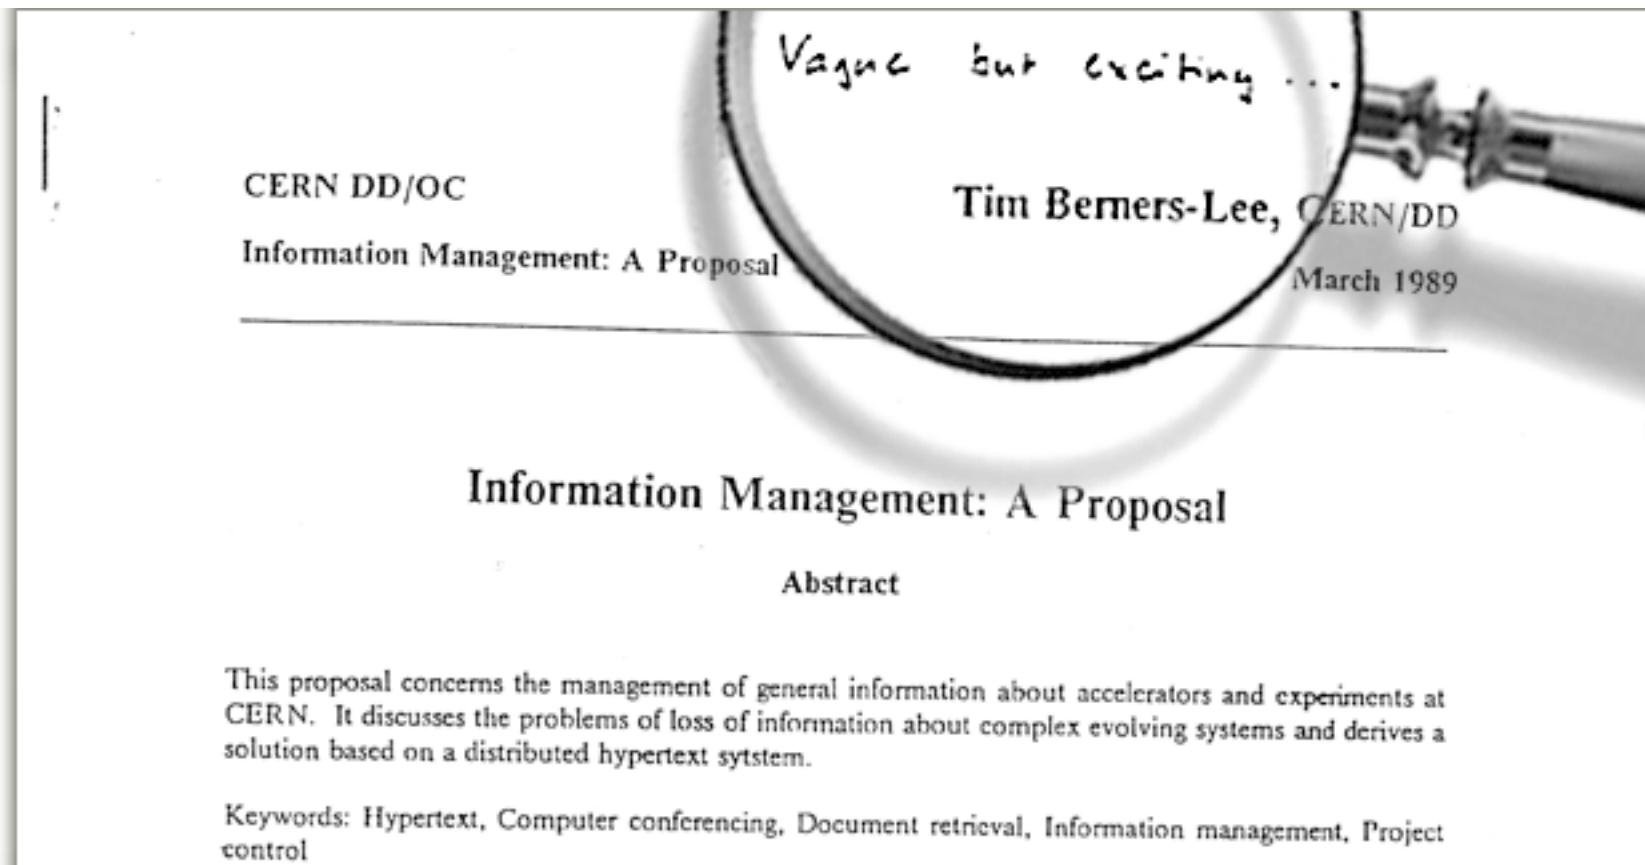
Ever have the feeling your “job title” understates your role?

Sir Timothy John "Tim" Berners-Lee is best known as the inventor of the World Wide Web

"Vague but exciting ..."

In March 1989, Tim Berners-Lee submitted a proposal for an information management system to his boss, Mike Sendall.

"Vague, but exciting", were the words that Sendall wrote on the proposal, allowing Berners-Lee to continue.



Web and HTTP

First some jargon

- Web page consists of objects
 - ❖ Object can be HTML file, JPEG image, Java applet, audio file, ...
 - ❖ An object can be ANY file
- Web page consists of base HTML file which includes several referenced objects
 - ❖ Each object is addressable by a URL
 - Uniform Resource Locator
- Example URL:

www.mit.edu/img/MIT_logo.gif

The diagram shows the URL "www.mit.edu/img/MIT_logo.gif" with three curly braces underneath it. The first brace covers "www.mit.edu", the second covers "img", and the third covers "MIT_logo.gif". Below each brace is a label: "host name" under the first, "path" under the second, and "filename" under the third.

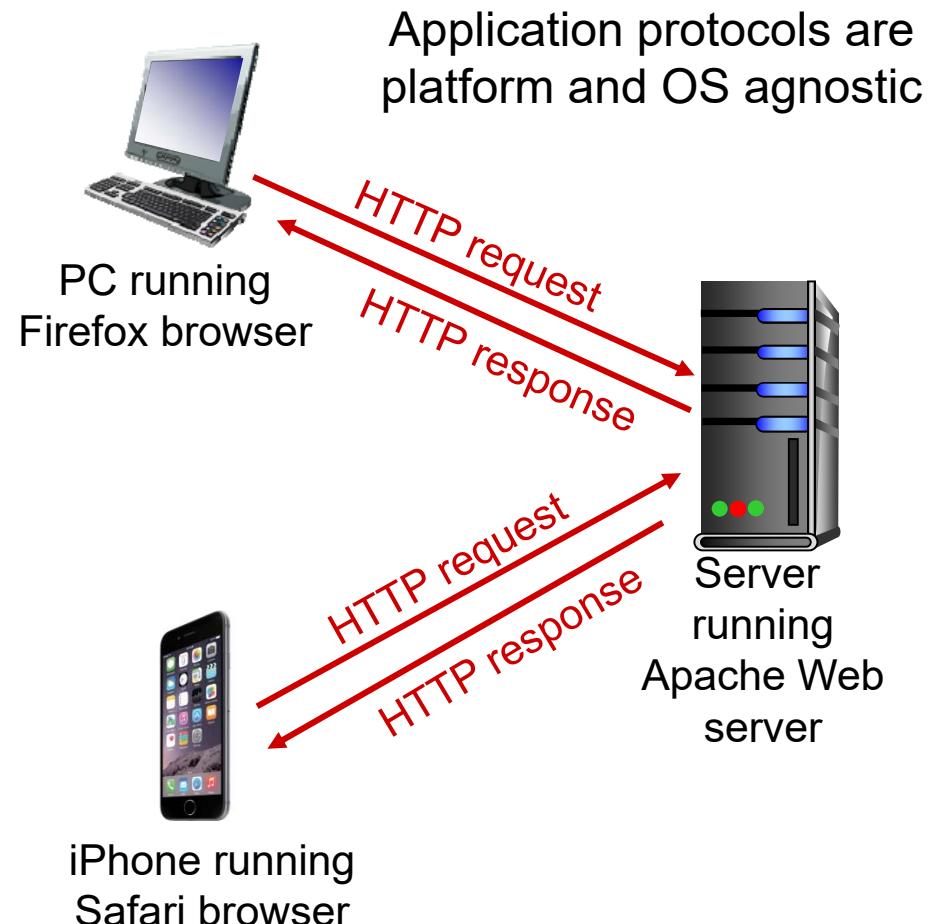
host name path filename

HTTP: Hypertext Transfer Protocol

- Web's application layer protocol

Client/Server model

- *Client*: browser that requests, receives, and "displays" Web objects
- *Server*: Web server sends objects in response to requests
- HTTP documented in
 - ❖ HTTP 1.0: RFC 1945
 - ❖ HTTP 1.1: RFC 2616

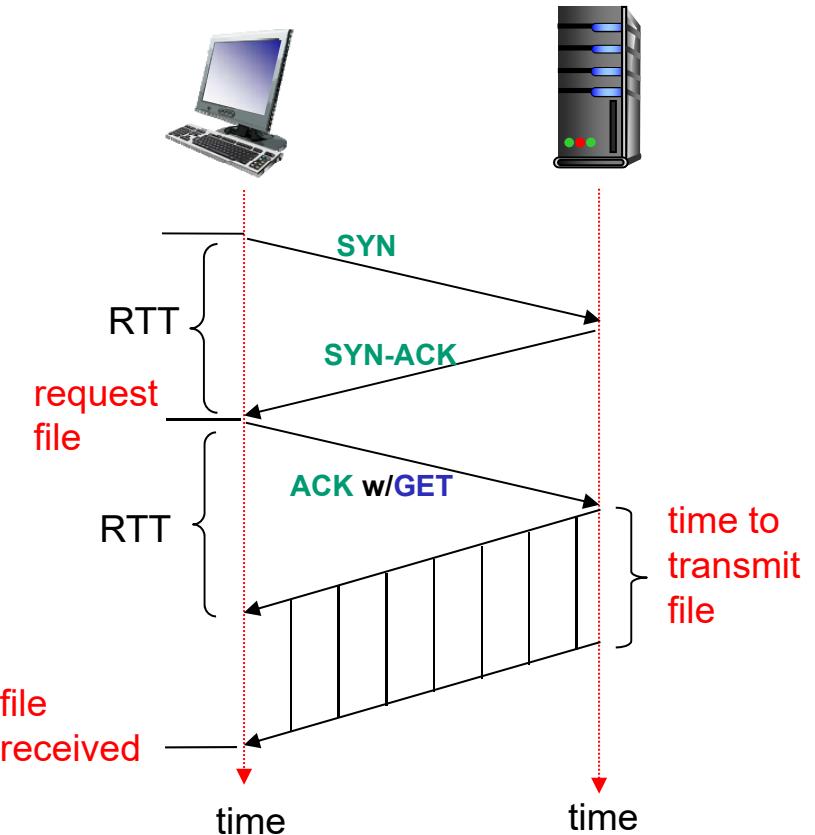


HTTP Overview

Uses TCP:

- ❑ Client initiates **TCP** connection (creates socket) to server on port 80
 - ❖ Sends a "**SYN**" packet
- ❑ Server accepts TCP connection from client
 - ❖ Sends a "**SYN-ACK**" packet
- ❑ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❑ TCP connection closed (not shown on figure)

RTT: Time required for a small packet to travel from client to server and back



File transmission time:

$$\text{Total} = 2\text{RTT} + \text{file transmit time}$$

HTTP Connections

Non-persistent HTTP

- At most only **one** object is sent over a TCP connection

- HTTP/1.0 uses non-persistent HTTP

Persistent HTTP

- Multiple objects can be sent over single TCP connection between client and server

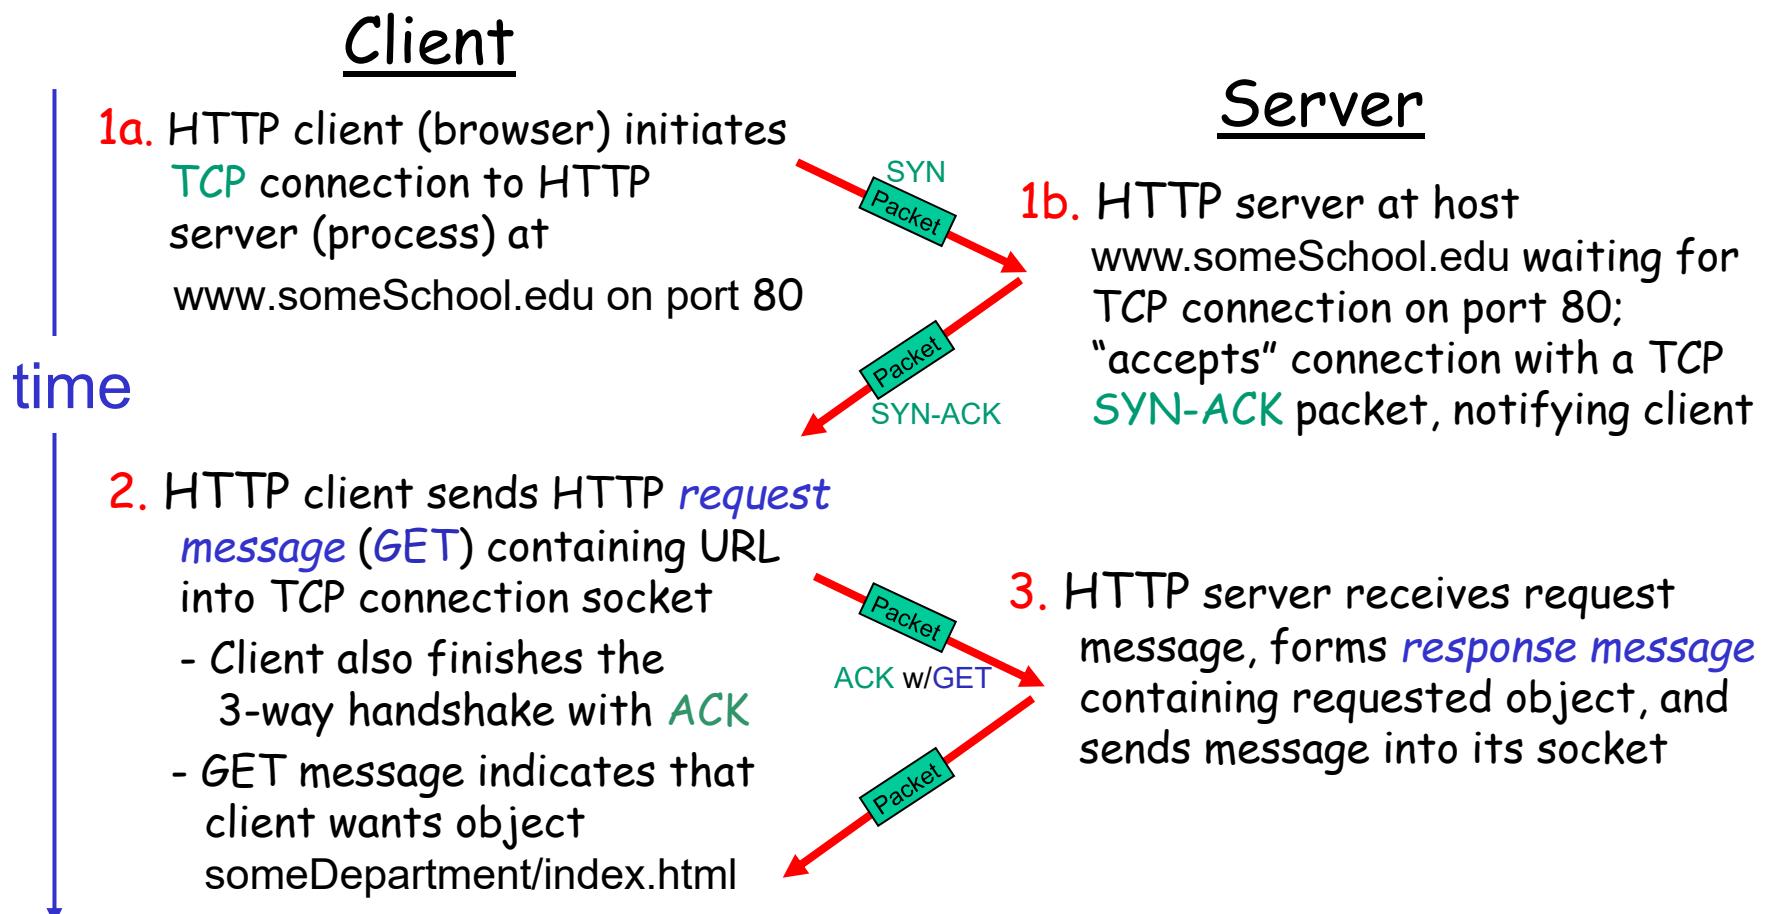
- HTTP/1.1 uses persistent connections by default

Non-persistent HTTP

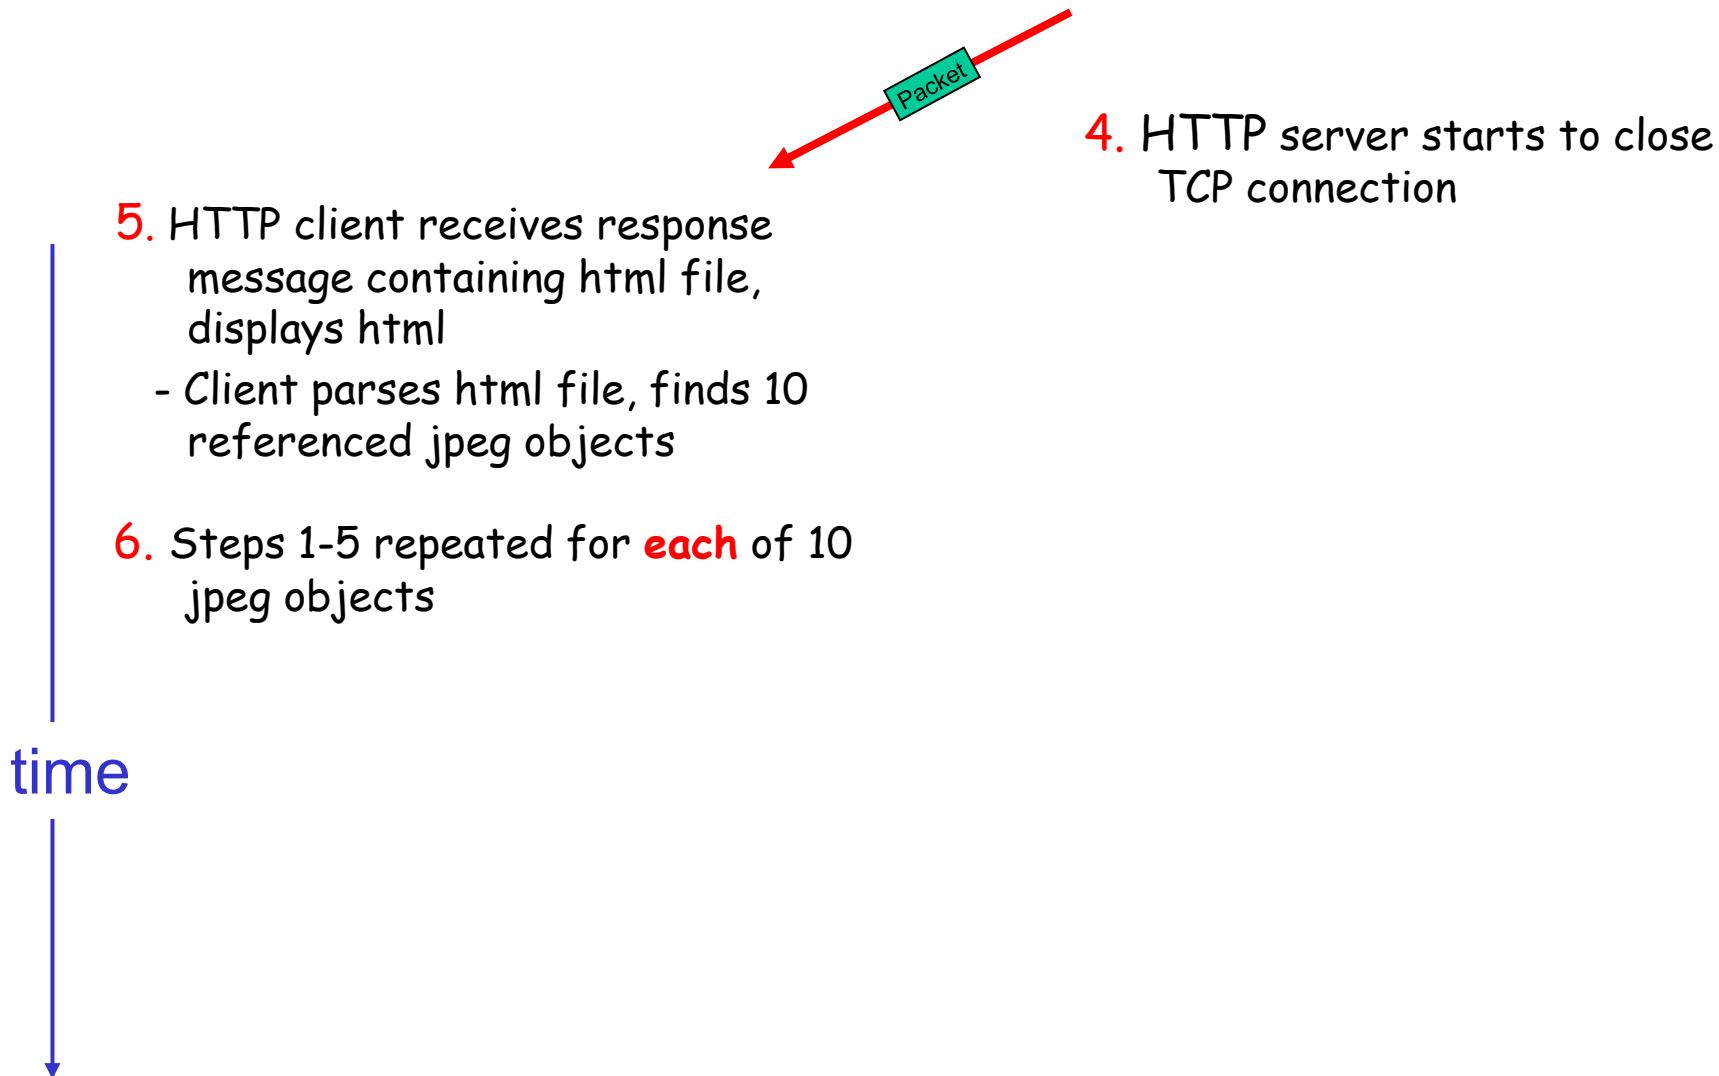
Suppose user enters URL

`www.someSchool.edu/someDepartment/index.html`

(contains base html and references to 10 jpeg images)



Non-persistent HTTP (cont.)



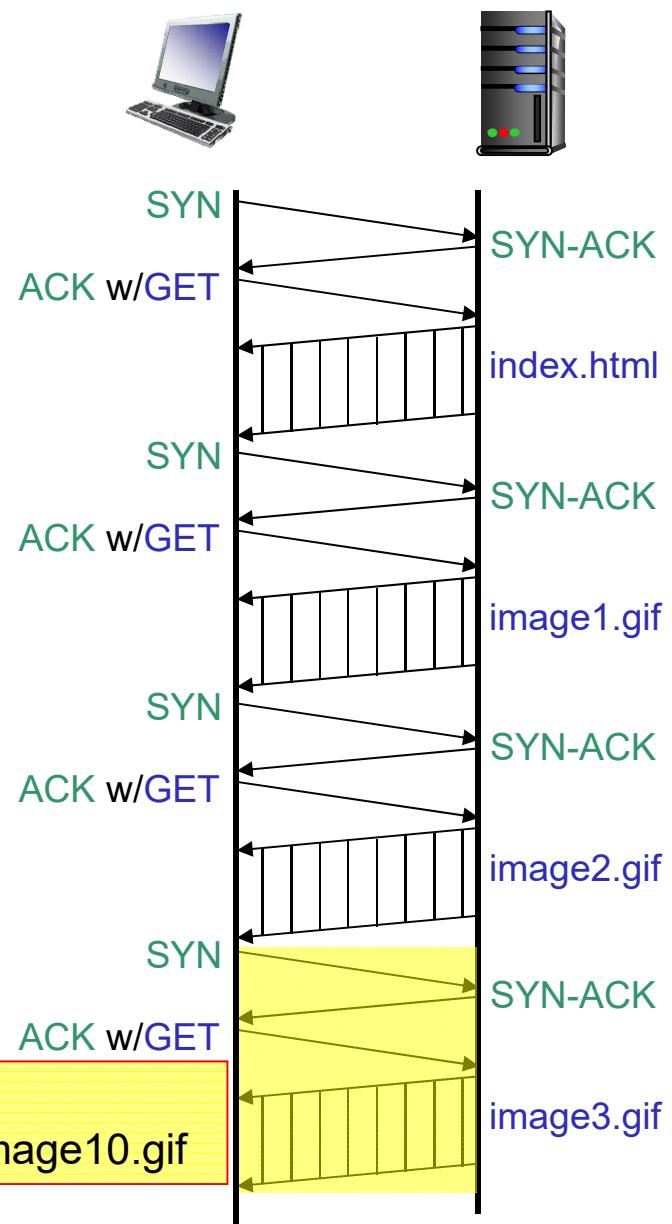
Visualizing Non-persistent Connections

Non-persistent HTTP issues:

- Requires 2 RTTs per object
- OS overhead for each TCP connection
 - ❖ TCP buffers and variables



Repeat for
image4.gif → image10.gif



Persistent HTTP

Persistent HTTP

- ❑ Server leaves TCP connection open after sending response
- ❑ Subsequent HTTP messages between same client/server sent over open connection

Persistent without pipelining:

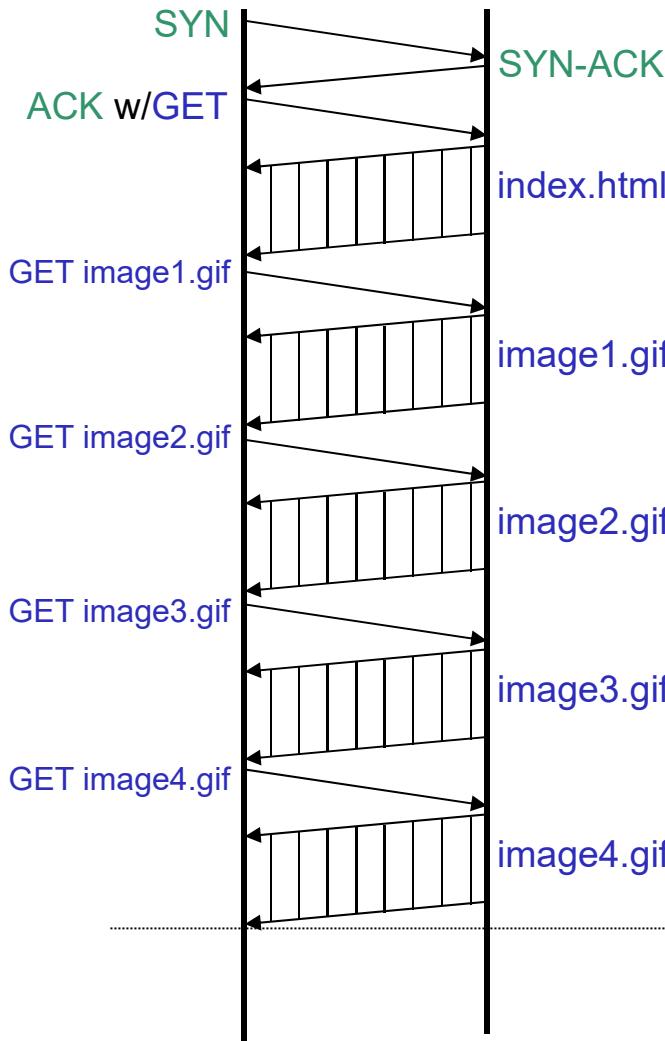
- ❑ Client issues new request only when previous response has been received
- ❑ One RTT for each referenced object

Persistent with pipelining:

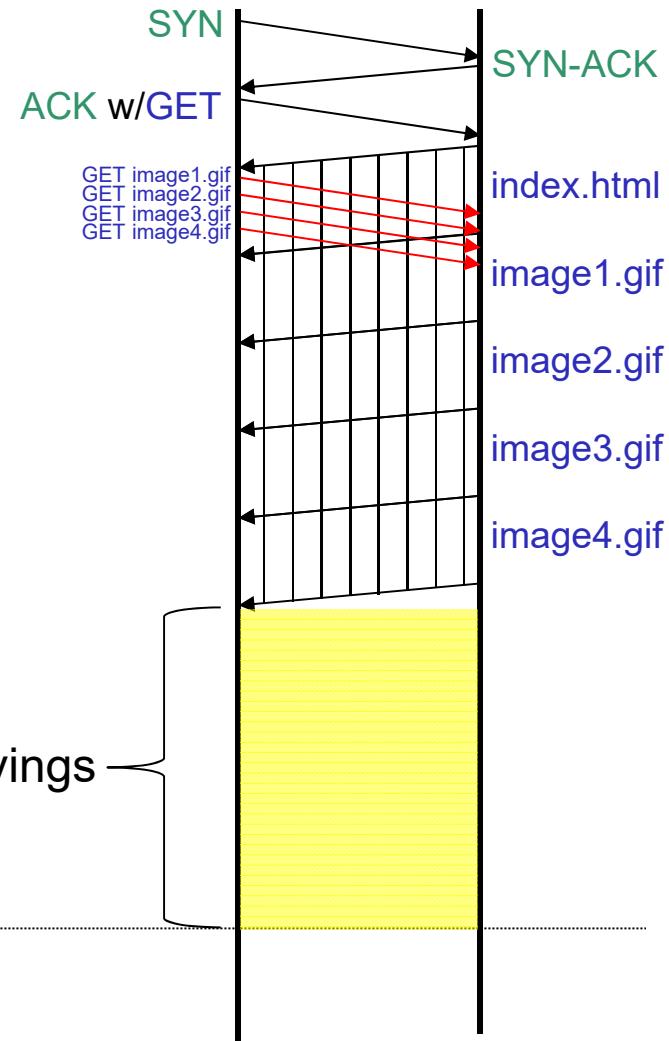
- ❑ Default in HTTP/1.1
- ❑ Client sends requests as soon as it encounters a referenced object
- ❑ Specification requires that the server respond to the requests in the order in which they are received
 - ❖ Must use FIFO
 - ❖ As a result, one large object can back up many smaller objects.
 - Known as *head of line blocking*

Visualizing Pipelining - Persistent Connections

Without pipelining



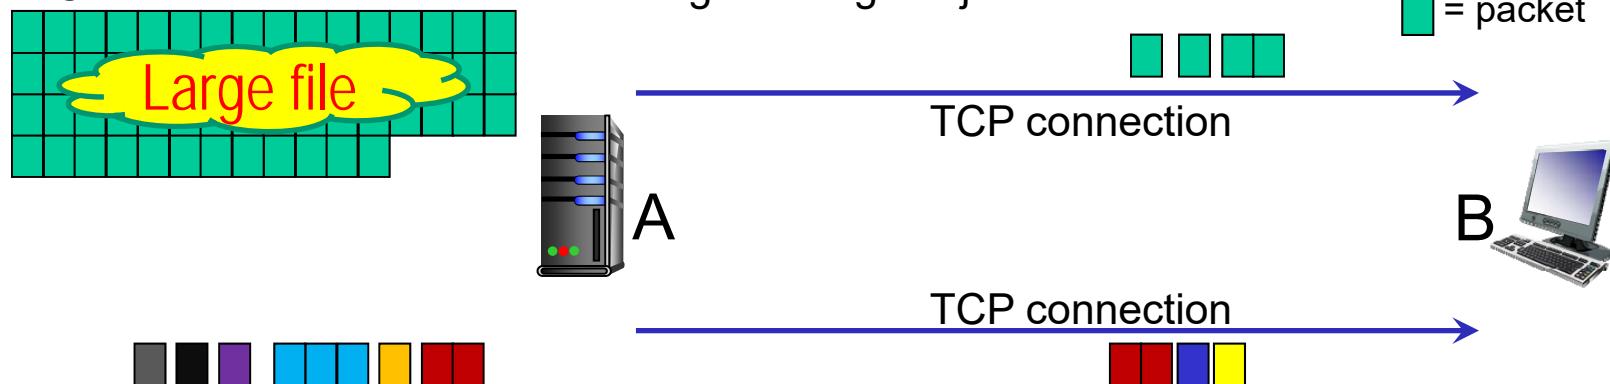
With pipelining



Multiple Parallel TCP Connections

- As an alternative to, or in addition to, pipelining, many clients can open more than one TCP session (connection) to the server and then make requests through each connection
- This mitigates head of line blocking as shown below
- The HTTP 1.1 standard specifies that each client should open no more than two connections to any one server
 - ❖ In practice, this is basically ignored

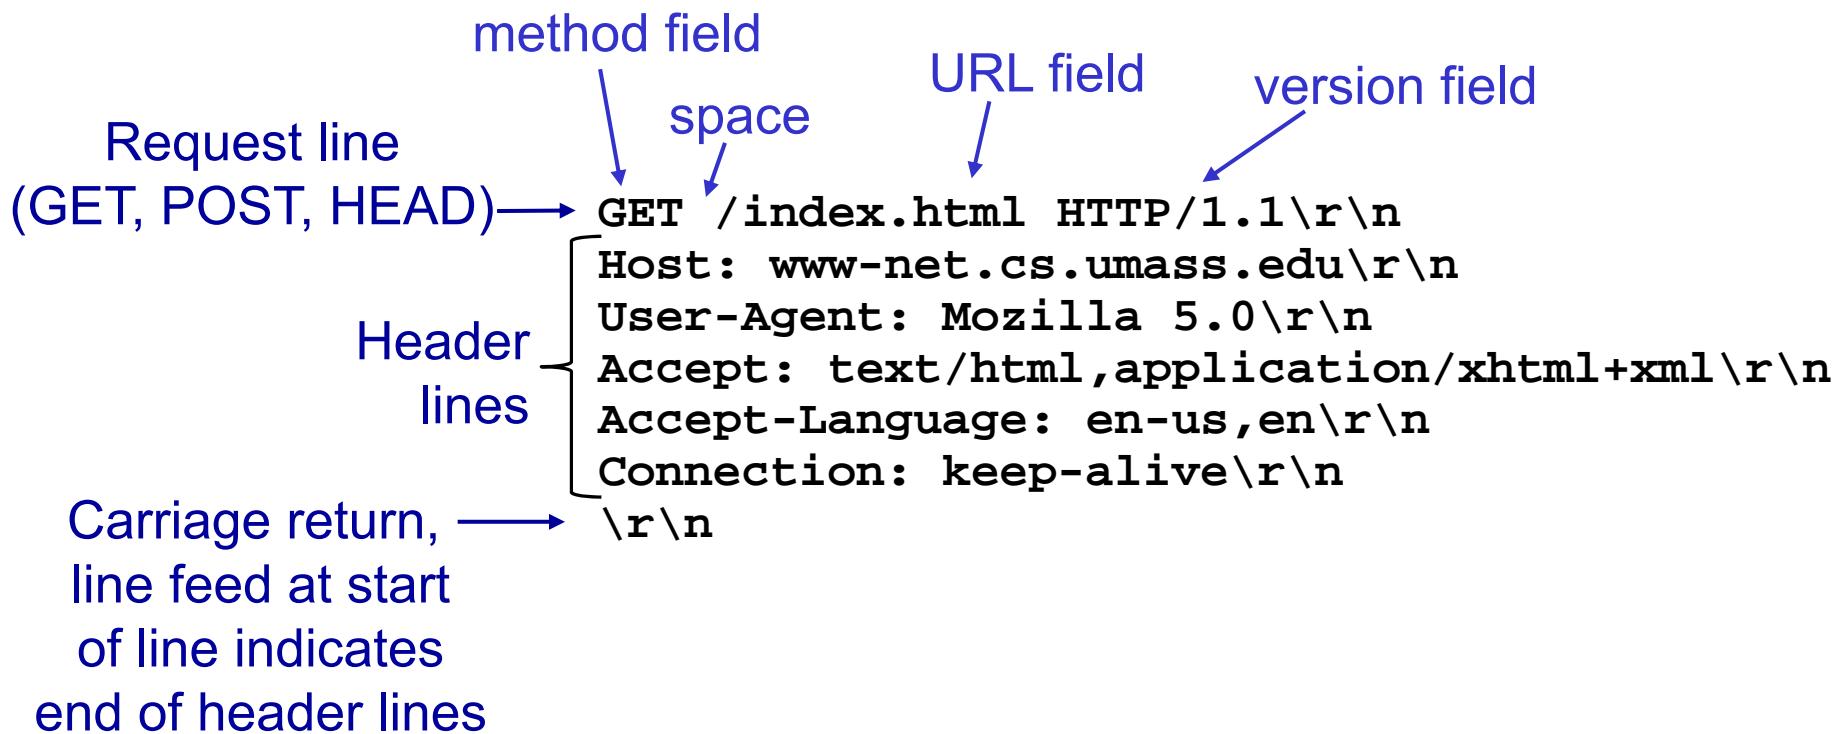
First TCP connection from A to B sending one large object



Second TCP connection from A to B sending several small objects

HTTP Request Message

- Two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ❖ ASCII (human-readable format)



User agent: a client application implementing a network protocol

HTTP User-Agent Trickery

- Since 1990s, most major browsers use
 - ❖ **User-Agent: Mozilla/5.0**
 - ❖ Code name for the now-defunct Netscape Web browser
 - ❖ Netscape was viewed as the cutting-edge browser of its day
- Browsers pretended to be Netscape to trick Web servers into sending it the most sophisticated version of their site even though most of them aren't derived from Mozilla software
 - ❖ Internet Explorer
 - ❖ Chrome
 - ❖ Firefox
 - ❖ Torch
- Microsoft Edge → **User-Agent: Microsoft-CryptoAPI/10.0**

<http://www.washingtonpost.com/blogs/the-switch/wp/2013/09/23/this-hacker-might-seem-shady-but-throwing-him-in-jail-is-bad-for-everyone/>

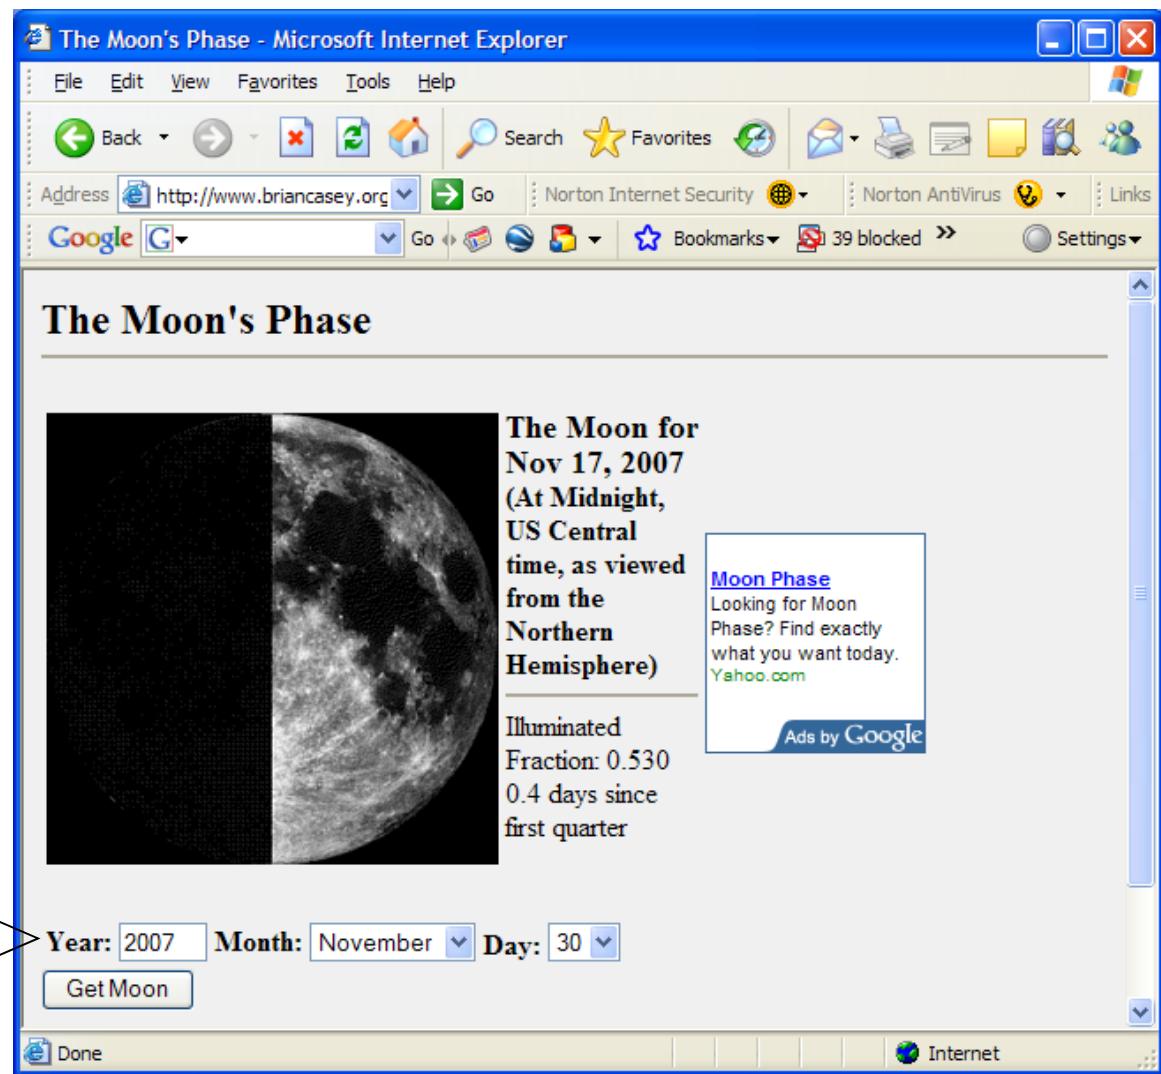
Uploading Form Input - Get Method

- Input is uploaded in **URL** field of request line:

```
GET http://www.mapquest.com/maps?address=2950+Hobson+Way&zipcode=45433 HTTP/1.1
<<snipped some header information>>
Host: www.mapquest.com
<CRLF>
```

Uploading Form Input - Post Method

- ❑ A web page may include form input
 - ❖ Address to locate
- ❑ Input is uploaded to server in **body** of request
 - ❖ aka entity body
- ❑ Currently viewing 17 Nov 07
- ❑ Want to see 30 Nov 07 ...



Uploading Form Input - Post Method

```
POST http://www.briancasey.org/artifacts/astro/moon.cgi HTTP/1.1
```

```
Content-Length: 25
```

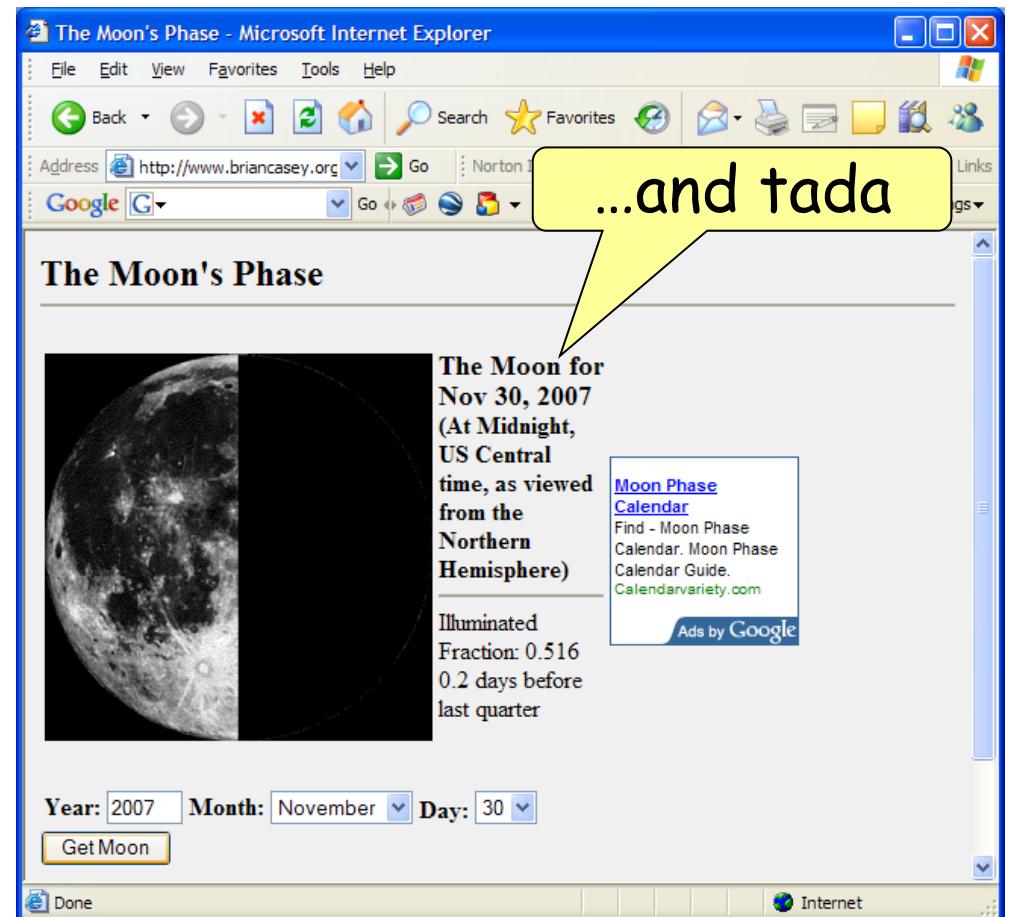
```
Host: www.briancasey.org
```

```
<<snipped some header information>>
```

```
<CRLF>
```

```
year=2007&month=11&day=30
```

25 characters



Method Types

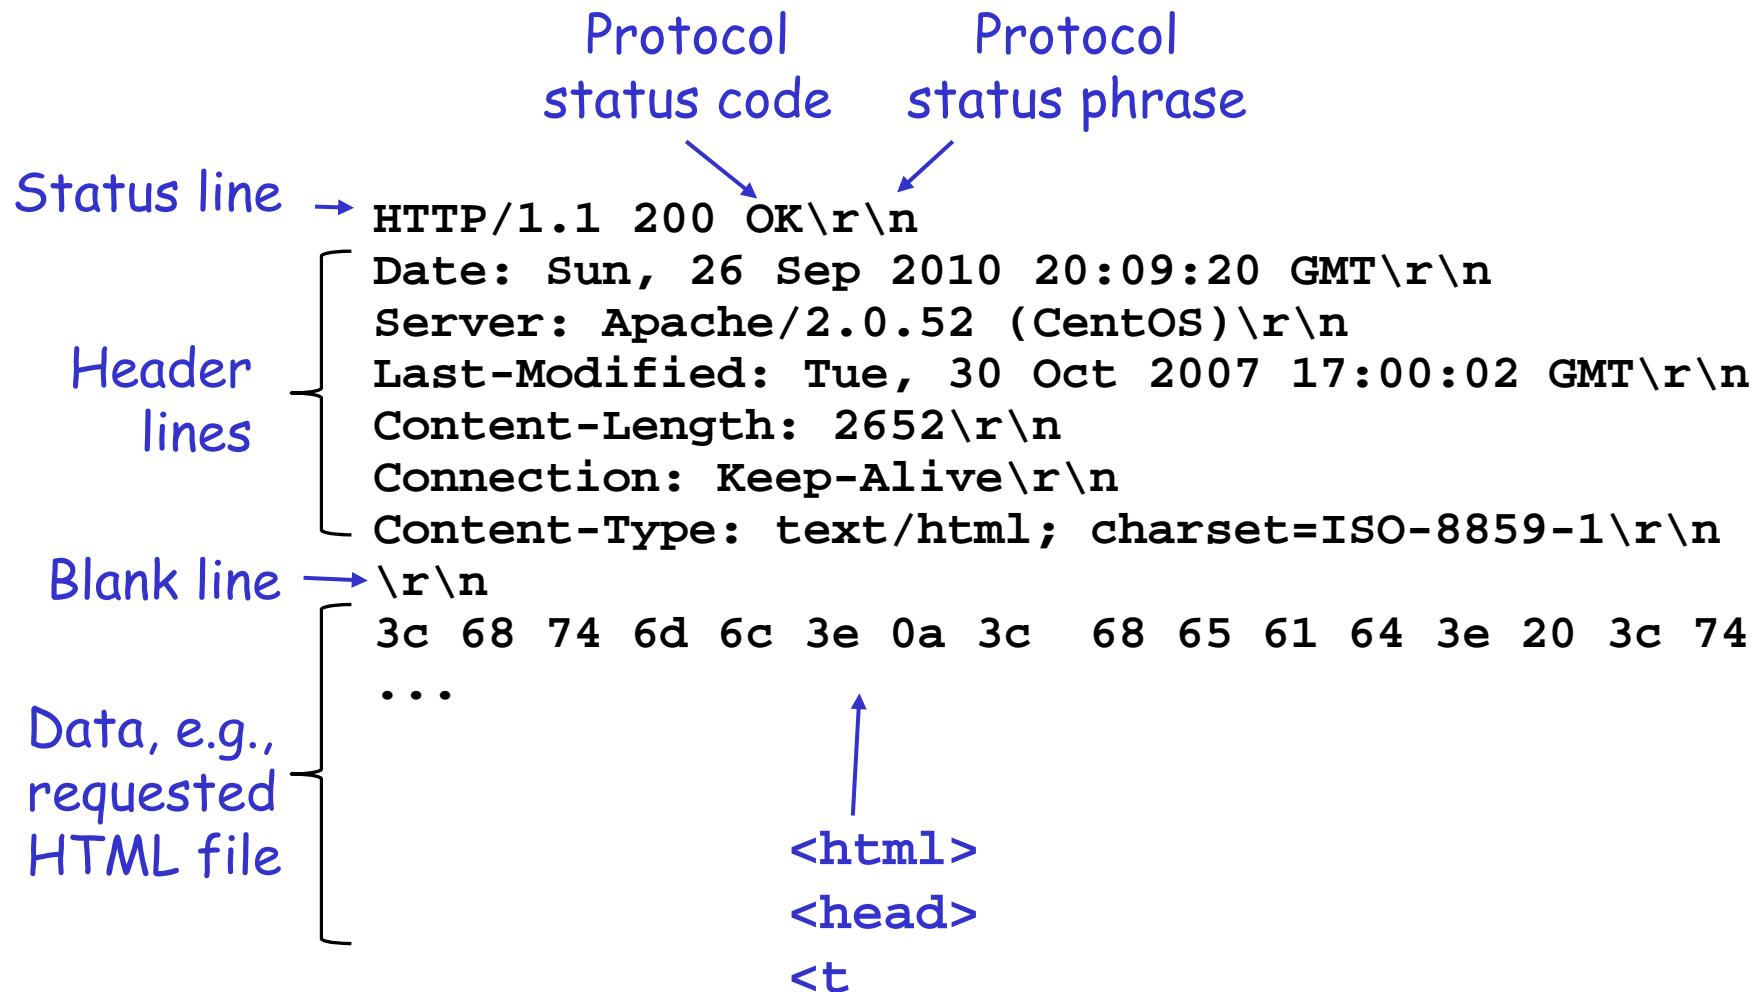
HTTP/1.0

- GET
- POST
- HEAD
 - ❖ Asks server to leave requested object out of response
 - ❖ Used for debugging

HTTP/1.1

- GET, POST, HEAD
- PUT
 - ❖ Uploads file in entity body to path specified in URL field
- DELETE
 - ❖ Deletes file specified in the URL field

HTTP Response Message



HTTP Response Status Codes

In first line of server->client response message

A few sample codes:

200 OK

- ❖ Request succeeded, requested object later in this message

301 Moved Permanently

- ❖ Requested object moved, new location specified later in this message (Location:)

400 Bad Request

- ❖ Request message not understood by server

505 HTTP Version Not Supported

1xx Intermediate Status
2xx Successful Response
3xx Redirects
4xx Request Errors
5xx Server Errors



More exhaustive list

<http://www.w3.org/Protocols/rfc2616/rfc2616-sec10.html>

HTTP Response Status Codes

404 Not Found

- ❖ Requested document not found on this server

Flavor not found?!?!



1xx Intermediate Status
2xx Successful Response
3xx Redirects
4xx Request Errors
5xx Server Errors



HTTP Response Status Codes

404 Not Found

2018 Quidditch
National Championship



Manual HTTP (GET)

Notice the server does not echo the characters

Open a command shell and type (or paste) the **RED** text

```
> telnet gaia.cs.umass.edu 80      Start telnet session to Web service (port 80) on gaia.cs.umass.edu  
Trying 128.119.245.12...  
Connected to gaia.cs.umass.edu. May not see this notice  
Escape character is '^]'.
```

Request line

```
GET http://gaia.cs.umass.edu/wireshark-labs/HTTP-wireshark-file2.html HTTP/1.0 [ENTER]  
Connection: keep-alive [ENTER] ← Header line  
[ENTER] ← Carriage return, line feed at start of line indicates end of header lines
```

Empty entity body

HTTP/1.1 200 OK

Status line

Date: Mon, 12 Oct 2015 13:31:14 GMT
Server: Apache/2.4.6 (CentOS) OpenSSL/1.0.1e-fips PHP/5.4.16 mod_perl/2.0.9dev Perl/v5.16.3
Last-Modified: Mon, 12 Oct 2015 05:59:01 GMT
ETag: "173-521e206aa287e"
Accept-Ranges: bytes
Content-Length: 371
Keep-Alive: timeout=5, max=100
Connection: Keep-Alive
Content-Type: text/html; charset=UTF-8

Response Header

<html>

Blank line – separates header from data
Data (Webpage)

Congratulations again! Now you've downloaded the file lab2-2.html.

This file's last modification date will not change. <p>
Thus if you download this multiple times on your browser, a complete copy

will only be sent once by the server due to the inclusion of the IN-MODIFIED-SINCE

field in your browser's HTTP GET request to the server.

</html>

Manual HTTP (POST)

Notice the server does not echo the characters

Open a command shell and type (or paste) the **RED** text

```
> telnet www.briancasey.org 80    Start telnet session to Web service (port 80) on www.briancasey.org  
Trying 192.155.194.19...  
Connected to www.briancasey.org. May not see this notice  
Escape character is '^]'.  
[REDACTED]
```

```
POST http://www.briancasey.org/artifacts/astro/moon.cgi HTTP/1.1 [ENTER]  
Content-Length: 25 [ENTER]
```

Host: www.briancasey.org [ENTER]

[ENTER] ← Carriage return, line feed at start of line indicates end of header lines

```
year=2007&month=11&day=30 < Entity body
```

HTTP/1.1 200 OK Status line

Date: Mon, 12 Oct 2015 13:51:07 GMT

Server: Apache/2.2.15 (CentOS)

Connection: close

Response Header

Transfer-Encoding: chunked

Content-Type: text/html; charset=ISO-8859-1

Blank line – separates header from data

```
<html>  
<head> <title> The Moon's Phase </title> </head>  
<body bgcolor="#F0F0F0">
```

Data (Webpage)

<<snipped>>

Web Server Scripting

- A URL may refer to a static web page or a server-side script
 - ❖ Script is just a program that is run in response to a HTTP request
- Server-side scripts create “dynamic” web pages based on the input
- Common Gateway Interface (CGI)
 - ❖ Standard argument passing convention between web server and clients
- CGI scripts may be written in any language
 - ❖ PHP, Python, Perl, Ruby, sh, csh, Java
- CGI scripts are commonly used to produce responses to Web page form input from client browsers
- <http://www.briancasey.org/artifacts/astro/moon.cgi>

Client Side Web Page Scripts/Programs

- Web pages may also contain scripts or programs within the HTML code to be run on the client
- Unlike server scripts, web page scripts and programs run on the browser machine's processor, not on the server's processor
- Examples:
 - ❖ Javascript
 - ❖ VBScript
 - ❖ Java applets
- http://media.pearsoncmg.com/aw/aw_kurose_network_2/applets/message/messagessegmentation.html

HTTP is “Stateless”

- ❑ Server maintains no information about past client requests

Protocols that maintain “state” are complex!

- ❑ Past history (state) must be maintained
- ❑ If server/client crashes, their views of “state” may be inconsistent and must be reconciled

How to keep “state”

- ❑ Protocol endpoints
 - ❖ Maintain state at sender/receiver over multiple transactions
- ❑ **Cookies:** HTTP headers carry state

Cookies

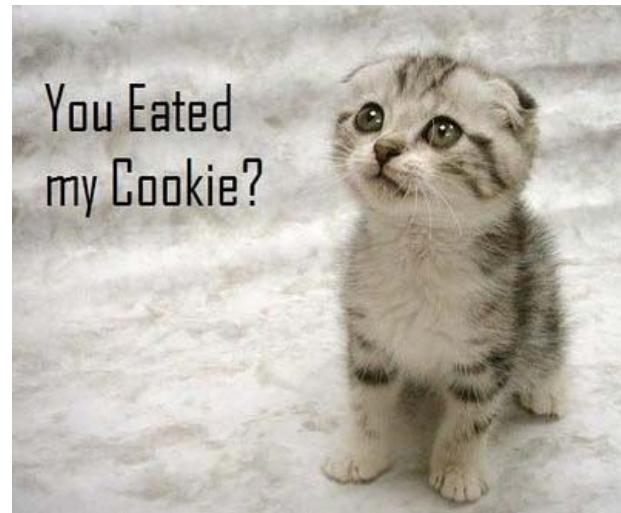
Most major Web sites use cookies

What cookies can bring:

- Authorization
- Shopping carts
- Recommendations
- User session state
 - ❖ Web email

aside
Cookies and privacy:

- Cookies permit sites to learn a lot about you
- You may supply name and email to sites
- Search engines use redirection & cookies to learn even more



Providing State Information: Cookies

Four components:

- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

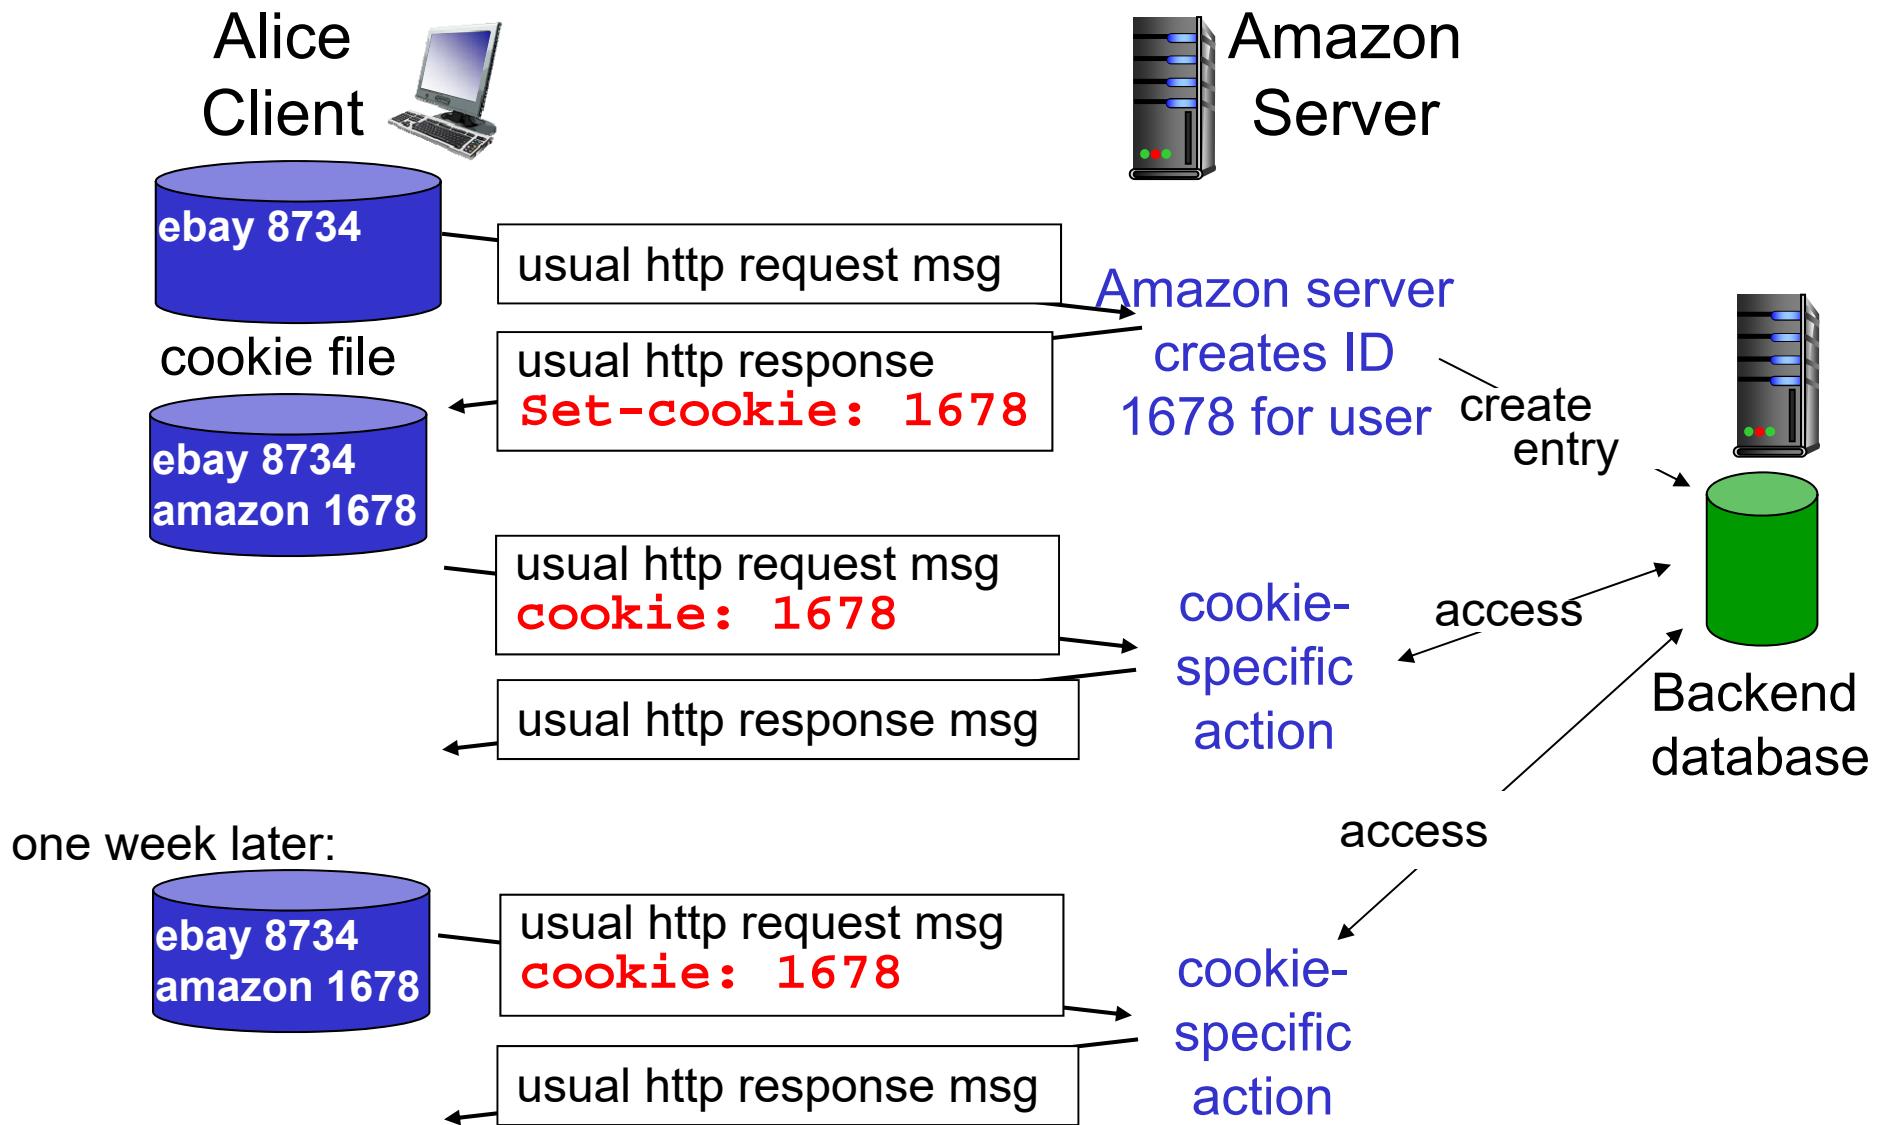
- Alice accesses Internet always from same PC
- She visits a specific e-commerce site for first time
- When initial HTTP requests arrives at site, site creates
 - ❖ a unique ID and
 - ❖ an entry in backend database for ID

"A cookie is a well-known computer science term that is used when describing an opaque piece of data held by an intermediary."

<http://cookiecentral.com/faq/#1.2>

Windows 10 cookies folder can be displayed using
Press the Windows key and R together. Type shell:cookies and click OK.

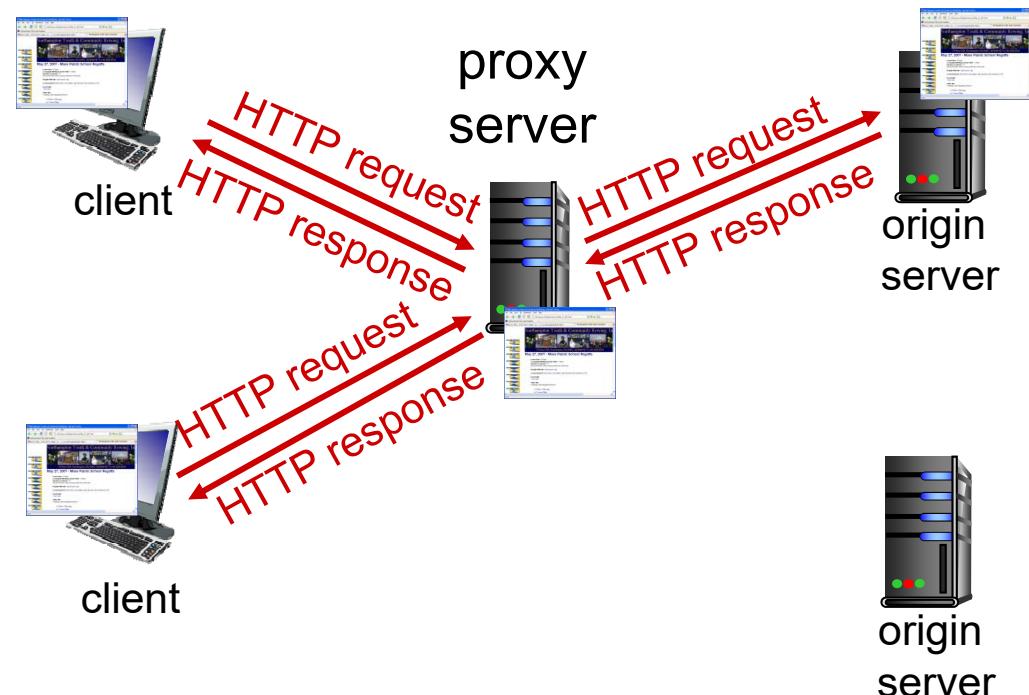
Cookies: Keeping "State"



Web Caches (Proxy Server)

Goal: satisfy client requests without involving origin server

- ❑ User starts browser:
Web accesses via proxy
- ❑ Browser sends all HTTP
requests to proxy
 - ❖ If object in cache:
cache returns object if
"fresh"
 - ❖ else proxy requests
object from origin
server, then returns
object to client
 - Keeps a copy in cache



Fresh: Object has an expiry time or other age-controlling header set and is still within the fresh period

Expires: 31 Dec 2024 14:19:41 GMT

More About Web Caching

- Cache acts as both client and server
 - ❖ Server for original requesting client
 - ❖ Client to origin server
- Typically cache installed by ISP
 - ❖ University, company, residential ISP

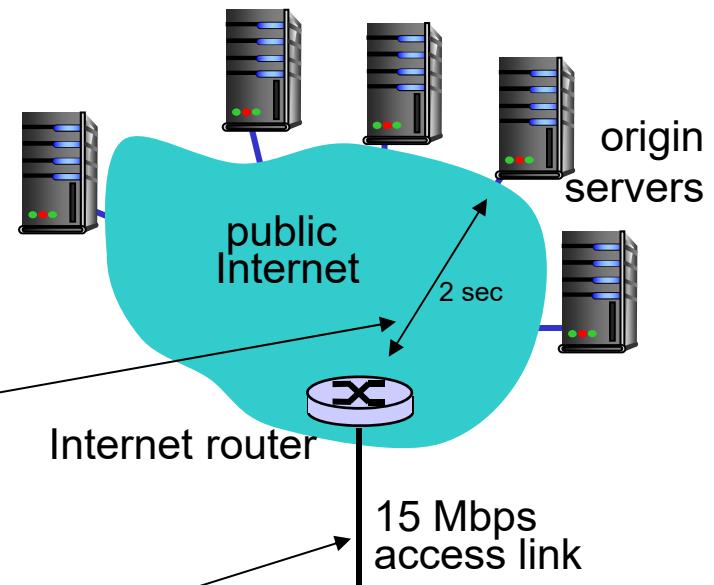
Why Web caching?

- Reduce response time for client request
- Reduce traffic on an institution's access link
- Reduce web traffic on the Internet

Caching Example - Baseline

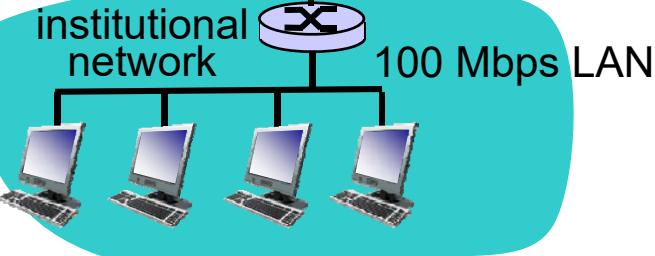
Assumptions

- ❑ Average object size = 1,000,000 bits
- ❑ Avg. request rate from institution's browsers to origin servers = 15/sec
 - ❖ 15Mbps required
- ❑ RTT from Internet router to any origin server and back to router = 2 sec (Internet delay)



Consequences

- ❑ Utilization on access link
 $15\text{Mbps}/15\text{Mbps} = 100\%$
 - ❖ Traffic intensity = 1
- ❑ LAN Utilization = $15\text{Mbps}/100\text{Mbps} \rightarrow 15\%$
- ❑ Total delay = $\text{Internet delay} + \text{access delay (delay between routers)} + \text{LAN delay}$
 $= 2 \text{ sec} + \text{minutes} + \text{milliseconds}$



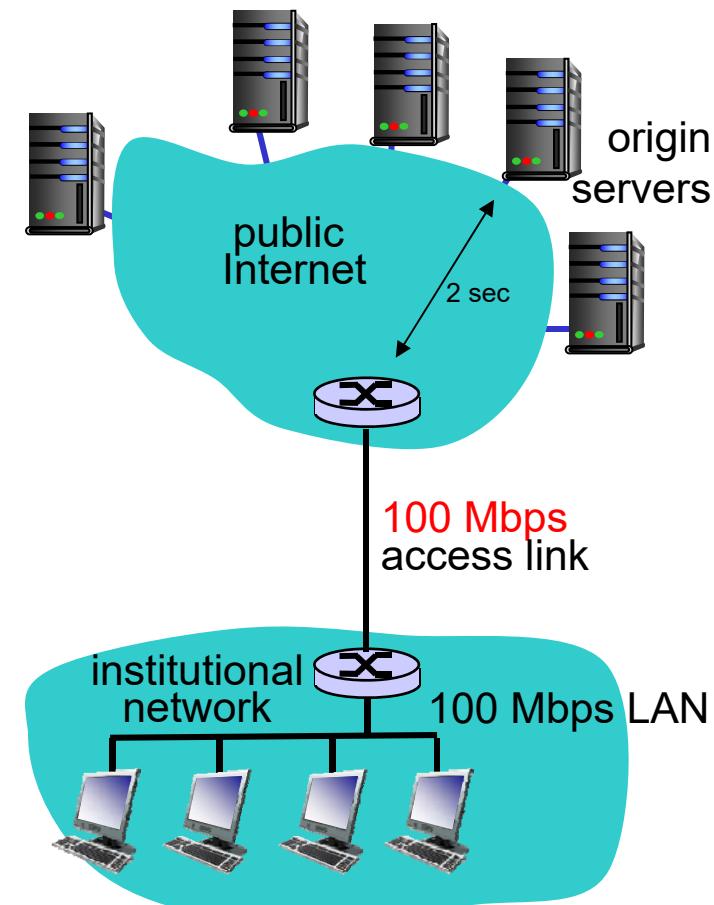
Caching Example - Increase Bandwidth

Possible solution

- ❑ Increase bandwidth of access link to **100 Mbps**

Consequences

- ❑ Utilization on LAN = 15%
- ❑ Utilization on access link = 15%
- ❑ Total delay = Internet delay + access delay + LAN delay
= 2 sec + msec + msec
- ❑ **Often a costly upgrade**



Caching Example - Install a Cache

Install a cache

- ❑ Suppose hit rate is 0.4

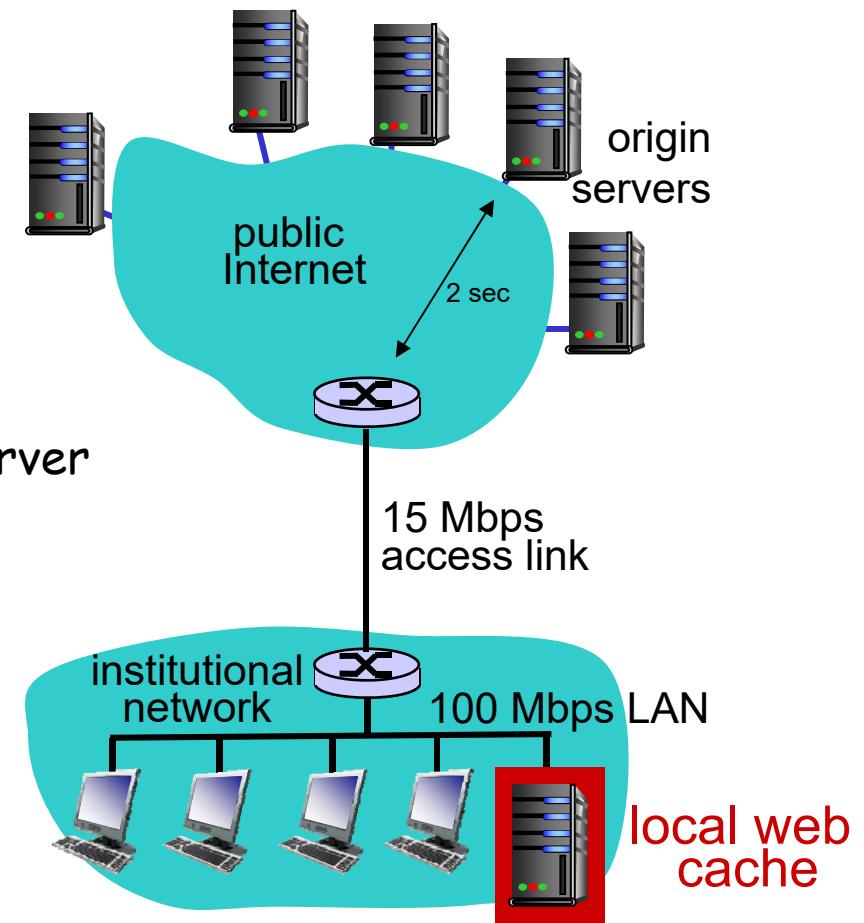
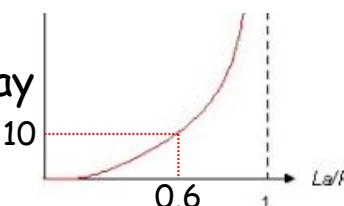
Consequence

- ❑ 40% requests are fresh and can be satisfied almost immediately (10 msec)
 - ❖ 60% requests satisfied by origin server

- ❑ Access link utilization:

- ❖ 60% of requests use access link
 - ❖ Data rate over access link
 $= 0.6 * (15)(1,000,000) \text{Mbps} = 9 \text{ Mbps}$
 - ❖ With 0.6 utilization
 - Assume 10 msec delay

- ❑ Total avg delay =
 - ❖ $0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied from cache})$
 - ❖ $0.6 * (2.01)$
 - ❖ $0.4 * (10 \text{ msec})$



Expires: 31 Dec 2024 14:19:41 GMT

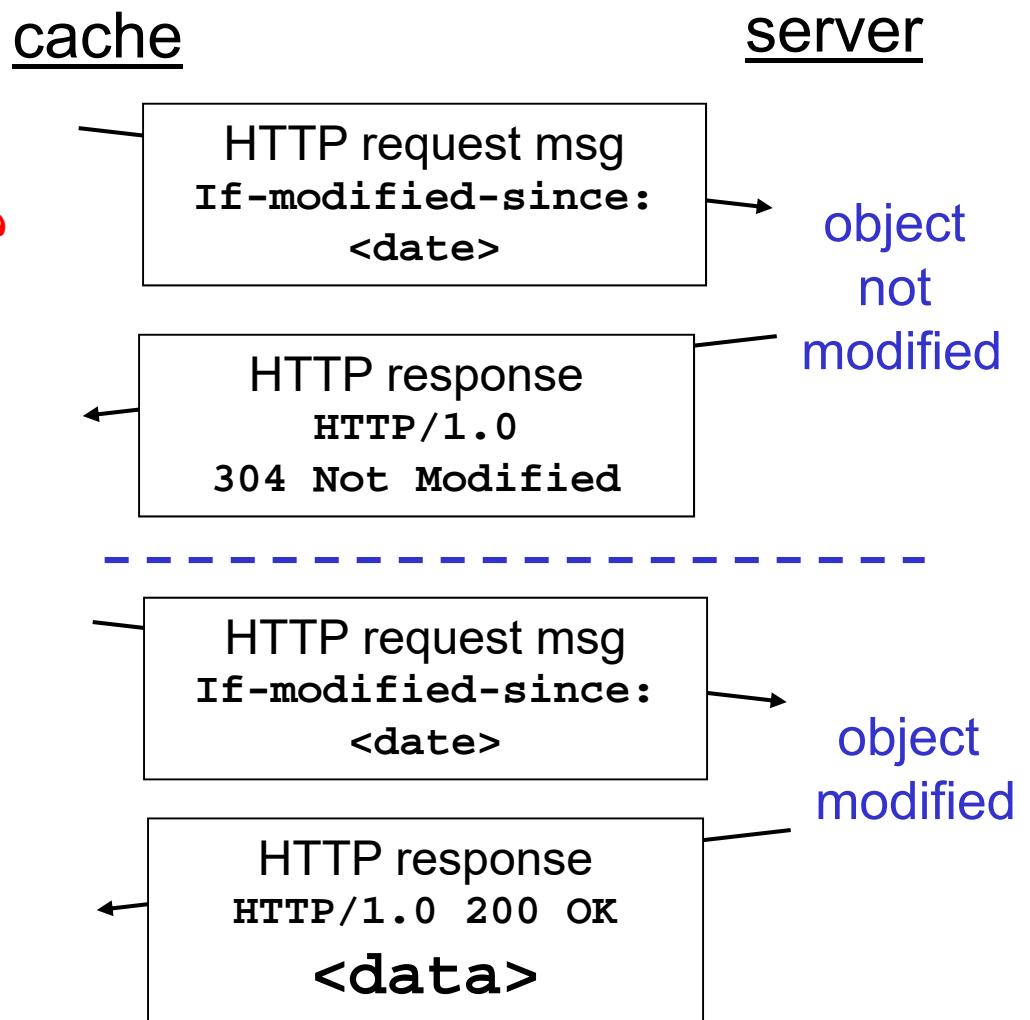
Should everything be cached?

Bank transactions?

Conditional GET

- Q: How do we know the cached object is the latest version (i.e., "fresh") if it has "expired"?
- A: Use a conditional GET to validate the object

- Goal: don't send object if cache already has up-to-date version
- Cache: specify date of cached copy in HTTP request
 If-modified-since: <date>
- Server: response contains no object if cached copy is up-to-date:
 HTTP/1.0 304 Not Modified



Conditional Get Example

1. First cache request

```
GET /fruit/kiwi.gif HTTP/1.1  
Host: www.exotiquecuisine.com
```

2. First server response

```
HTTP/1.1 200 OK  
Date: Mon, 7 Jul 2003 15:39:29  
Server: Apache/1.3.0 (Unix)  
Last-Modified: Wed, 2 Jul 2003  
09:23:24  
Content-Type: image/gif  
  
(data data data ...)
```

4. Subsequent cache request

```
GET /fruit/kiwi.gif HTTP/1.1  
Host: www.exotiquecuisine.com  
If-modified-since: Wed, 2 Jul 2003 09:23:24
```

3. Cache saves Last-Modified time of the object

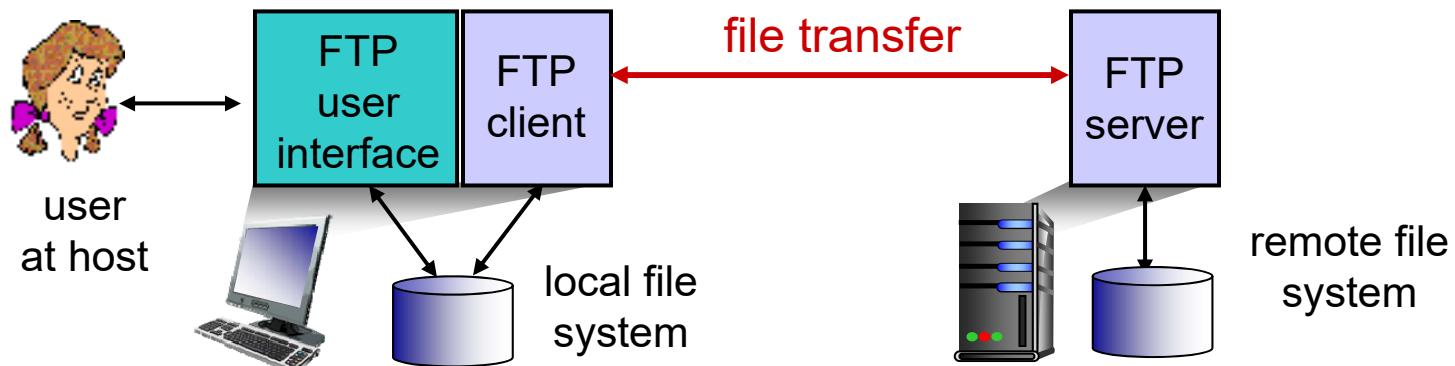
5. Subsequent server response

```
HTTP/1.1 304 Not Modified  
Date: Mon, 14 Jul 2003 15:39:29  
Server: Apache/1.3.0 (Unix)  
  
(empty entity body ...)
```

Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- **FTP**
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications

FTP: File Transfer Protocol



- Transfer file to/from remote host
- Client/Server model across TCP
 - ❖ **Client**: side that initiates transfer
 - ❖ **Server**: remote host
- FTP: RFC 959
- FTP server listens on port 21

FTP Commands, Responses

Sample commands:

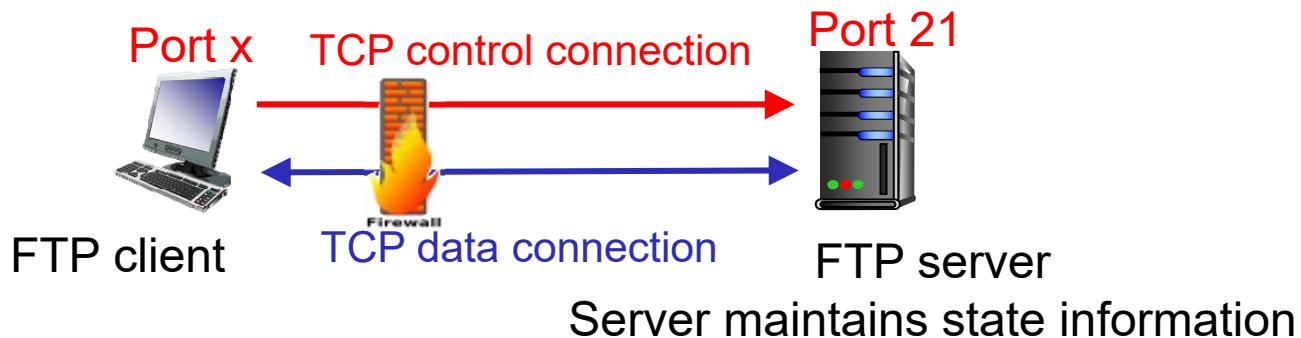
- Sent as ASCII text over control channel
- USER username**
- PASS password**
- LIST / LS** return list of files in current directory
- GET filename** retrieves a file
- PUT filename** sends a file to the remote host (server)

Sample return codes

- Status code and phrase (as in HTTP)
- Sent as ASCII text over control channel
- 331 Username OK, password required**
- 125 data connection already open; transfer starting**
- 425 Can't open data connection**
- 452 Error writing file**

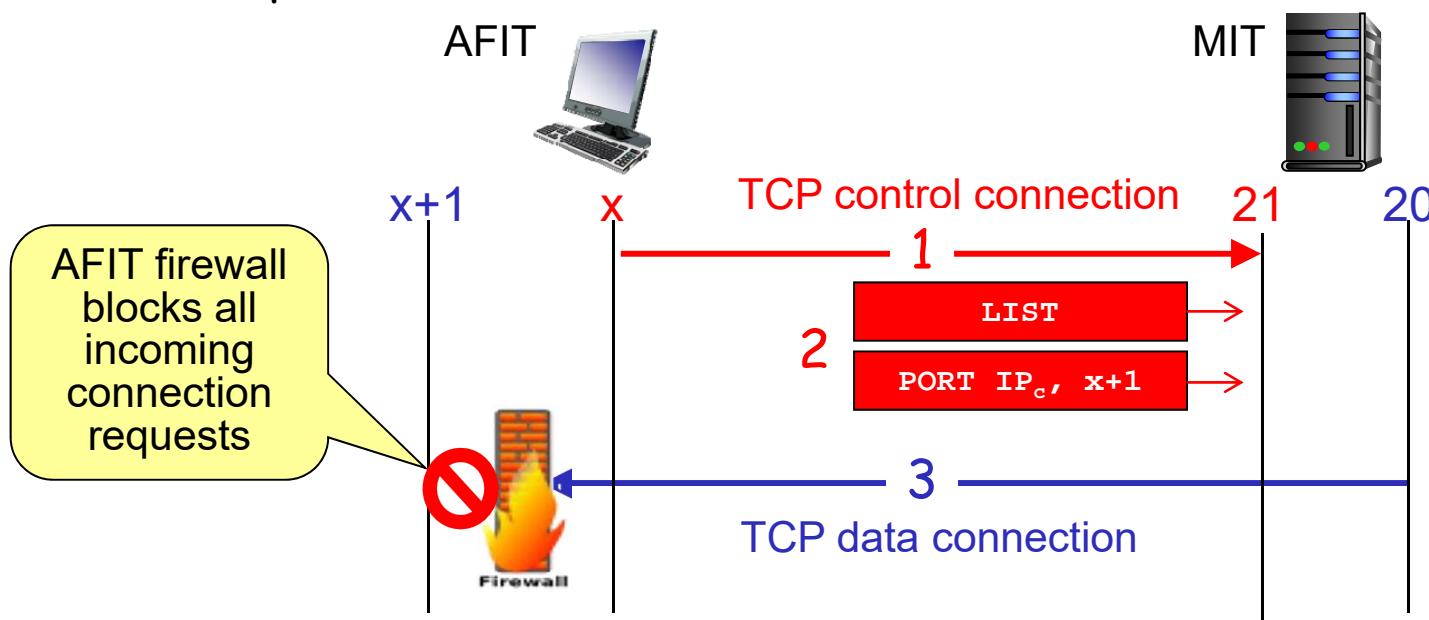
FTP: Separate Control, Data Connections

- ❑ Uses TWO connections
- ❑ **Control** connection is used to
 - ❖ Obtain authorization (username and password)
 - ❖ Browse remote directories on server
 - ❖ Control connection is called "**out of band**"
- ❑ **Data** connection is used to transfer data including directory listings
 - ❖ FTP can run in active or passive mode, which determines how the data connection is established
 - ❖ Use passive mode if incoming connections are blocked by firewall



Active FTP (Server → Client)

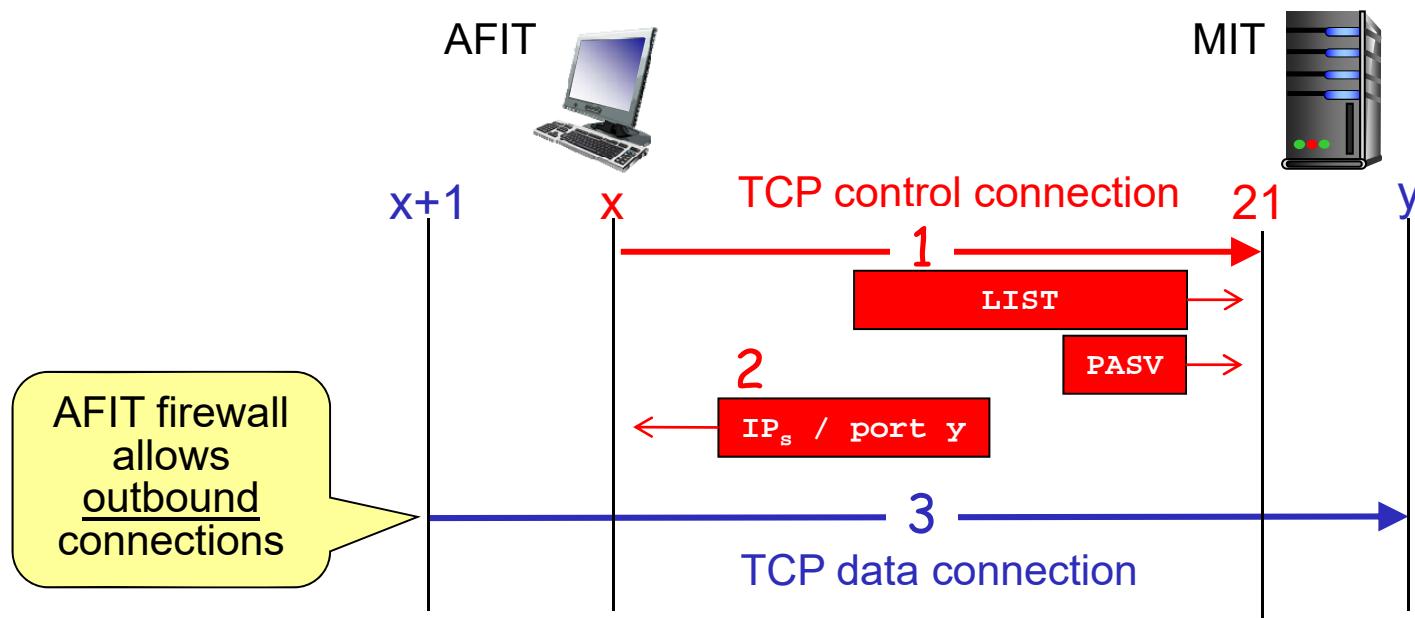
1. AFIT client (port x : $x > 1023$) connects to FTP server (MIT) on TCP port 21
2. When client sends a command for a file transfer / directory listing (LIST) ...
 - ❖ Client starts listening on port $x+1$ and sends PORT command across command connection to server to provide client's IP address (IP_c) and listening port ($x+1$)
3. **Server** opens a TCP data connection from port 20 to port $x+1$ on client
 - ❖ After transferring one file/listing, server closes data connection
4. Subsequent file transfers/listings → repeat steps 2 and 3 with different client ports (i.e., not $x+1$)



Passive FTP (Client → Server)

Most Web browsers
(which act as FTP clients)
use passive FTP by default

1. AFIT client opens two ports (port x and $x+1$; $x > 1023$) and connects the first port (x) to FTP server (MIT) on TCP port 21
2. When client sends a command for a file transfer / directory listing (LIST) ...
 - ❖ Client sends PASV and waits for server response
 - ❖ Server responds with its IP address (IP_s) / listening port (y : $y > 1023$)
3. **Client** then opens a TCP data connection from port $x+1$ to port y on IP_s
4. Subsequent file transfers/listings → repeat steps 2 and 3 with different client ports (i.e., not $x+1$)



Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications

*Stupid computer
Keeps saying
"you got mail"*



Email

- Email is transferred from one host to “another” using the Simple Mail Transfer Protocol (SMTP)
 - ❖ SMTP is a delivery protocol only
- Like HTTP, SMTP has a similar set of ASCII commands and replies to transfer messages between machines
- Recipient of SMTP traffic must always be on
 - ✓ SMTP traffic can flow from:
 - ✓ Sending host → dedicated email server
 - ✓ Dedicated email server → another dedicated email server
 - ✗ SMTP traffic cannot flow from:
 - ✗ Dedicated email server → a receiving host
 - We use POP or IMAP for this

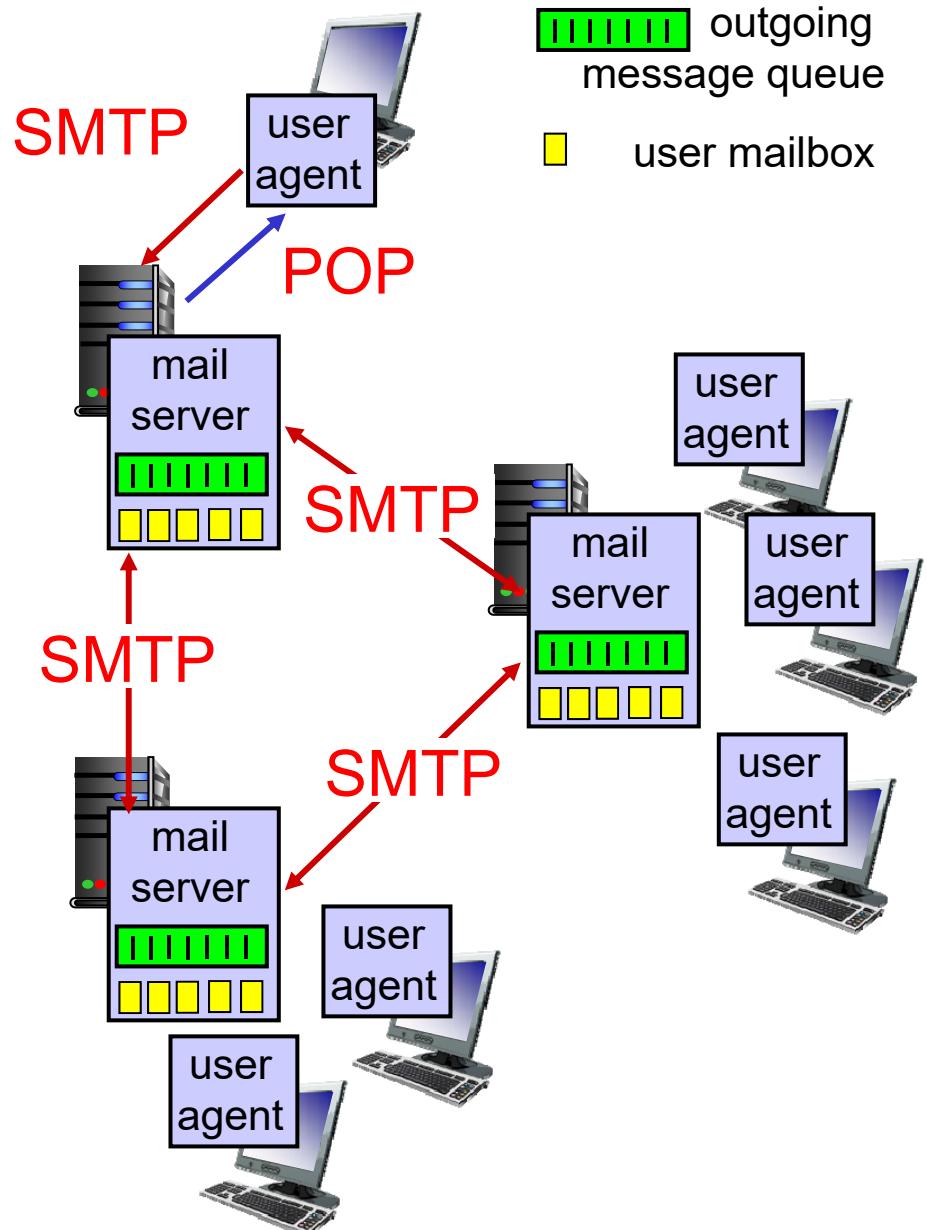
Electronic Mail

Three major components:

- User agents
- Mail servers
- Simple mail transfer protocol: SMTP

User Agent

- a.k.a. "mail reader"
- Composing, editing, reading mail messages
- e.g., Outlook, Mail (Linux)
- Outgoing, incoming messages stored on server



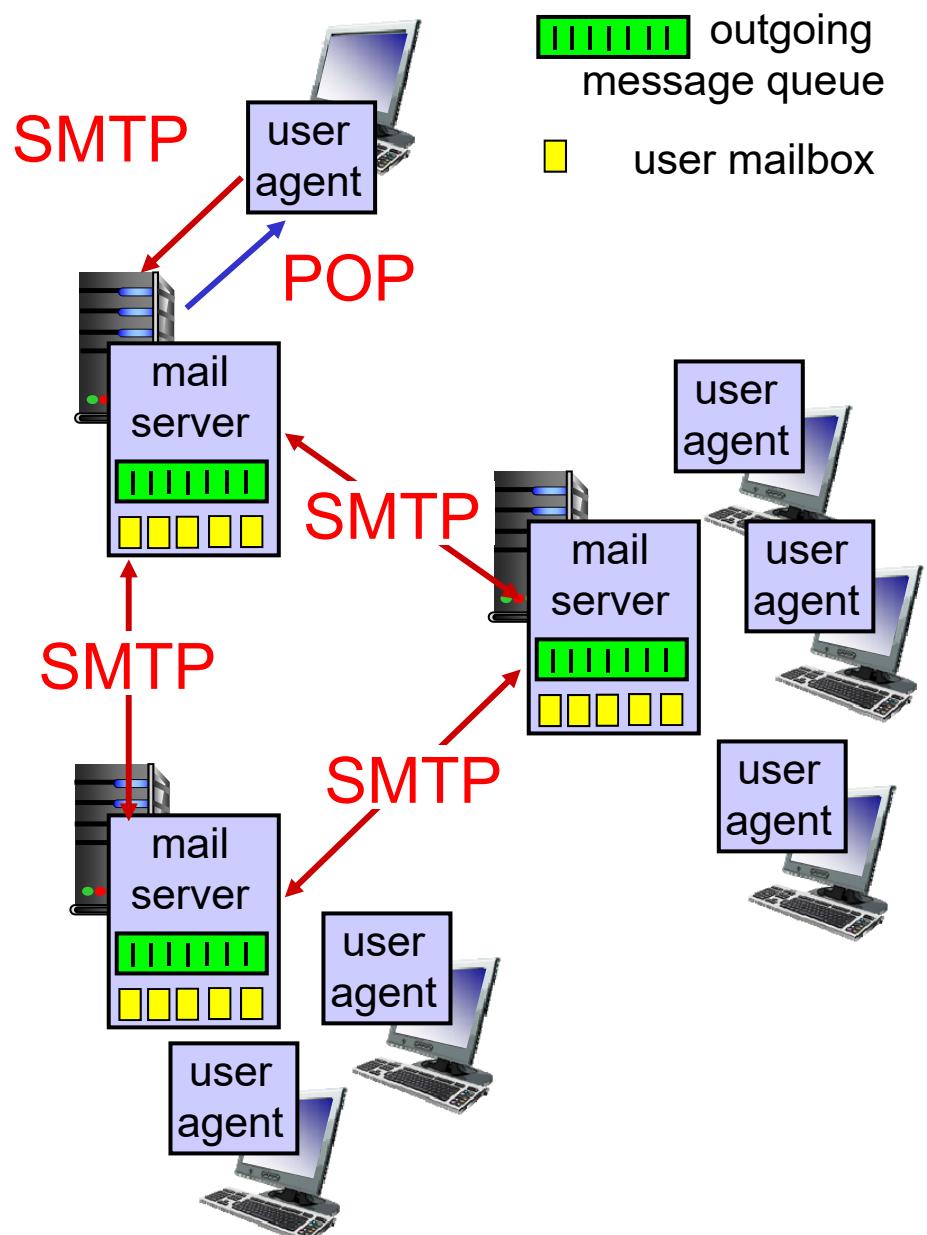
Electronic Mail

Mail Servers

- ❑ **Mailbox** contains incoming messages for user
- ❑ **Message queue** of outgoing (to be sent) mail messages

SMTP protocol between mail servers to send email messages

- ❖ "client": sending mail
- ❖ "server": receiving mail

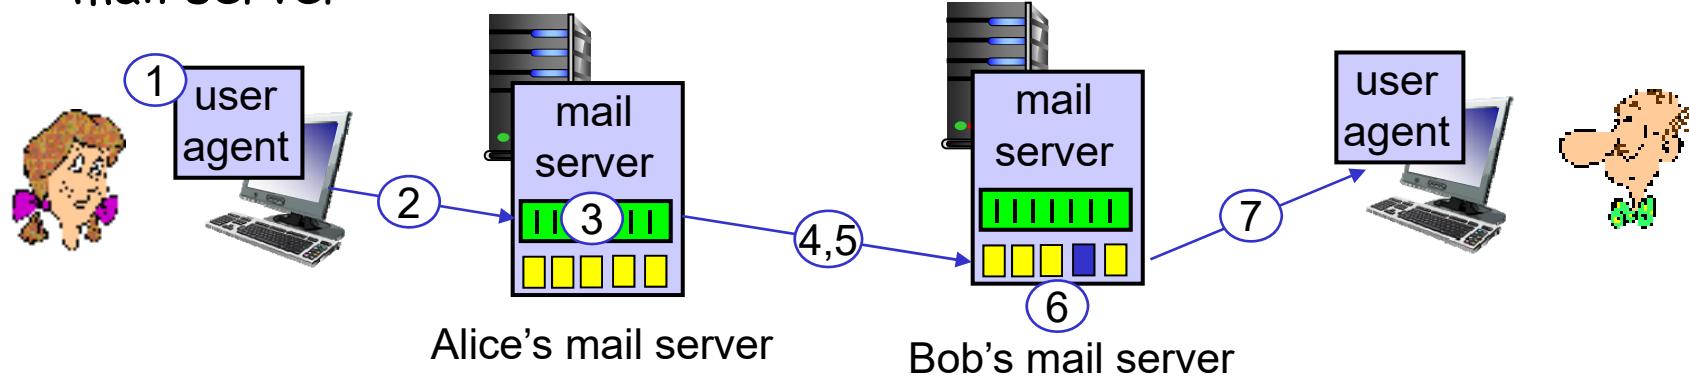


Electronic Mail: SMTP [RFC 822/2821]

- Uses TCP to transfer email message
 - ❖ From client to server port 25
 - ❖ From sending server to receiving server (direct transfer)
- Three phases of transfer
 - ❖ Handshaking (greeting)
 - ❖ Transfer of messages
 - ❖ Closure
- Command/Response interaction
 - ❖ Commands: ASCII text
 - ❖ Response: ASCII status code and phrase
- **Messages must be in 7-bit ASCII**

Scenario: Alice Sends Message to Bob

- 1) Alice uses UA to compose message "to" `bob@someschool.edu`
- 2) Alice's UA sends message to her mail server (SMTP or HTTP)
- 3) Message placed in message queue
- 4) Client side of SMTP opens TCP connection with Bob's mail server
- 5) SMTP client sends Alice's message over the TCP connection
- 6) Bob's mail server places the message in Bob's mailbox
- 7) Bob invokes his user agent to read message (POP, HTTP, IMAP)

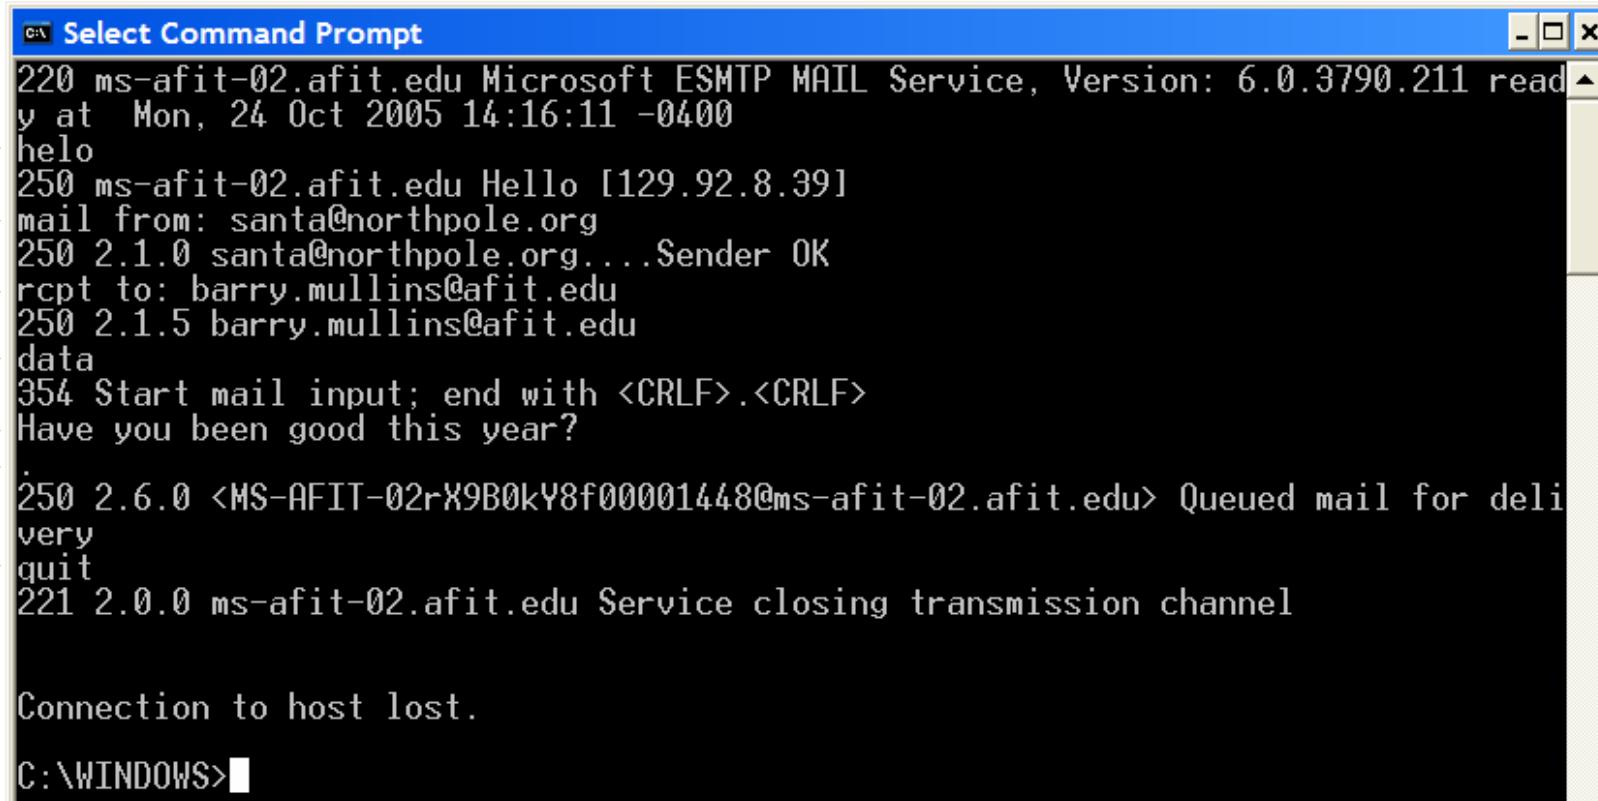


Sample SMTP Interaction

Assumes you are connected
to the AFIT network (i.e., VPN)

❑ telnet servername 25

❖ telnet ms-afit-02 25



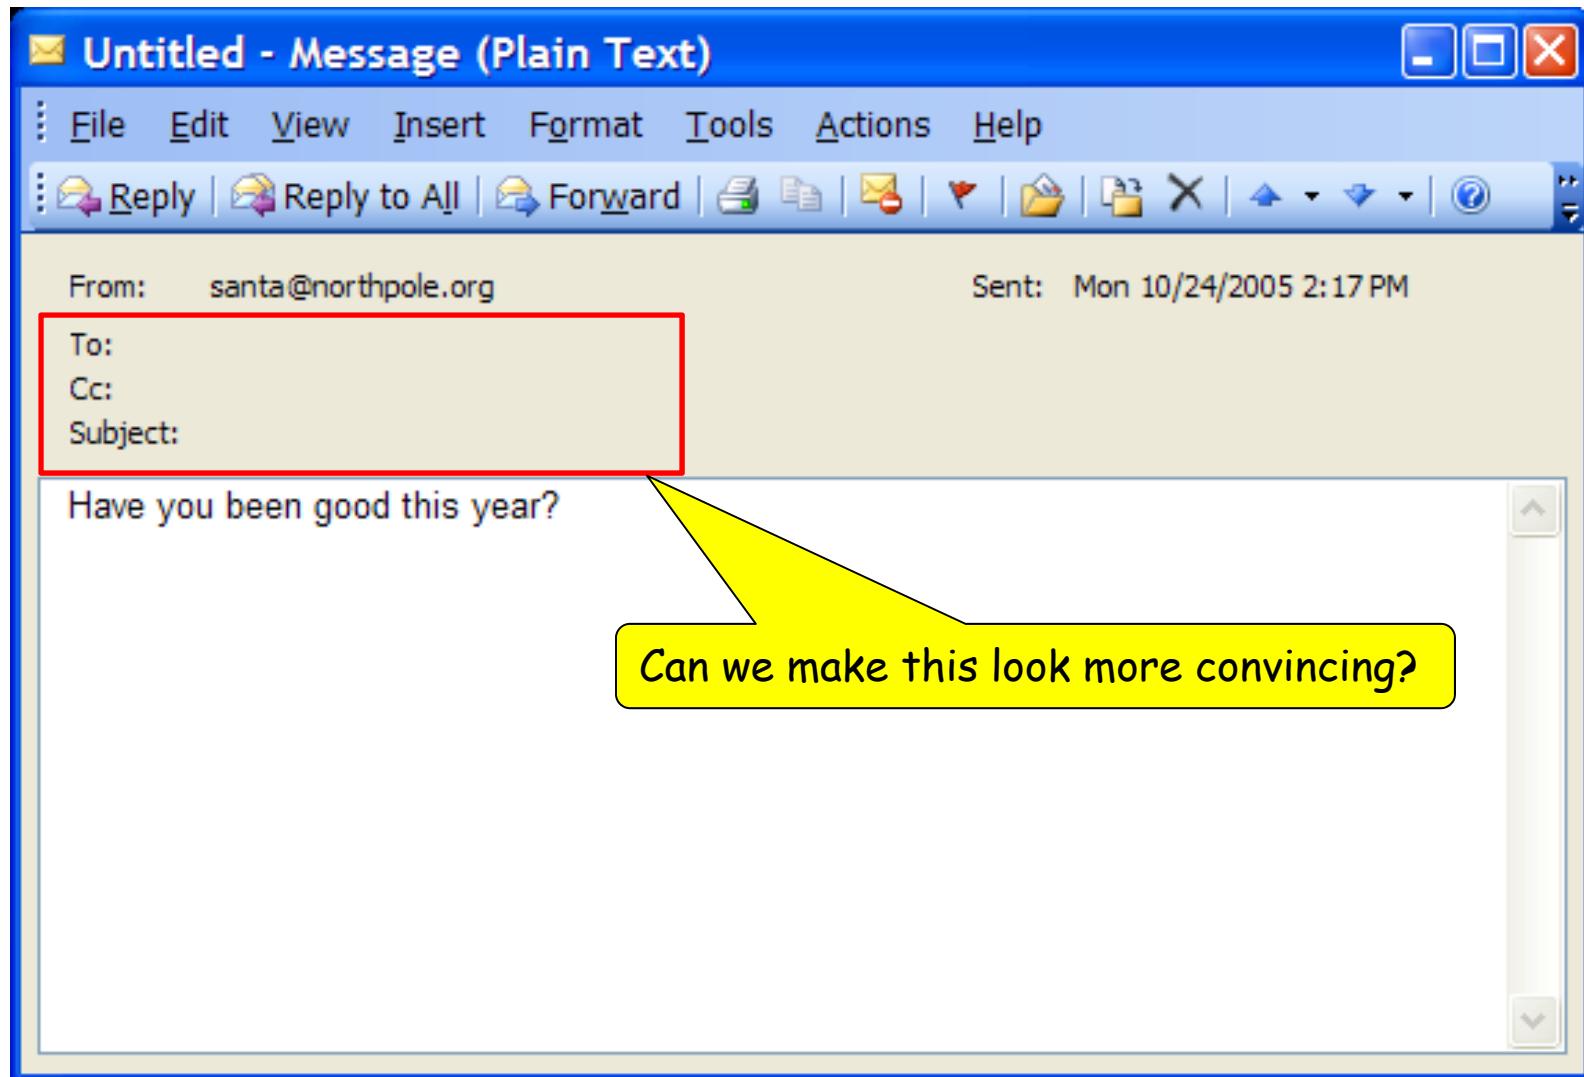
```
220 ms-afit-02.afit.edu Microsoft ESMTP MAIL Service, Version: 6.0.3790.211 ready at Mon, 24 Oct 2005 14:16:11 -0400
→ helo
→ 250 ms-afit-02.afit.edu Hello [129.92.8.39]
→ mail from: santa@northpole.org
→ 250 2.1.0 santa@northpole.org... Sender OK
→ rcpt to: barry.mullins@afit.edu
→ 250 2.1.5 barry.mullins@afit.edu
→ data
→ 354 Start mail input; end with <CRLF>.<CRLF>
→ Have you been good this year?
→ .
→ 250 2.6.0 <MS-AFIT-02rX9B0kY8f00001448@ms-afit-02.afit.edu> Queued mail for delivery
→ quit
→ 221 2.0.0 ms-afit-02.afit.edu Service closing transmission channel

Connection to host lost.

C:\WINDOWS>
```

Above lets you send email without using email client (reader)

Sample SMTP Interaction



Mail Message Format

RFC 822

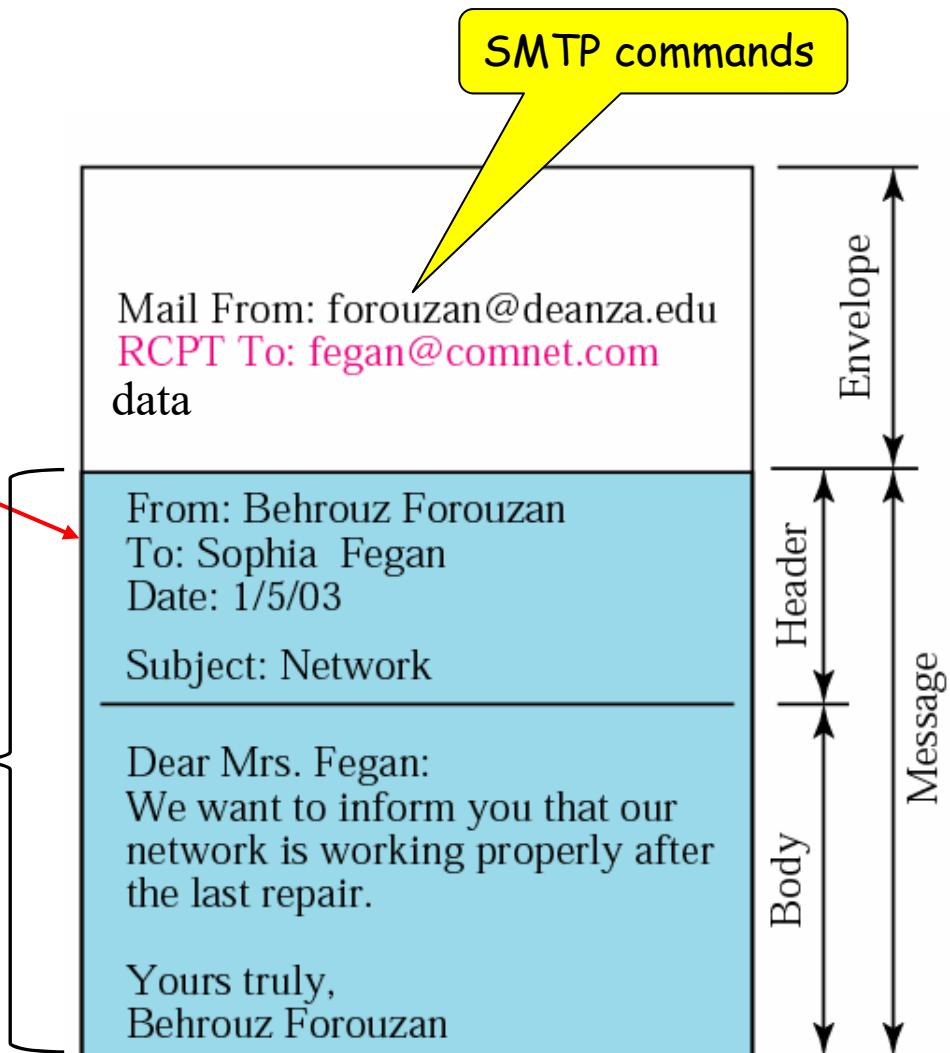
standard for text message format:

- Header lines, e.g.,

- ❖ To:
- ❖ From:
- ❖ Subject:

different from SMTP
commands!

- "Message"
 - ❖ ASCII characters only



An Example Mail Session

```
telnet ms-afit-02 25
```

```
he1o
```

```
mail from: anything
```

```
rcpt to: barry.mullins@afit.edu
```

```
data
```

```
to: drevil
```

```
from: Mullins Barry CSCE 560 rocks
```

```
subject: Request for Help
```

```
cc: Anyone else
```

```
date: 4 Jul 27 08:34 EDT
```

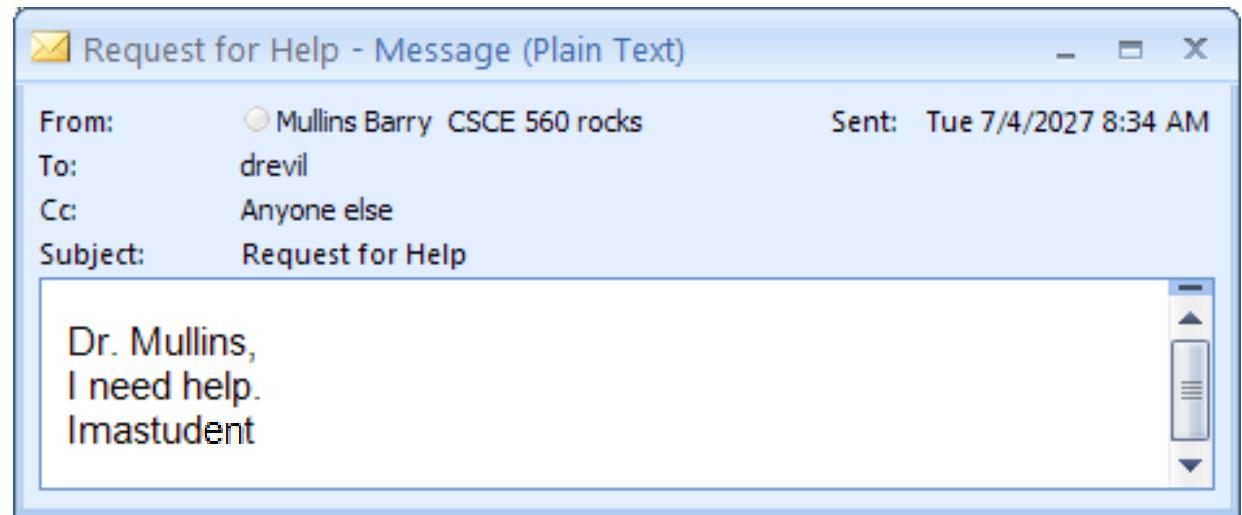
```
Dr. Mullins,
```

```
I need help.
```

```
Imastudent
```

```
.<CR><LF>
```

A period followed by the Enter key



Now you do it!

Your assignment is to send me a spoofed email. Include your real name in the email body. Also include the word "spoofed" somewhere in the body. If you don't, I won't know the email is from you and the department head may get involved.

Don't ask; it's a long story. ☺

SMTP: Comparison with HTTP

- ❑ HTTP: pull
- ❑ SMTP: push
- ❑ Both have ASCII command/response interaction, status codes
- ❑ HTTP: each object encapsulated in its own response msg
- ❑ SMTP: multiple objects sent in multipart msg

Message Format: Multimedia Extensions

- Multipurpose Internet Mail Extensions (MIME)
 - ❖ Multimedia mail extension, RFC 2045, 2046
- Additional lines in msg header declare MIME content type

The diagram illustrates the structure of a MIME message. On the left, four categories are listed with red arrows pointing to specific parts of the message on the right:

- MIME version → **MIME-Version: 1.0**
- method used to encode data → **Content-Transfer-Encoding: base64**
- multimedia data type, subtype, parameter declaration → **Content-Type: image/jpeg**
- encoded data → The base64 encoded data block, which contains three dots indicating continuation.

```
From: alice@crepes.fr
To: bob@hamburger.edu
Subject: Picture of yummy crepe.
MIME-Version: 1.0
Content-Transfer-Encoding: base64
Content-Type: image/jpeg

base64 encoded data .....
.....
.....base64 encoded data
```


Base64 Index Table

- Arrange these characters in an index table

| Index | Char | Index | Char | Index | Char | Index | Char |
|-------|------|-------|------|-------|------|-------|------|
| 0 | A | 16 | Q | 32 | g | 48 | w |
| 1 | B | 17 | R | 33 | h | 49 | x |
| 2 | C | 18 | S | 34 | i | 50 | y |
| 3 | D | 19 | T | 35 | j | 51 | z |
| 4 | E | 20 | U | 36 | k | 52 | 0 |
| 5 | F | 21 | V | 37 | l | 53 | 1 |
| 6 | G | 22 | W | 38 | m | 54 | 2 |
| 7 | H | 23 | X | 39 | n | 55 | 3 |
| 8 | I | 24 | Y | 40 | o | 56 | 4 |
| 9 | J | 25 | Z | 41 | p | 57 | 5 |
| 10 | K | 26 | a | 42 | q | 58 | 6 |
| 11 | L | 27 | b | 43 | r | 59 | 7 |
| 12 | M | 28 | c | 44 | s | 60 | 8 |
| 13 | N | 29 | d | 45 | t | 61 | 9 |
| 14 | O | 30 | e | 46 | u | 62 | + |
| 15 | P | 31 | f | 47 | v | 63 | / |

Base64 Example

- Text to encode: Man
- Encode using ASCII table
- Convert to binary
- Group into 6 bits and convert to decimal
- Look up decimal number in index table
- Send 7-bit printable, non-special ASCII character

77 4D M

| | |
|----|---|
| 19 | T |
| 20 | U |
| 21 | V |
| 22 | W |

84 54 T

| Source | Text (ASCII) | M | a | n |
|----------------|---|-----------|-----------|------------|
| | Octets | 77 (0x4d) | 97 (0x61) | 110 (0x6e) |
| Bits | 0 1 0 0 1 1 0 1 0 1 1 0 0 0 0 1 0 1 1 0 1 1 1 1 1 0 | | | |
| Base64 encoded | Sextets | 19 | 22 | 5 |
| | Character | T | W | F |
| | Octets | 84 (0x54) | 87 (0x57) | 70 (0x46) |
| | | | | 117 (0x75) |

MIME Types

Content-Type: type/subtype; parameters

Text

- Example subtypes:
 - ❖ plain, html

Image

- Example subtypes:
 - ❖ jpeg, gif

Audio

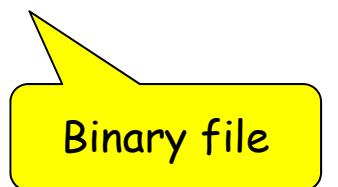
- Example subtypes:
 - ❖ basic (8-bit mu-law encoded)
 - ❖ 32kadpcm (32 kbps coding)

Video

- Example subtypes:
 - ❖ mpeg, quicktime

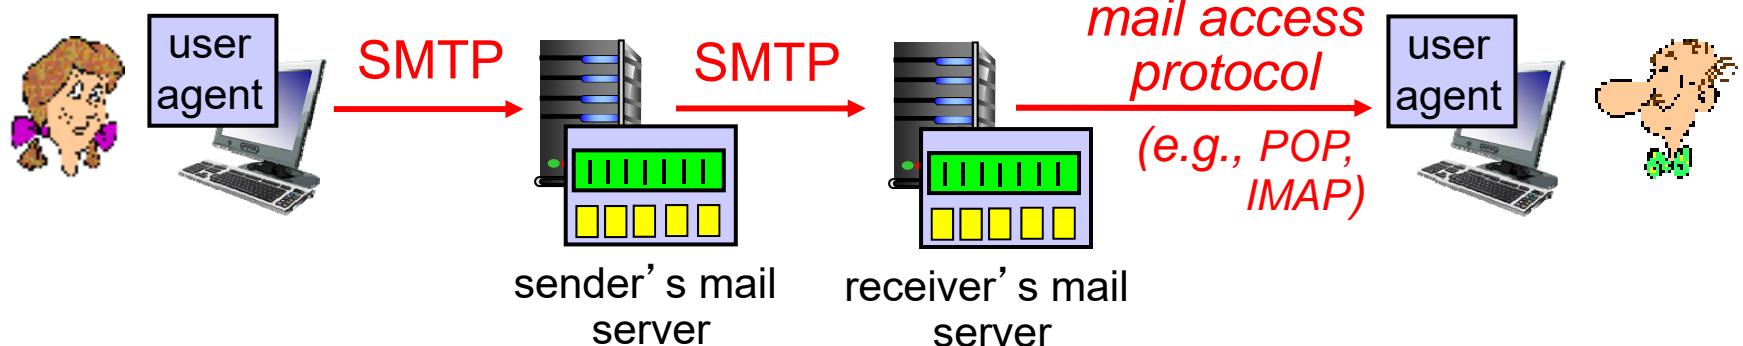
Application

- Other data that must be processed by reader before "viewable"
- Example subtypes:
 - ❖ msword, octet-stream



Binary file

Mail Access Protocols



- Mail access protocol: retrieve mail from server
 - ❖ POP: Post Office Protocol [RFC 1939]
 - Authorization (agent <--> server)
 - Download
 - ❖ IMAP: Internet Mail Access Protocol [RFC 1730]
 - More features (more complex)
 - Manipulation of stored msgs on server
 - ❖ HTTP: gmail, Hotmail, Yahoo! Mail, etc.

POP3 Protocol

Authorization phase

- Client commands:
 - ❖ user: declare username
 - ❖ pass: password
- Server responses
 - ❖ +OK
 - ❖ -ERR

Transaction phase, client:

- list: list message numbers
- retr: retrieve message by number
- dele: delete
- Quit

Update phase

- Server deletes marked messages

```
telnet mailserver 110
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on
C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

Sent in plain text!

POP3 (more) and IMAP

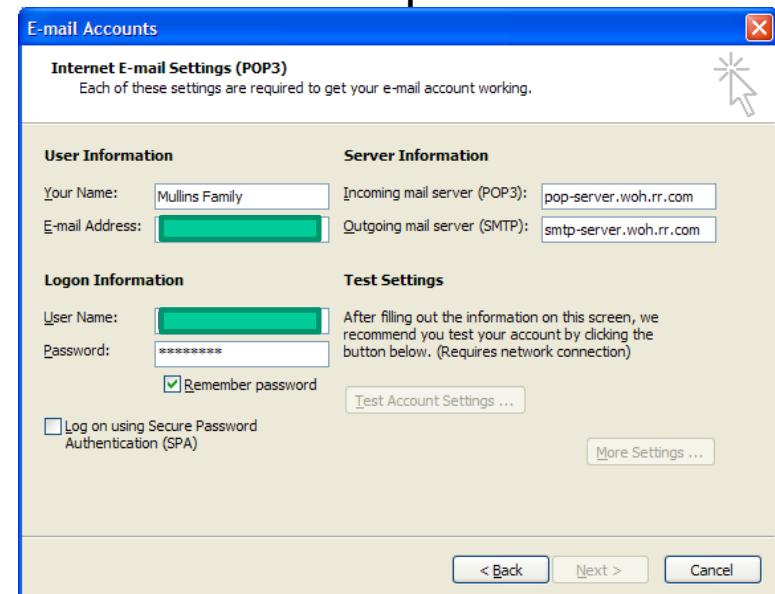
More about POP3

- Previous example uses "download-and-delete" mode
- Bob cannot re-read email if he changes client
- "Download-and-keep": copies of messages on different clients
- POP3 is stateless across sessions
- Server listens on port 110

- Use SMTP to **send mail to your mail server**
- Use POP to **download mail from the mail server**

IMAP

- Keep all messages in 1 place: server
- Allows user to organize messages in folders
- IMAP keeps user state across sessions:
 - ❖ Names of folders and mappings between message IDs and folder name
- Server listens on port 143



IMAP - Internet Message Access Protocol

Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications

Domain Name System (DNS) Motivation

□ Problem statement:

- ❖ Internet hosts and routers identified by:
 - 32-bit IP address (e.g., 129.92.6.150) - up to 12 digits
 - Used for addressing datagrams
- ❖ Average brain can easily remember 7 digits
- ❖ We need an easy way to remember IP addresses or provide a mapping between IP addresses and host names

□ Solution:

- ❖ Use alphanumeric names (e.g., www.yahoo.com) to refer to hosts
- ❖ Add a **distributed, hierarchical protocol** (called DNS) to map between alphanumeric host names and binary IP addresses
- ❖ We call this *Address Resolution* →

DNS: Domain Name System

Domain Name System is a:

- ❑ Distributed database implemented in a hierarchy of many name servers
- ❑ Application-layer protocol that allows hosts to query the distributed database to resolve names (address/name translation)
- ❑ Interesting note: A core Internet function is implemented as an application-layer protocol
 - ❖ Complexity pushed to network's "edge"
- ❑ Address resolution adds additional delay to an application
 - ❖ Must find correct IP address before setting up TCP connection

DNS Services

- Hostname to IP address translation
- Host aliasing
 - ❖ Canonical and alias names
 - ❖ Relay1.west-coast.enterprise.com = enterprise.com
 - ❖ Relay1.west-coast.enterprise.com = www.enterprise.com
- Mail server aliasing
 - ❖ We prefer bob@hotmail.com instead of bob@relay1.west-coast.hotmail.com
- Load distribution
 - ❖ Replicated web servers: set of IP addresses for one canonical name
 - ❖ DNS responds with the list of IP addresses but rotates order with each request

DNS in Action - nslookup

```
C:\WINDOWS>nslookup www.afit.edu
Server: dc-afit-01.afit.edu
Address: 129.92.1.34
Name: www.afit.edu
Address: 129.92.253.61
From authoritative server

C:\WINDOWS>nslookup www.ebay.com
Server: dc-afit-01.afit.edu
Address: 129.92.1.34
Non-authoritative answer:
Name: hp.core.ebay.com
Addresses: 66.135.208.89, 66.135.192.124, 66.135.192.123, 66.135.208.90
Aliases: www.ebay.com
Not from authoritative server – from a cache

C:\WINDOWS>nslookup www.ebay.com
Server: dc-afit-01.afit.edu
Address: 129.92.1.34
Non-authoritative answer:
Name: hp.core.ebay.com
Addresses: 66.135.192.124, 66.135.192.123, 66.135.208.90, 66.135.208.89
Aliases: www.ebay.com
Load distribution

C:\WINDOWS>nslookup www.ebay.com
Server: dc-afit-01.afit.edu
Address: 129.92.1.34
Non-authoritative answer:
Name: hp.core.ebay.com
Addresses: 66.135.192.123, 66.135.208.90, 66.135.208.89, 66.135.192.124
Aliases: www.ebay.com
```

DNS - Basically How Does It Work?

- Your application needs to setup a connection with a server (www.mit.edu)
- Since your app does not already know the IP address of the destination, your app invokes DNS
 - ❖ Requests an IP address for a host name (www.mit.edu)
- DNS on your host
 - ❖ Checks HOSTS file for static entry
 - 74.125.228.116 google.com
 - 127.0.0.1 annoyingadsite.com
 - ❖ If not in HOSTS file, checks DNS cache on host
 - ipconfig /displaydns
- If no hit on host, sends a query “**into the network**” using UDP port 53
 - ❖ Waits ms to seconds for the DNS response
- IP address is passed to calling application and used to initiate connection to the server

C:\Windows\System32\drivers\etc
/etc/hosts

DNS

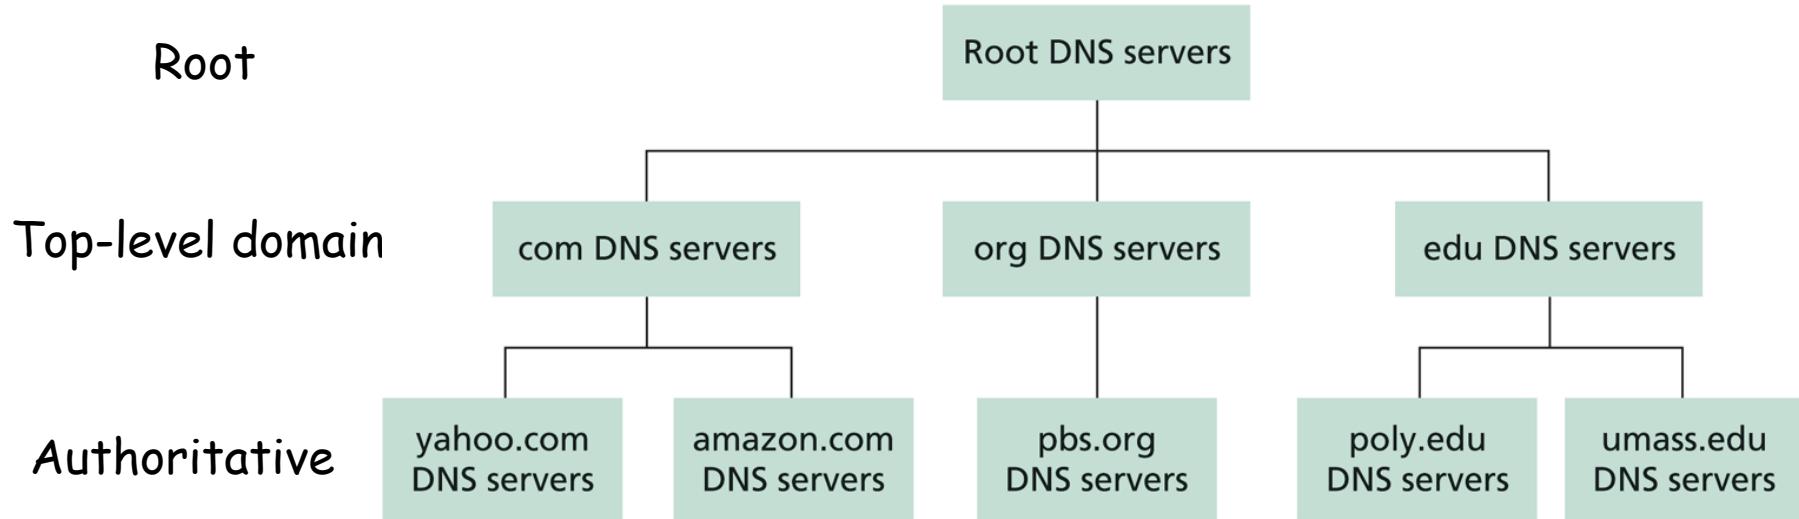
- Where is this “**into the network**”?
 - ❖ A single centralized DNS server?

NO!!

Why not centralize DNS?

- Single point of failure
- Traffic volume
- Distant centralized database
 - ❖ Distance does matter - remember propagation delay?
- Maintenance
 - ❖ Database management would be unreal
- Doesn't scale!

Distributed, Hierarchical Database



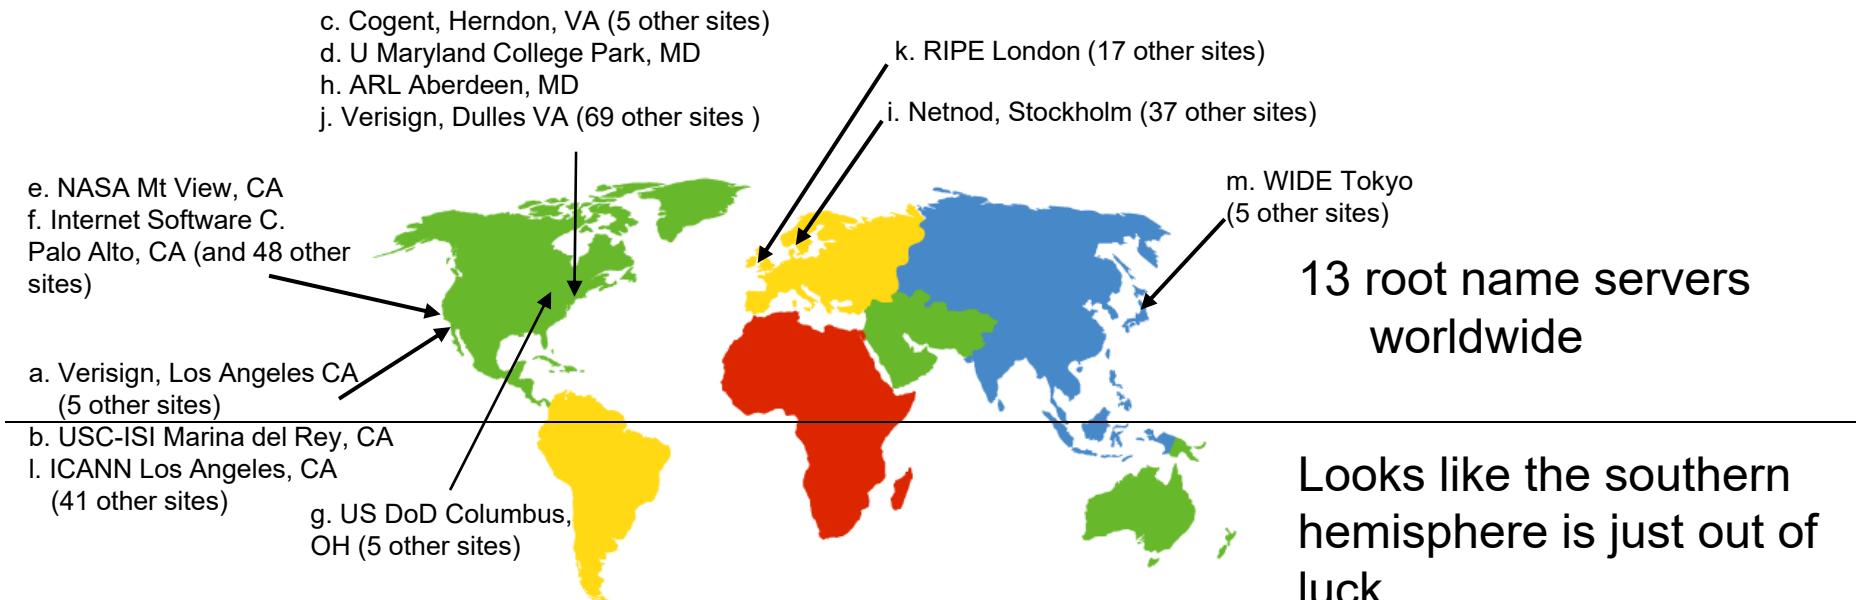
- ❑ No single DNS server has all the mappings!!

Client wants IP address for www.amazon.com; 1st approx:

- ❑ Client queries a root server to find .com DNS server
- ❑ Client queries .com DNS server to get amazon.com DNS server
- ❑ Client queries amazon.com DNS server to get IP address for www.amazon.com

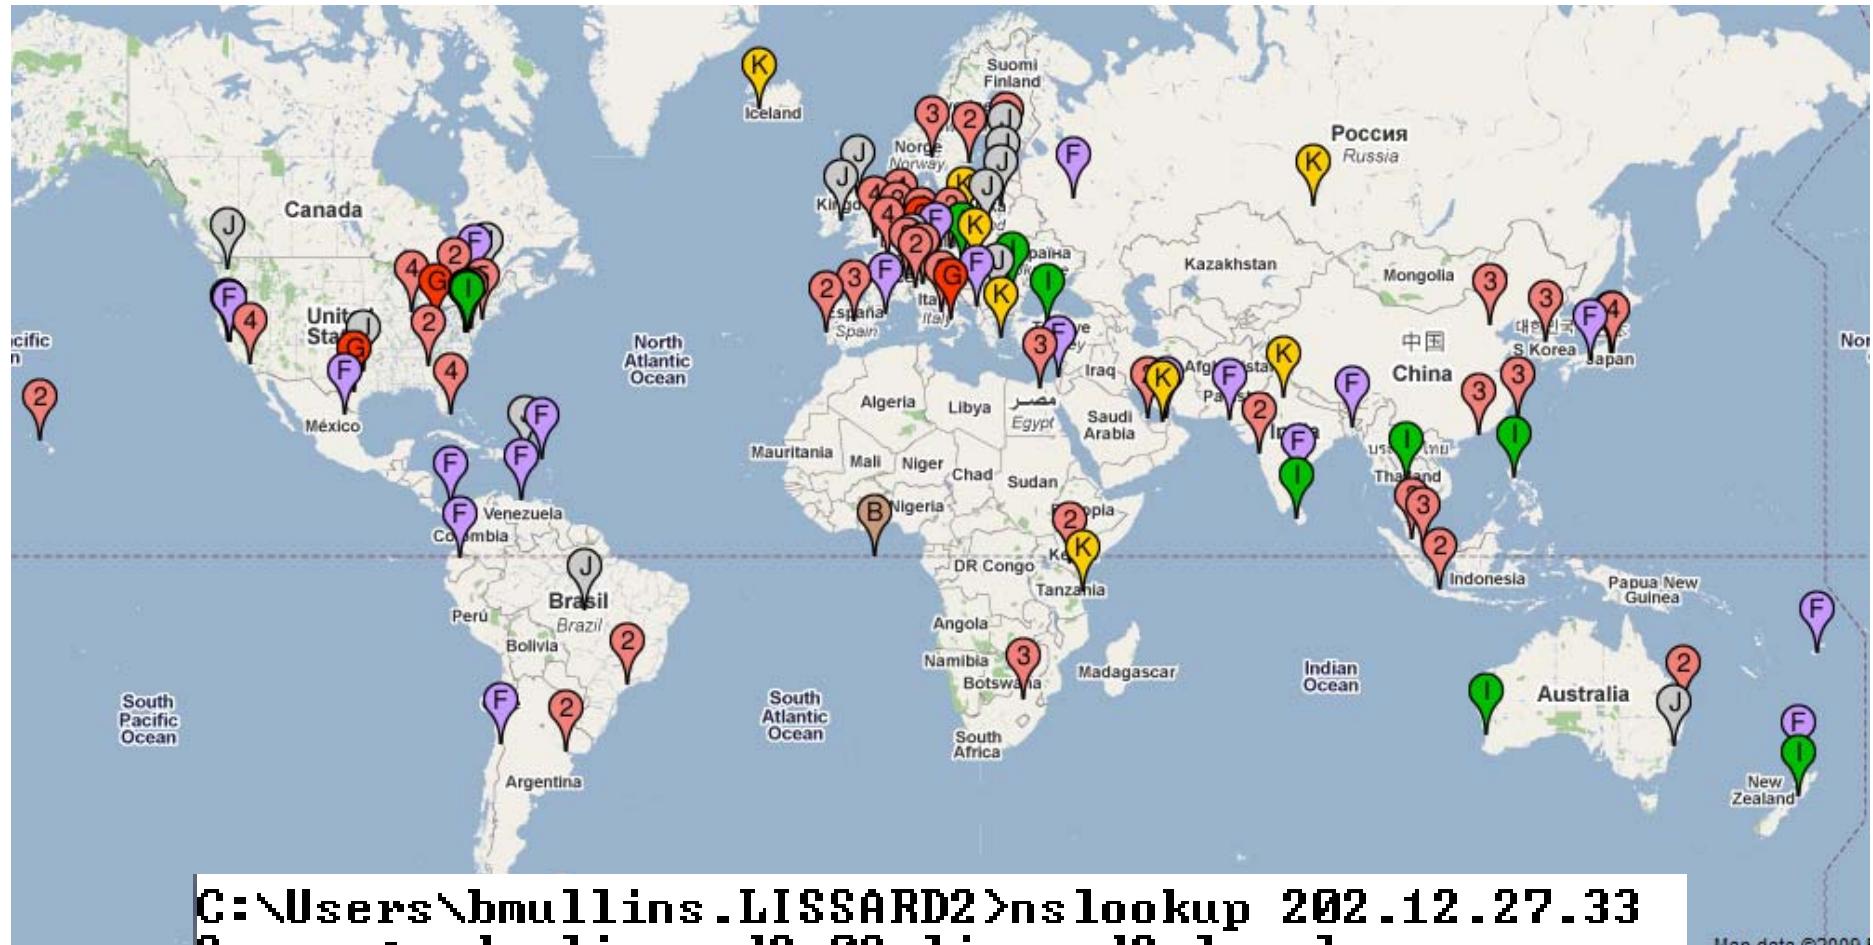
DNS: Root Name Servers

- **ICANN: Internet Corporation for Assigned Names and Numbers**
 - ❖ Manages DNS root servers
 - ❖ IP addresses → <https://www.iana.org/domains/root/servers>
 - ❖ Names → **k.root-servers.net**



Remember
propagation delay?

DNS: Root Name Servers Replication



```
C:\Users\bmullins.LISSARD2>nslookup 202.12.27.33
Server: dc-lissard2-02.lissard2.local
Address: 10.1.2.2
```

```
Name: m.root-servers.net
Address: 202.12.27.33
```

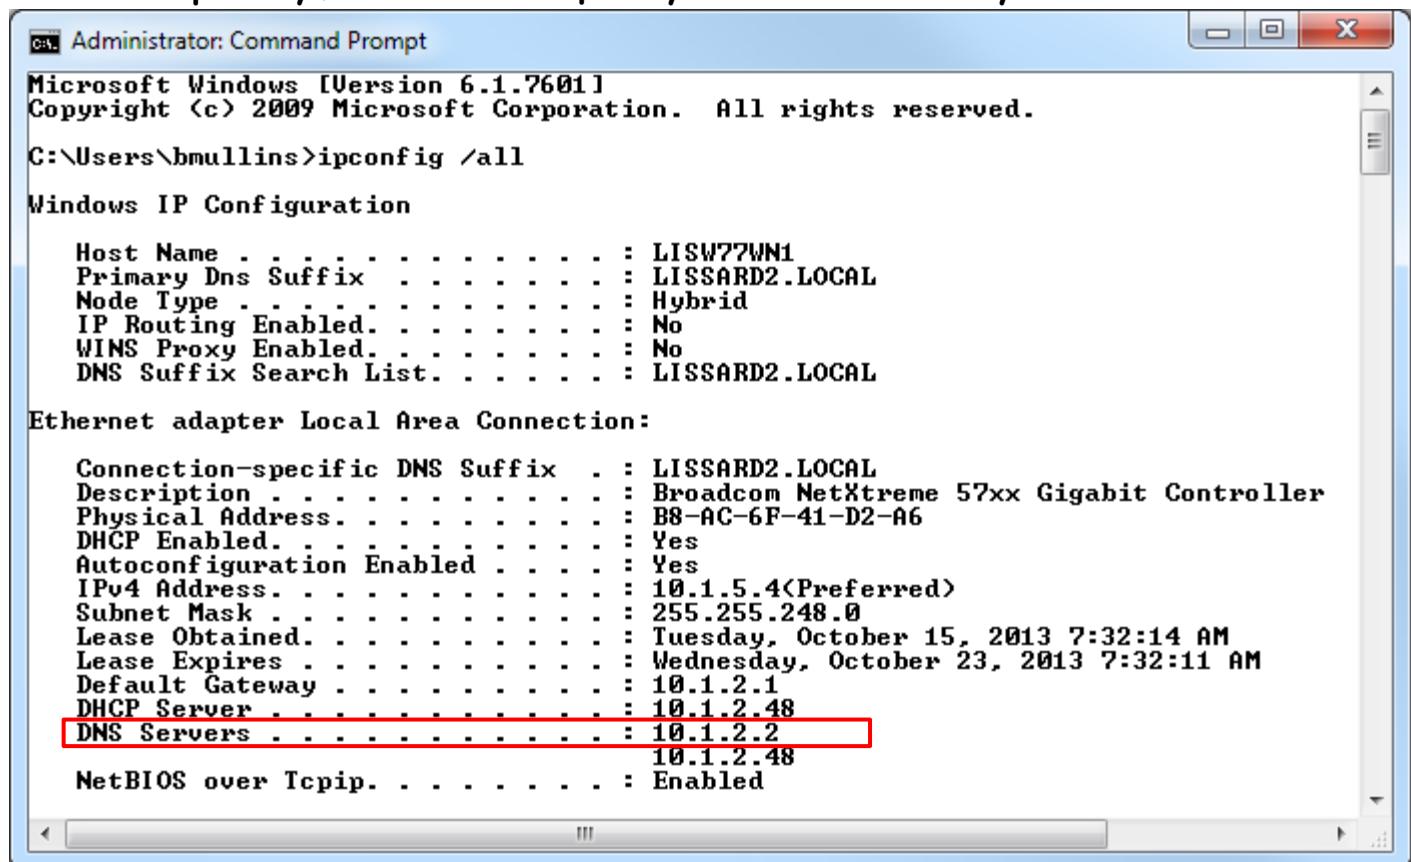
TLD and Authoritative Servers

- Top-level domain (TLD) servers:
 - ❖ Responsible for .com, .org, .net, .edu, .mil, .gov, etc, and all top-level country domains .uk, .fr, .ca, .jp
 - ❖ Network Solutions maintains TLD servers for .com
 - ❖ Educause maintains .edu TLD
 - ❖ As of June 2011, ICANN will allow any TLD → .barrymullins
 - At that time, only cost \$185,000 ☺

- Authoritative DNS servers:
 - ❖ Organization's DNS servers provide authoritative hostname to IP mappings for organization's servers (e.g., Web and mail)
 - ❖ Can be maintained by organization or service provider

Local Name (DNS) Server

- ❑ Does not strictly belong within the hierarchy
- ❑ Each ISP (residential ISP, company, university) has one though
 - ❖ Also called "default name server"
- ❑ When host makes a DNS query, query is sent to its local DNS server
 - ❖ Server acts as a proxy, forwards query into hierarchy



```
c:\ Administrator: Command Prompt
Microsoft Windows [Version 6.1.7601]
Copyright <c> 2009 Microsoft Corporation. All rights reserved.

C:\Users\bmullins>ipconfig /all

Windows IP Configuration

Host Name . . . . . : LISW77WN1
Primary Dns Suffix . . . . . : LISSARD2.LOCAL
Node Type . . . . . : Hybrid
IP Routing Enabled. . . . . : No
WINS Proxy Enabled. . . . . : No
DNS Suffix Search List. . . . . : LISSARD2.LOCAL

Ethernet adapter Local Area Connection:

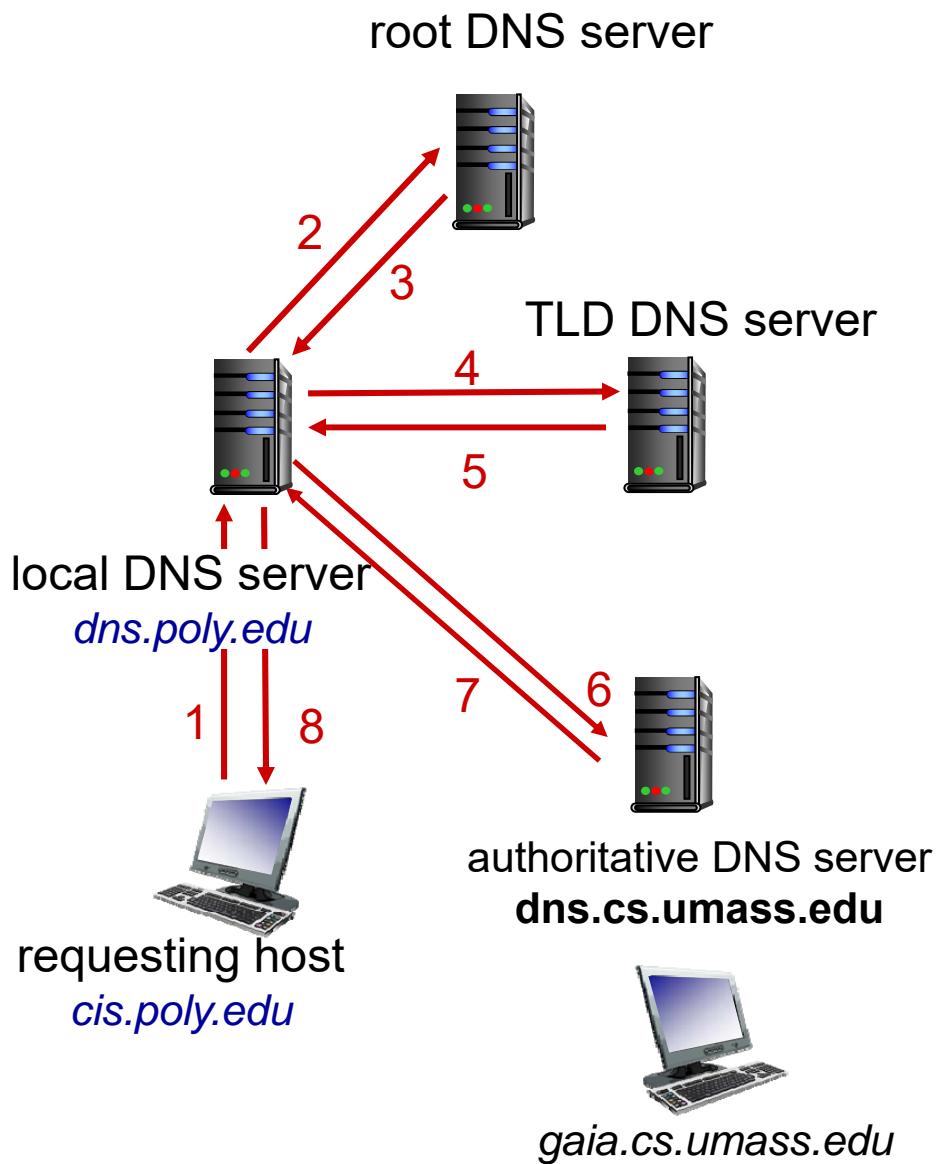
Connection-specific DNS Suffix . . . . . : LISSARD2.LOCAL
Description . . . . . : Broadcom NetXtreme 57xx Gigabit Controller
Physical Address. . . . . : B8-AC-6F-41-D2-A6
DHCP Enabled. . . . . : Yes
Autoconfiguration Enabled . . . . . : Yes
IPv4 Address . . . . . : 10.1.5.4<Preferred>
Subnet Mask . . . . . : 255.255.248.0
Lease Obtained. . . . . : Tuesday, October 15, 2013 7:32:14 AM
Lease Expires . . . . . : Wednesday, October 23, 2013 7:32:11 AM
Default Gateway . . . . . : 10.1.2.1
DHCP Server . . . . . : 10.1.2.48
DNS Servers . . . . . : 10.1.2.2
                                         10.1.2.48
NetBIOS over Tcpip. . . . . : Enabled
```

Example

- ❑ Host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

Iterated query:

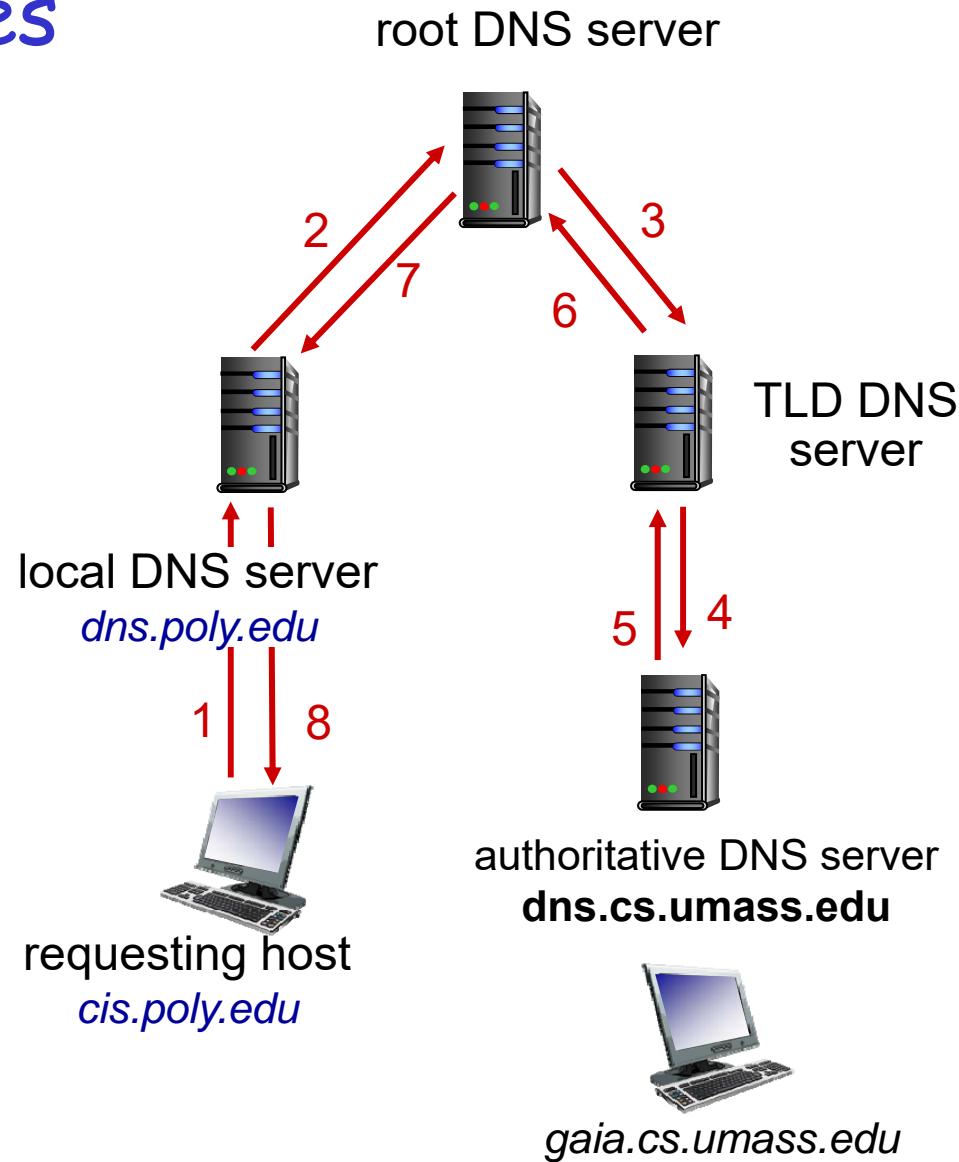
- ❑ Contacted server replies with name of server to contact
 - ❖ “I don’t know this name, but ask this server”
- ❑ In practice, use this pattern
 - ❖ One recursive
 - (1,8)
 - ❖ and subsequent iterative
 - 2 through 7



Recursive Queries

Recursive query:

- Query asks server to obtain mapping on its behalf
- Puts burden of name resolution on contacted name server
- Heavy load?



DNS: Caching and Updating Records

- Once (any) name server learns mapping, it *caches* mapping
 - ❖ Cache entries timeout (disappear) after some time (TTL)
 - Time varies greatly
 - ❖ TLD servers typically cached in local name servers
 - Thus root name servers not often visited
- Your host also maintains a DNS cache
 - ❖ `ipconfig /displaydns`
- Can delete current cache
 - ❖ `ipconfig /flushdns`

DNS Records

DNS: distributed db storing resource records (RR)

RR format: `(name, value, type, ttl)`

- ❑ Type=A
 - ❖ **name** is hostname
 - ❖ **value** is IPv4 address
 - ❖ (relay1.bar.foo.com, 145.37.93.126, A)
 - ❖ IPv6 uses a AAAA record
- ❑ Type=NS
 - ❖ **name** is domain (e.g., foo.com)
 - ❖ **value** is IP address of authoritative name server for this domain
 - ❖ (foo.com, dns.foo.com, NS)
- ❑ Type=CNAME
 - ❖ **name** is alias name for some “canonical” (the real) name
 - ❖ www.ibm.com is really servereast.backup2.ibm.com
 - ❖ **value** is canonical name
 - ❖ (foo.com, relay1.bar.foo.com, CNAME)
- ❑ Type=MX
 - ❖ **value** is name of mailserver associated with **name**
 - ❖ (foo.com, mail.bar.foo.com, MX)

DNS Protocol, Messages

Query & reply messages both have the same message format

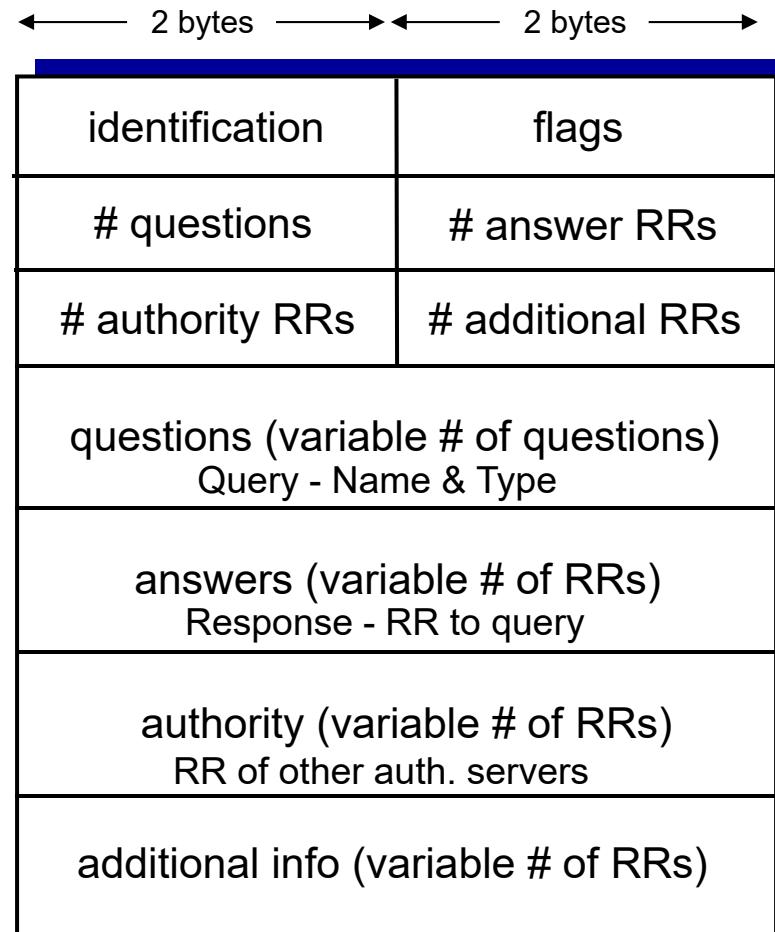
msg header

□ Identification:

- ❖ 16 bit # for query
- ❖ reply to query uses same #

□ Flags:

- ❖ query or reply
- ❖ recursion desired
- ❖ recursion available
- ❖ reply is authoritative



DNS Protocol, Message Example

C:\>nslookup www.mit.edu
Server: dc-lissard2-02.lissard2.local
Address: 10.1.2.2

Non-authoritative answer:
Name: e9566.b.akamaiedge.net
Address: 23.212.28.78
Aliases: www.mit.edu
www.mit.edu.edgekey.net

Next slide

42 1.068984000 10.1.0.65 10.1.2.2 DNS 71 Standard query 0x0004 A www.mit.edu

+ Frame 42: 71 bytes on wire (568 bits), 71 bytes captured (568 bits) on interface 0

+ Ethernet II, Src: b8:ac:6f:41:d2:a6 (b8:ac:6f:41:d2:a6), Dst: 00:50:56:a9:27:8e (00:50:56:a9:27:8e)

+ Internet Protocol version 4, Src: 10.1.0.65 (10.1.0.65), Dst: 10.1.2.2 (10.1.2.2)

+ User Datagram Protocol, Src Port: 59697 (59697), Dst Port: 53 (53)

+ Domain Name System (query)
[Response In: 45]

Transaction ID: 0x0004

+ Flags: 0x0100 Standard query

Questions: 1

Answer RRs: 0

Authority RRs: 0

Additional RRs: 0

+ Queries
+ www.mit.edu: type A, class IN
Name: www.mit.edu
[Name Length: 11]
[Label Count: 3]
Type: A (Host Address) (1)
class: IN (0x0001)

| | | |
|------|---|---------------------|
| 0000 | 00 50 56 a9 27 8e b8 ac 6f 41 d2 a6 08 00 45 00 | .PV. OA.... E. |
| 0010 | 00 39 17 80 00 00 80 11 00 00 0a 01 00 41 0a 01 | .9..... A.. |
| 0020 | 02 02 e9 31 00 35 00 25 16 7b 00 04 01 00 00 01 | ...1.5% .{..... |
| 0030 | 00 00 00 00 00 00 03 77 77 77 03 6d 69 74 03 65 | w ww.mit.e |
| 0040 | 64 75 00 00 01 00 01 | du..... |

DNS Protocol, Message Example

45 1.322372000 10.1.2.2 10.1.0.65 DNS 157 Standard query response 0x0004 CNAME www.mit.edu.edgekey.net C...

+ Frame 45: 157 bytes on wire (1256 bits), 157 bytes captured (1256 bits) on interface 0

+ Ethernet II, Src: 00:50:56:a9:27:8e (00:50:56:a9:27:8e), Dst: b8:ac:6f:41:d2:a6 (b8:ac:6f:41:d2:a6)

+ Internet Protocol Version 4, Src: 10.1.2.2 (10.1.2.2), Dst: 10.1.0.65 (10.1.0.65)

+ User Datagram Protocol, src Port: 53 (53), Dst Port: 59697 (59697)

Domain Name System (response)

[Request In: 42] [Time: 0.253388000 seconds] This is generated by Wireshark.
Transaction ID: 0x0004 Not actually in the frame

Flags: 0x8180 Standard query response, No error

Questions: 1

Answer RRs: 3

Authority RRs: 0

Additional RRs: 0

Queries

www.mit.edu: type A, class IN

Name: www.mit.edu
[Name Length: 11]
[Label Count: 3]
Type: A (Host Address) (1)
Class: IN (0x0001)

Answers

www.mit.edu: type CNAME, class IN, cname www.mit.edu.edgekey.net

Name: www.mit.edu
Type: CNAME (Canonical NAME for an alias) (5)
Class: IN (0x0001)
Time to live: 1799
Data length: 25
CNAME: www.mit.edu.edgekey.net

www.mit.edu.edgekey.net: type CNAME, class IN, cname e9566.b.akamaiedge.net

e9566.b.akamaiedge.net: type A, class IN, addr 23.212.28.78

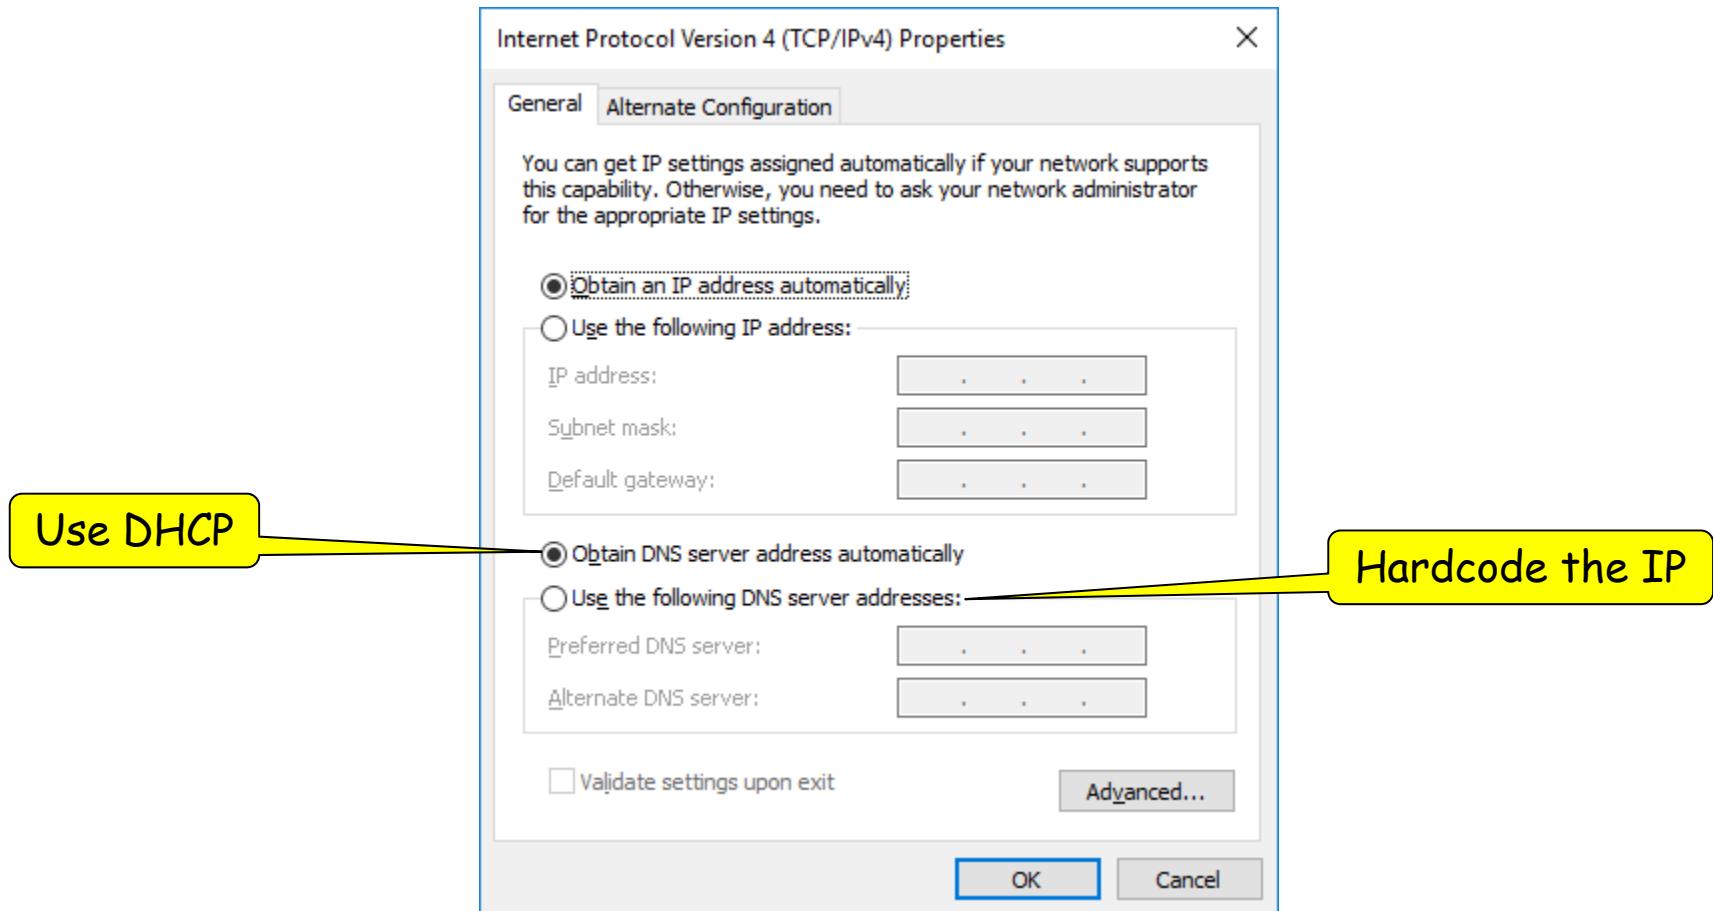
| Hex | Dec | ASCII |
|------|---|--------------------|
| 0000 | b8 ac 6f 41 d2 a6 00 50 56 a9 27 8e 08 00 45 00 | .oA...P V.'...E. |
| 0010 | 00 8f 4c 6e 40 00 80 11 97 ab 0a 01 02 02 0a 01 | ..Ln@... |
| 0020 | 00 41 00 35 e9 31 00 7b a7 0b 00 04 81 80 00 01 | .A.5.1.{ |
| 0030 | 00 03 00 00 00 00 03 77 77 77 03 6d 69 74 03 65 |w ww.mit.e |
| 0040 | 64 75 00 00 01 00 01 c0 0c 00 05 00 01 00 00 07 | du..... |
| 0050 | 07 00 19 03 77 77 03 6d 69 74 03 65 64 75 07 |www.mit.edu. |
| 0060 | 65 64 67 65 6b 65 79 03 6e 65 74 00 c0 29 00 05 | edgekey.net..). |
| 0070 | 00 01 00 00 00 3b 00 15 05 65 39 35 36 36 01 62 |;. e9566.b |
| 0080 | 0a 61 6b 61 6d 61 69 65 64 67 65 c0 3d c0 4e 00 | .akamaie.dae.=.N. |

Inserting Records into DNS

- Example: just created startup company "Network Utopia"
- Register name **networkuptopia.com** at a **DNS registrar** (e.g., GoDaddy, Network Solutions)
 - ❖ Need to provide registrar with names and IP addresses of your authoritative name server (primary and secondary)
 - ❖ Registrar inserts two RRs into the .com TLD servers:
(networkutopia.com, dns1.networkutopia.com, NS) ← alias
(dns1.networkutopia.com, 212.212.212.1, A) ← mapping
- Put in authoritative server
 - ❖ Type A record for www.networkuptopia.com
 - ❖ Type MX record for networkutopia.com mail server

Bootstrapping DNS

- ❑ How does a host contact the name server if all it has is the name and no IP address?
 - ❖ IP address of at least 1 name server must be given a priori
 - Hardcode the IP or use DHCP



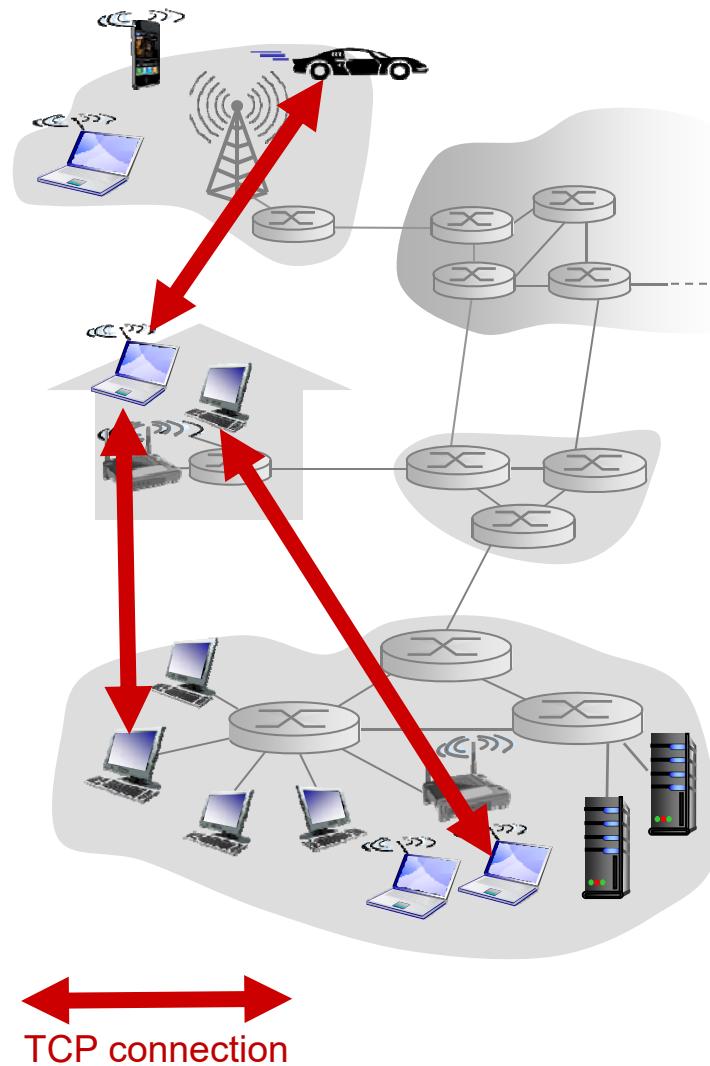
Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming: Creating Network Applications

Pure P2P Architecture

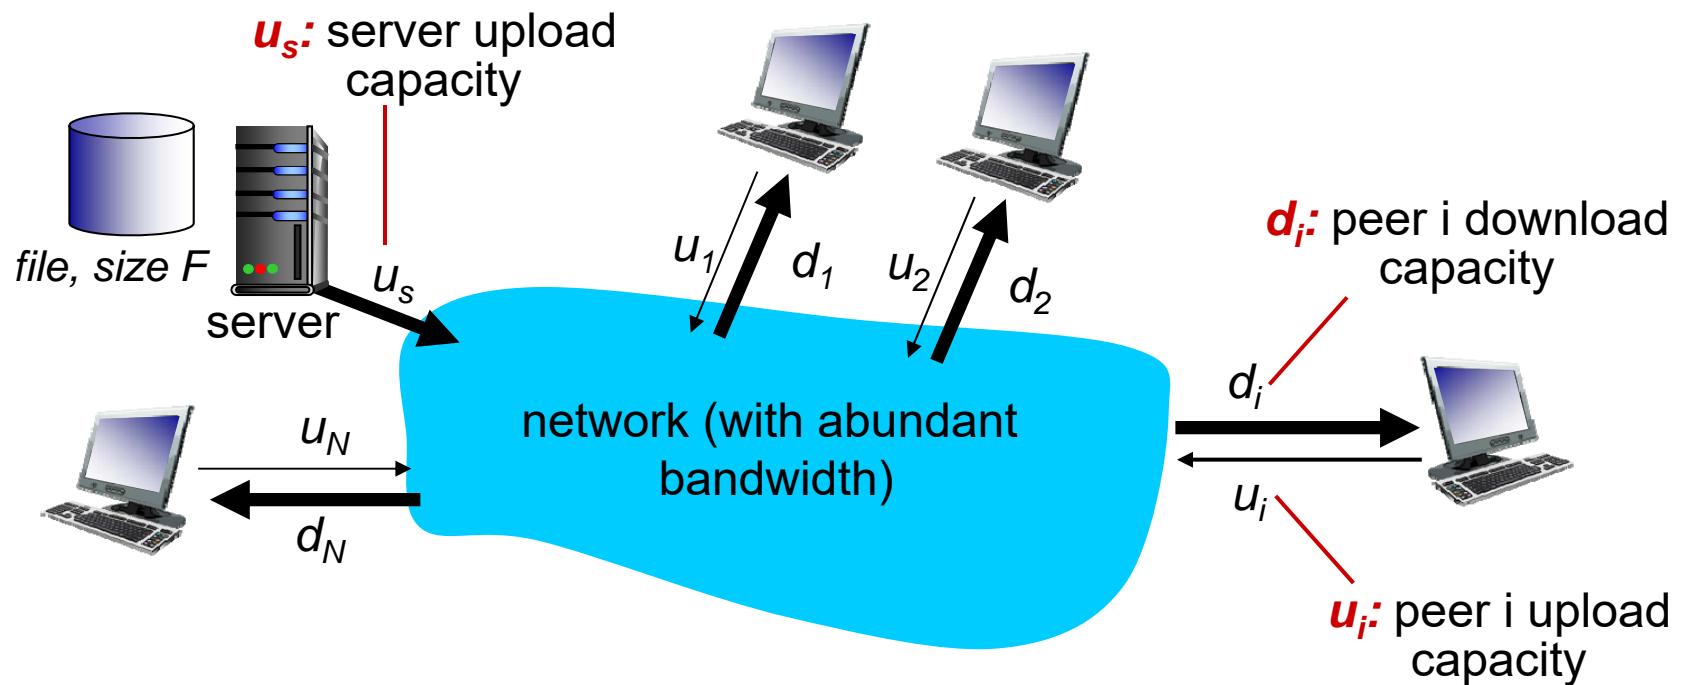
- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and change IP addresses
- Examples:
 - ❖ File distribution (BitTorrent)
 - ❖ VoIP (Skype)

Highly scalable



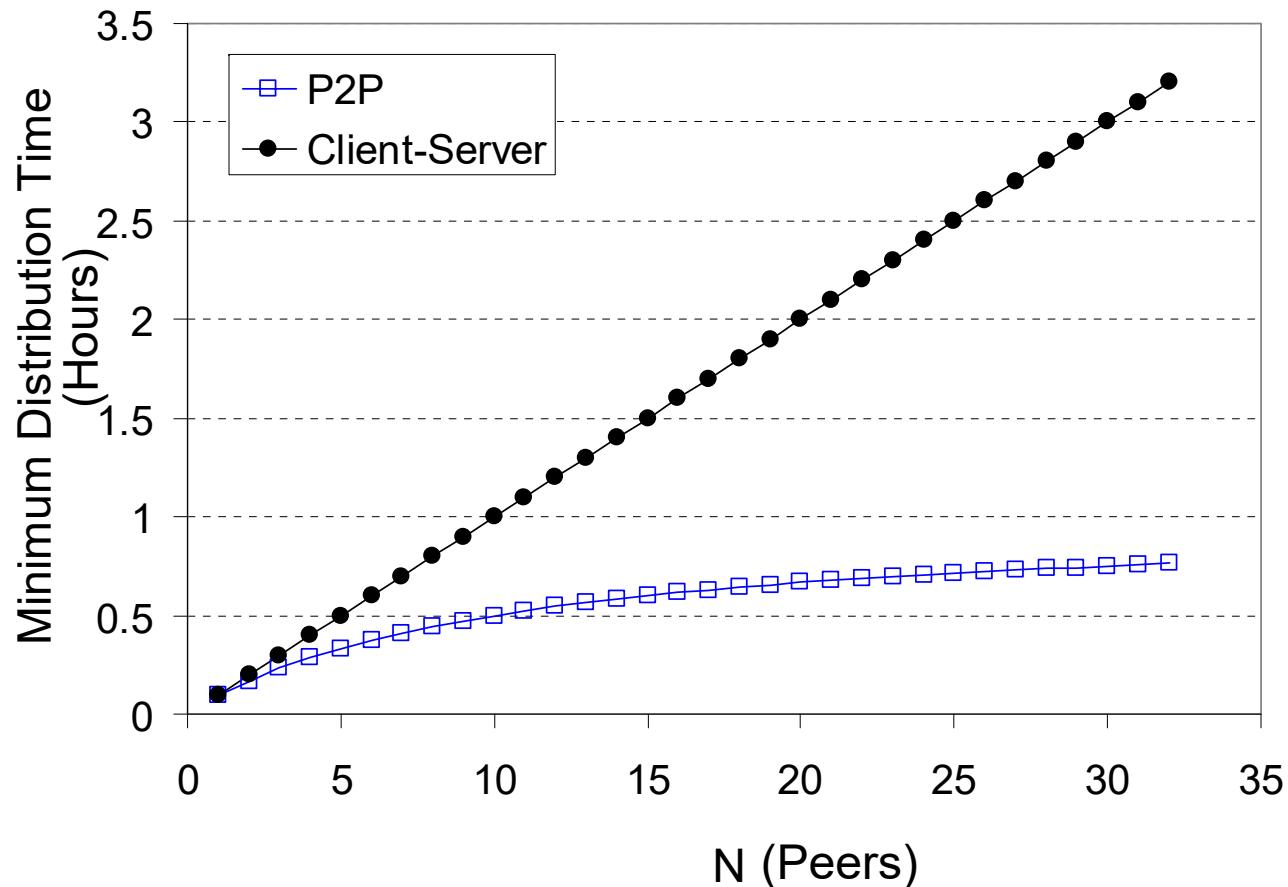
File distribution: Client-Server vs. P2P

- How much time is required to distribute file (size F) from one server to N peers?
 - ❖ Peer upload/download capacity is limited resource



Client-server vs. P2P: Example

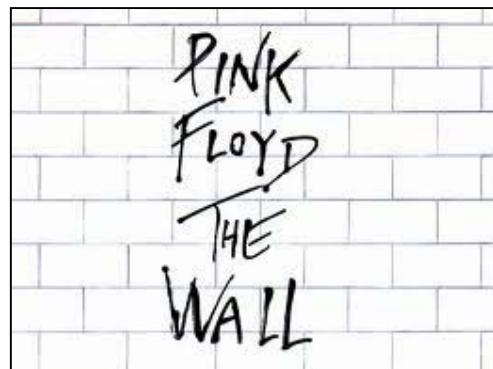
- Client upload rate: u , Assume: Upload file from client: $F/u = 1$ hour
- Server upload rate: $u_s = 10u \rightarrow$ file upload in 0.1 hour (6 min)
- Time to get file to all N peers



P2P File Sharing

Example

- Alice runs P2P client app on her computer
- Intermittently connects to Internet
 - ❖ New IP address for each connection
- Asks for "The Wall"
- Application displays other peers that have copy of "The Wall"



- Alice chooses one peer → Bob
- File is copied from Bob's PC to Alice's notebook using TCP
- While Alice downloads, she also uploads to other users
- Alice's peer is both a Web client and a transient Web server

All peers are servers = highly scalable!

BitTorrent: Multiple Uploaders

- Designed in 1999 by Bram Cohen to take advantage of high download speeds while mitigating slow upload rates
- Aggregation of many slow upload speeds = **FAST DOWNLOADS!**
- Users who download pieces of files automatically upload pieces they already have

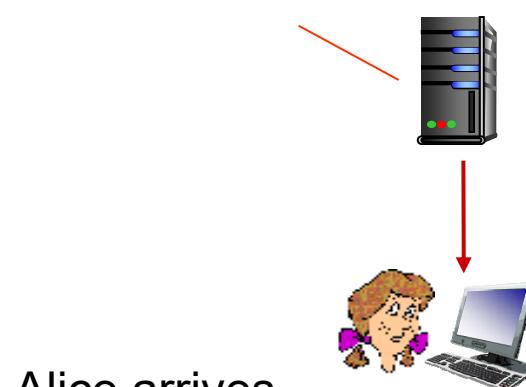


P2P File Distribution: BitTorrent

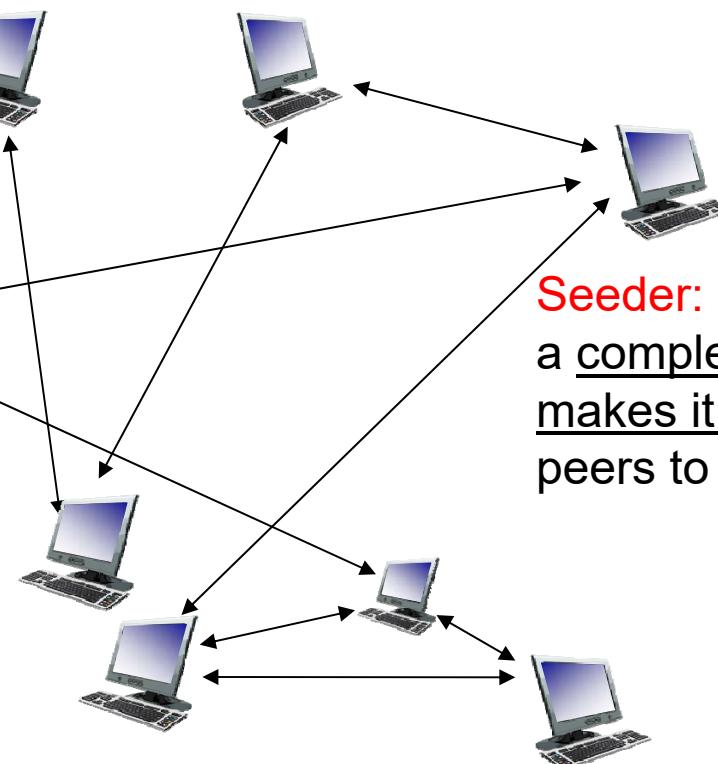
- File divided into 256KB chunks
- Peers in torrent send/receive file chunks



Tracker: tracks peers participating in torrent



Torrent or swarm: group of peers exchanging chunks of a file



P2P File Distribution: BitTorrent

- Peer joining torrent:
 - ❖ Registers with tracker to get list of peers, connects to subset of peers ("neighbors")
 - ❖ Initially has no chunks, but accumulates them over time from other peers
- While downloading, peer uploads chunks to other peers
- Peer may change peers with whom it exchanges chunks
 - ❖ Churn: peers may come and go
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

BitTorrent: Requesting, Sending Chunks

Requesting chunks:

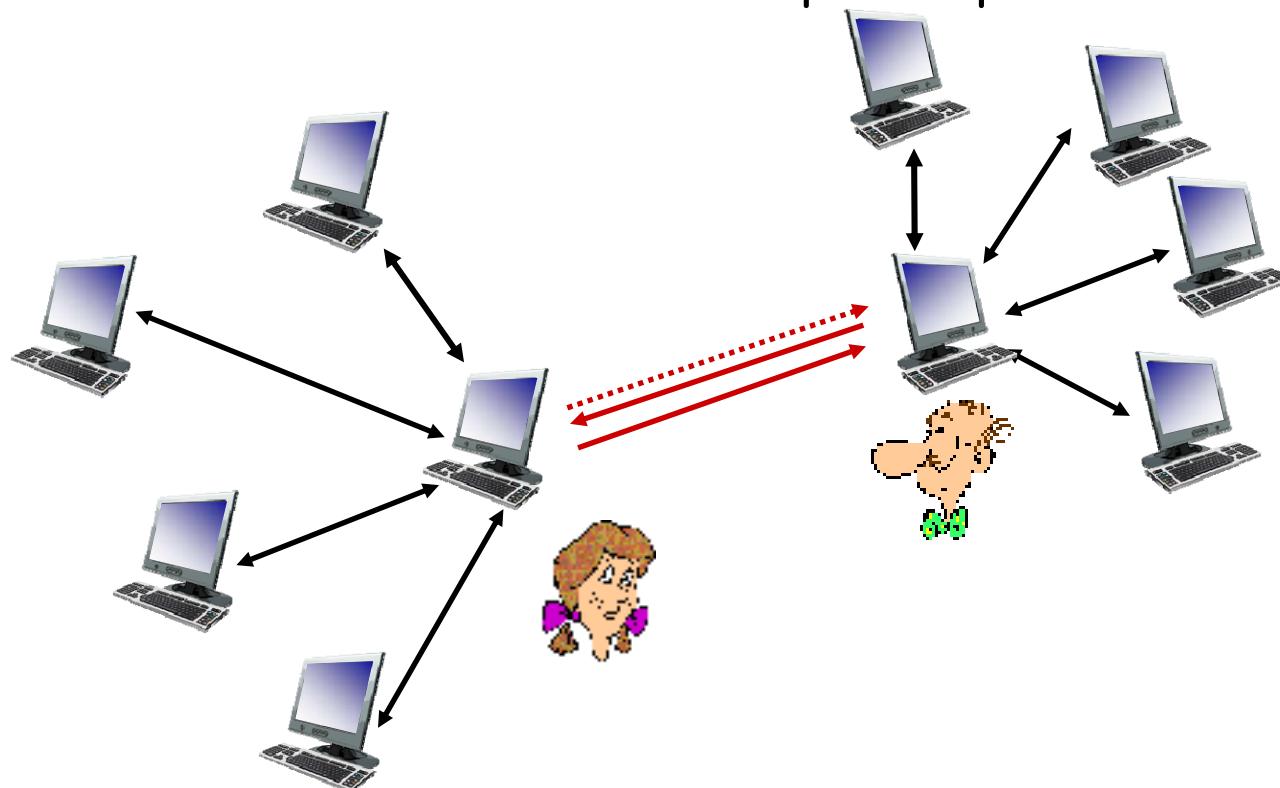
- At any given time, different peers have different subsets of file chunks
- Periodically, Alice asks each peer for list of chunks they have
- Alice requests missing chunks from peers, rarest first

Sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks at highest rate
 - ❖ Other peers are choked by Alice (do not receive chunks from her)
 - ❖ Re-evaluate top 4 every 10 secs

BitTorrent: tit-for-tat

- Every 30 secs: Alice randomly selects another peer, starts sending chunks
 - ❖ Alice “optimistically unchoke” Bob
 - ❖ Alice becomes one of Bob’s top-four providers; Bob reciprocates
 - ❖ Bob becomes one of Alice’s top-four providers



Chapter 2: Application Layer

- 2.1 Principles of Network Applications
- 2.2 Web and HTTP
- FTP
- 2.3 Electronic Mail
 - ❖ SMTP, POP3, IMAP
- 2.4 DNS
- 2.5 P2P Applications
- 2.6 Video Streaming and Content Distribution Networks
- 2.7 Socket Programming:
Creating Network Applications

Socket Programming

Goal: Learn how to build client/server applications that communicate using sockets

Socket API

- ❑ Socket is explicitly
 - ❖ created,
 - ❖ used,
 - ❖ then released by apps
- ❑ Client/Server paradigm

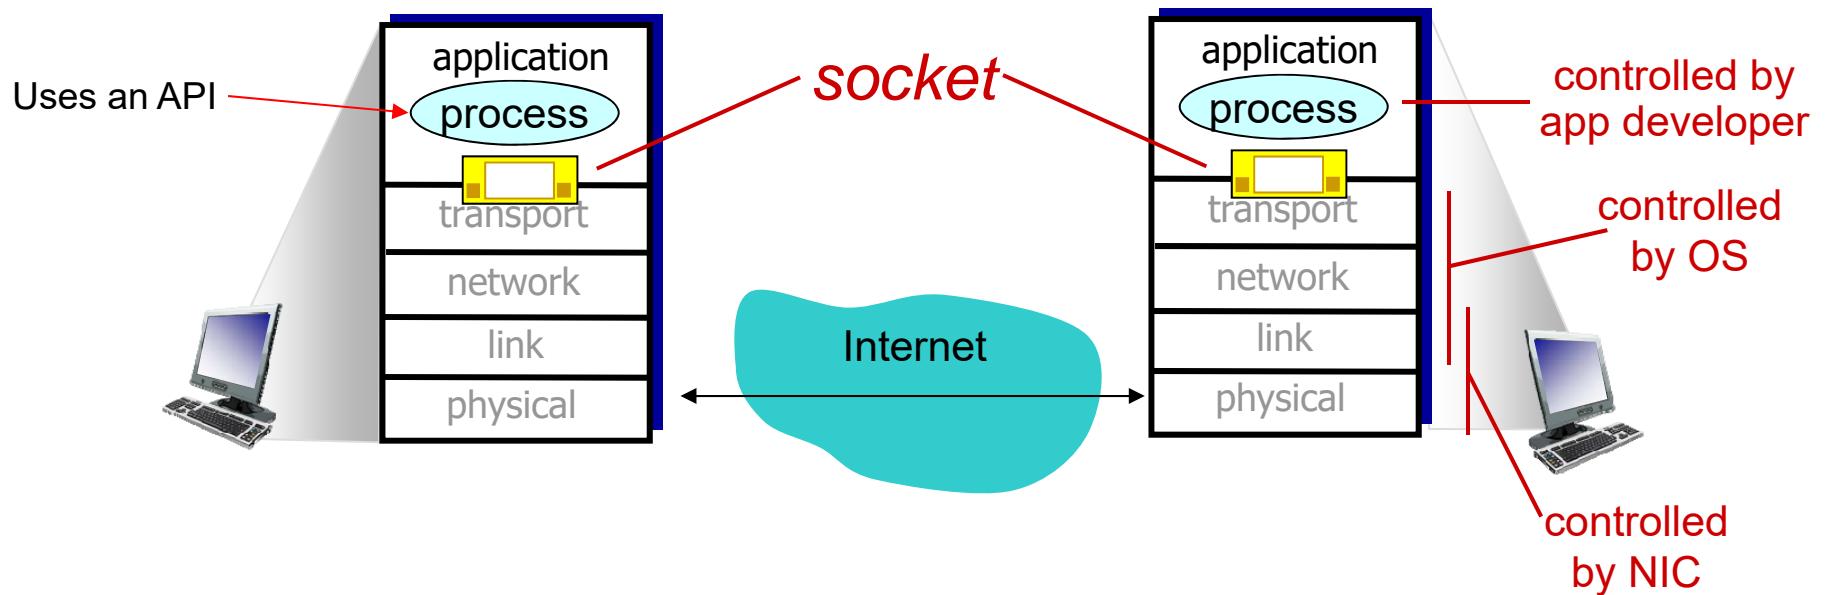
socket

A *host-local, application-created, OS-controlled* interface (a “door”) into which application process can both send and receive messages to/from another application process

Socket Programming

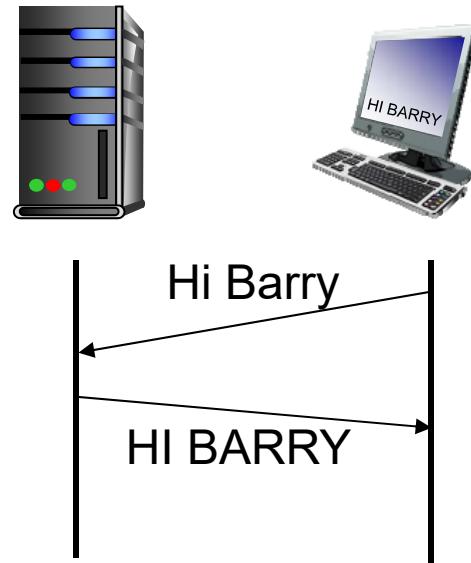
Socket: a door between application process and end-to-end transport protocol (UDP or TCP)

- Two socket types for two transport services:
 - ❖ UDP: unreliable datagram
 - ❖ TCP: reliable, byte stream-oriented



Application Example

- Client reads a line of characters (data) from its keyboard and sends the data to the server
- Server receives the data and converts characters to uppercase
- Server sends the modified data to the client
- Client receives the modified data and displays the line on its screen



Socket Programming With UDP

UDP: transmitted data may be lost or received out-of-order

UDP: no “connection” between client & server

- ❑ No handshaking before sending data
- ❑ Sender application explicitly attaches to each packet
 - ❖ IP destination address
 - ❖ Port #
- ❑ Rcvr extracts sender IP address and port # from packet and uses them to reply to the sender

Application viewpoint:

- ❑ UDP provides unreliable transfer of **groups** of bytes (“datagrams”) between client and server

Client/Server Socket Interaction: UDP

Server (running on serverIP)



```
create socket, port = x:  
serverSocket =  
    socket(AF_INET,SOCK_DGRAM)
```

wait for incoming
datagram

read datagram from
serverSocket

extract IP, port

create datagram with client IP
address, port number; send to
serverSocket

Client



```
create socket:  
clientSocket =  
    socket(AF_INET,SOCK_DGRAM)  
  
create datagram with serverIP and  
port=x; send datagram via  
clientSocket
```

```
read datagram from  
clientSocket  
close  
clientSocket
```

AF_INET - Address Family used for the socket you're creating (e.g., Internet Protocol address-INET)

Example App: UDP Server

Python UDPServer

```
from socket import *  
serverPort = 12000  
  
netstat -nao | find /i "listening"  
  
create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)  
bind socket to local port  
number 12000 → serverSocket.bind(("", serverPort))  
print ('The server is ready to receive')  
  
loop forever → while 1:  
  
Read from UDP socket into  
message, getting client's  
address (client IP and port) → message, clientAddress = serverSocket.recvfrom(2048)  
→ modifiedMessage = message.upper()  
  
send upper case string  
back to this client → serverSocket.sendto(modifiedMessage, clientAddress)  
  
Python uses indentation to  
delineate blocks of code
```



Notepad++ can parse
Python code and
automatically display code
blocks

Example App: UDP Client

Python UDPCClient

include Python's socket library

```
from socket import *
serverName = 'hostname'
serverPort = 12000
```

Change to the name or IP address of the server

create UDP socket for client

```
clientSocket = socket(AF_INET, SOCK_DGRAM)
```

Network layer - IP

Transport layer - UDP

get user keyboard input

```
message = input('Input lowercase sentence:')
```

Attach server name, port to message; send into socket

```
clientSocket.sendto(message,(serverName, serverPort))
```

read reply characters from socket into string.

serverAddress contains server IP and port

```
clientSocket.recvfrom(2048)
```

print out received string

```
print modifiedMessage
```

close socket

```
clientSocket.close()
```

Addressing

- ❑ For TCP/IP socket communication, generally need 4 parameters:
 - ❖ Source Identifier (32-bit IP address)
 - ❖ Source Port (16-bit)
 - ❖ Destination Identifier (32-bit IP address)
 - ❖ Destination Port (16-bit)
- ❑ UDP uses only Destination (IP+Port) for demultiplexing
- ❑ TCP uses Source + Destination
 - ❖ (quadruple: Src IP, Src Port, Dest IP, Dest Port)

Socket Programming with TCP

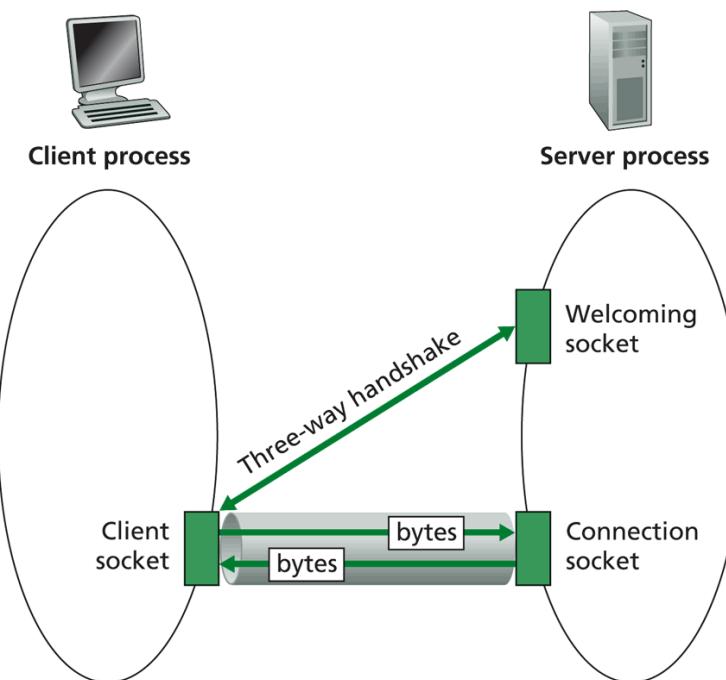
Client must contact server

- ❑ Server process must first be running
- ❑ Server must have created socket (door) that welcomes client's contact

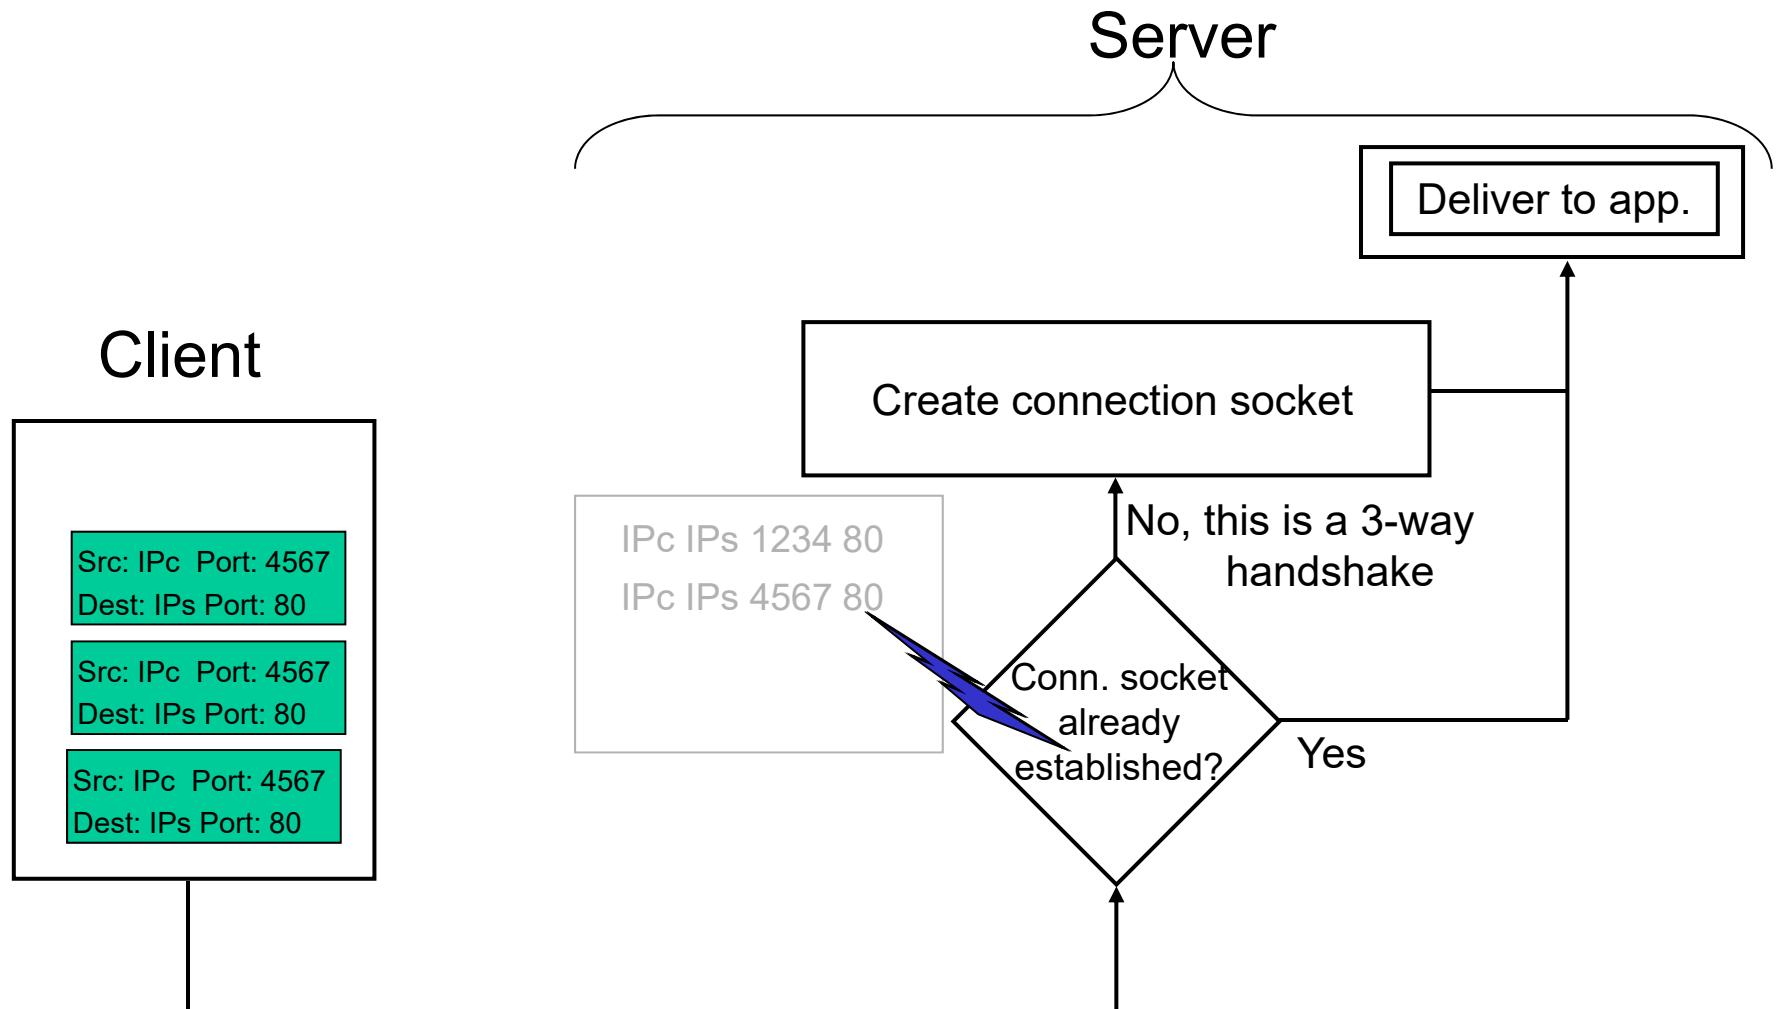
Client contacts server by:

- ❑ Creating client TCP socket
- ❑ Specifying IP address, port number of server process
- ❑ When **client creates socket**: client TCP establishes connection to server TCP

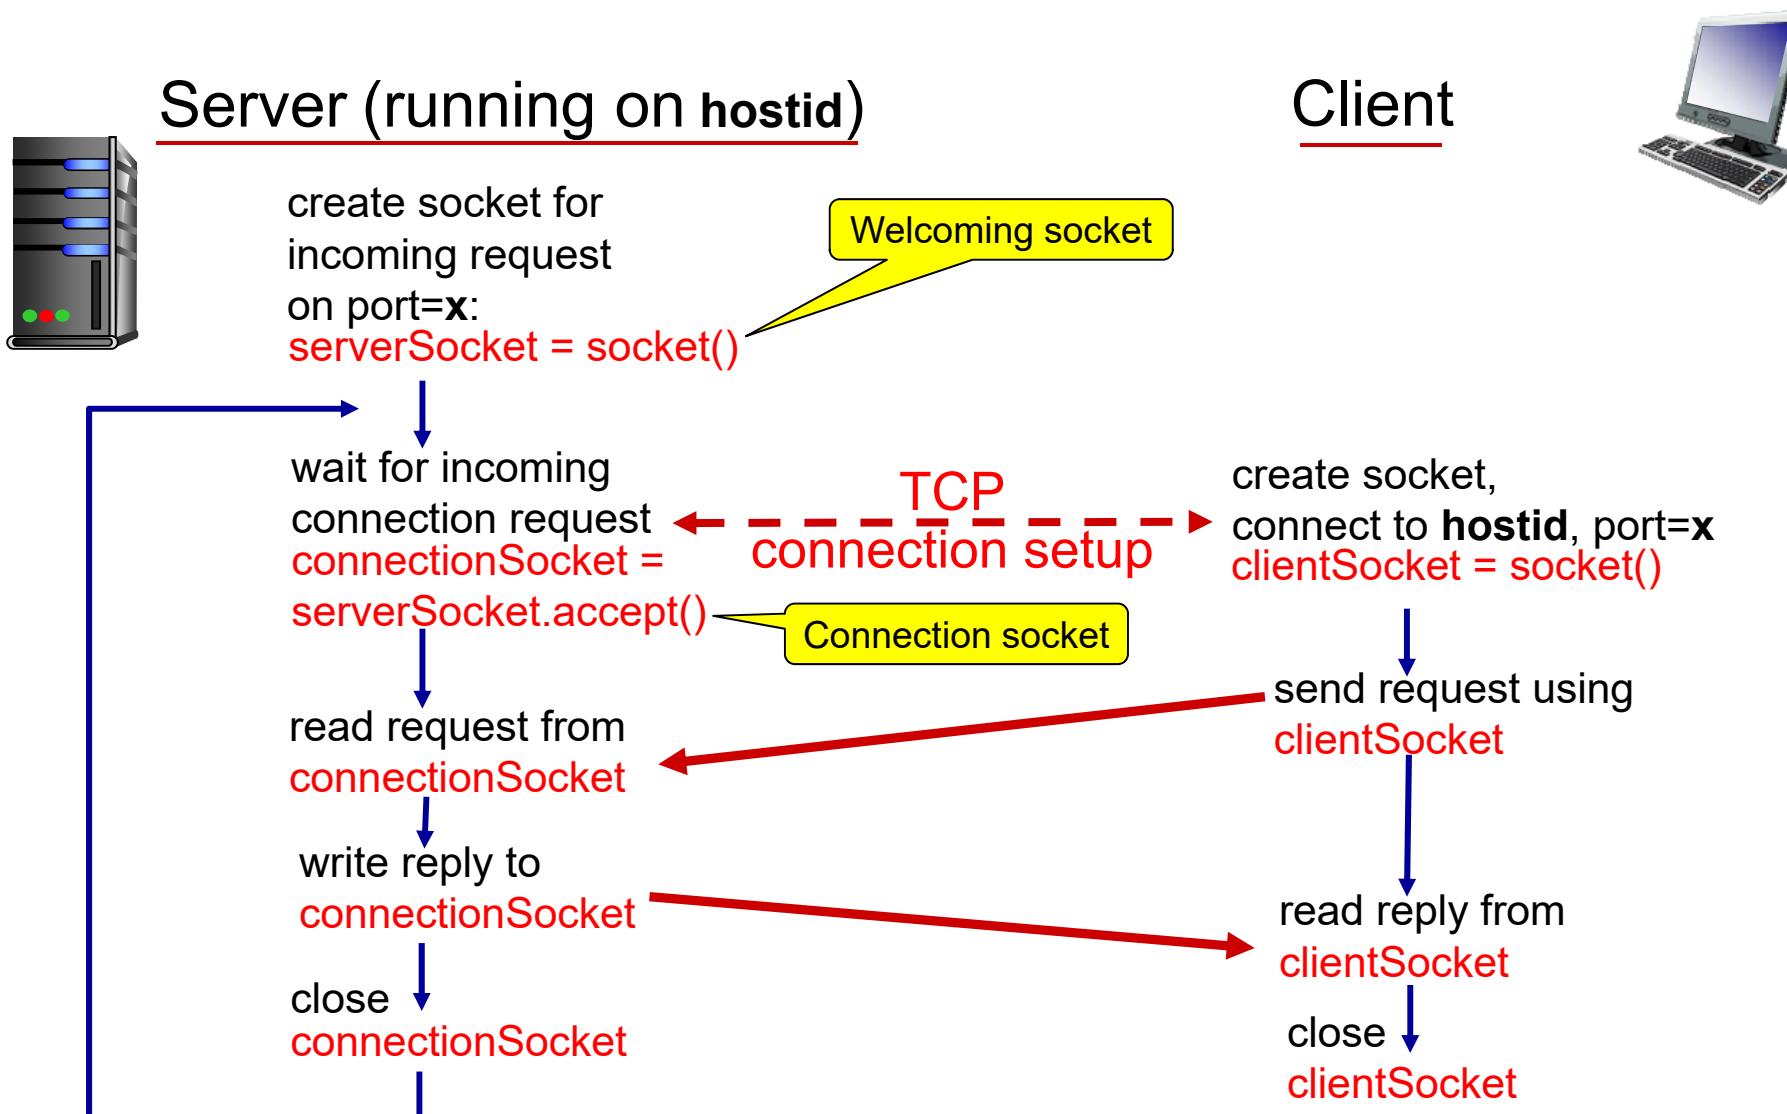
- ❑ When contacted by client, **server TCP creates new connection socket** for server process to communicate with client
 - ❖ Allows server to talk with multiple clients
 - ❖ Source port numbers used to distinguish clients



Socket Programming with TCP



Client/Server Socket Interaction: TCP



Example App: TCP Server

Python TCPServer

```
from socket import *
serverPort = 12000
create TCP welcoming  
socket → serverSocket = socket(AF_INET,SOCK_STREAM)
server begins listening for  
incoming TCP requests → serverSocket.bind(("",serverPort))
loop forever → while 1:
server waits on accept()  
for incoming requests, new  
socket created on return →     print ('The server is ready to receive')
→ connectionSocket, addr = serverSocket.accept()
read bytes from socket (not  
an address as in UDP) →         sentence = connectionSocket.recv(1024)
→                                         capitalizedSentence = sentence.upper()
close connection to this  
client (but not welcoming  
socket) →         connectionSocket.send(capitalizedSentence)
→                                         connectionSocket.close()
```

The diagram illustrates the state transitions of a TCP connection. It starts with a 'Welcoming socket' (represented by a yellow rounded rectangle) which undergoes a 'listen(1)' operation (indicated by a yellow arrow). This leads to a 'Connection socket' (also a yellow rounded rectangle). The connection socket then performs an 'accept()' operation, which creates a new socket (indicated by a yellow arrow). Finally, the connection socket sends data to this new socket (indicated by a yellow arrow).

Example App: TCP Client

Python TCPClient

create TCP socket for
server, remote port 12000

```
from socket import *
serverName = 'servername'
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = input('Input lowercase sentence:')
clientSocket.send(sentence)
modifiedSentence = clientSocket.recv(1024)
print 'From Server: ', modifiedSentence
sentence = input('Press enter to exit ...')
clientSocket.close()
```

No need to attach server
name, port

Change to the name or
IP address of the server

Transport layer - TCP

I added so you can see the
server's response