

# Logical Agents

- When we use search to solve a problem we must
  - Capture the knowledge needed to formalize the problem
  - Apply a search technique to solve problem
  - Execute the problem solution



# Knowledge Representation

- Knowledge-based agents
- Wumpus world
- Logic
  - Sentence, syntax, semantics, validity, satisfiability
  - Proof by Enumeration and Inference
- Proposition Logic
  - Inference Rules, Normal Forms
- Automating theorem proving
  - Forward chaining
  - Backward chaining
  - Resolution

# A Brief History of Reasoning

450 B.C	Stoics	Propositional logic, inference (maybe)
322 B.C.	Aristotle	"syllogisms" (inference rules), quantifiers
1565	Cardano	Probability theory (propositional logic + uncertainty)
1847	Boole	Propositional logic (again)
1879	Frege	First-order logic
1922	Wittgenstein	Proof by truth tables
1930	Godel	$\exists$ Complete algorithm for FOL
1930	Herbrand	Complete algorithm for FOL (reduce to propositional)
1931	Godel	$\neg\exists$ complete algorithm for arithmetic
1960	Davis/Putnam	"practical" algorithm for propositional logic
1965	Robinson	"practical" algorithm for FOL - resolution

# Role of Knowledge Representation

- We previously talked about applications of search but not about methods of formalizing the problem.
- Now we look at extended capabilities to general logical reasoning.
- The first step is the role of "knowledge representation" in AI.
- Formally:
  - The intended role of knowledge representation in artificial intelligence is to reduce problems of intelligent action to search problems.
  - "A good description, developed within the conventions of a good KR, is an open door to problem solving; a bad description, using a bad representation, is a brick wall preventing problem solving."

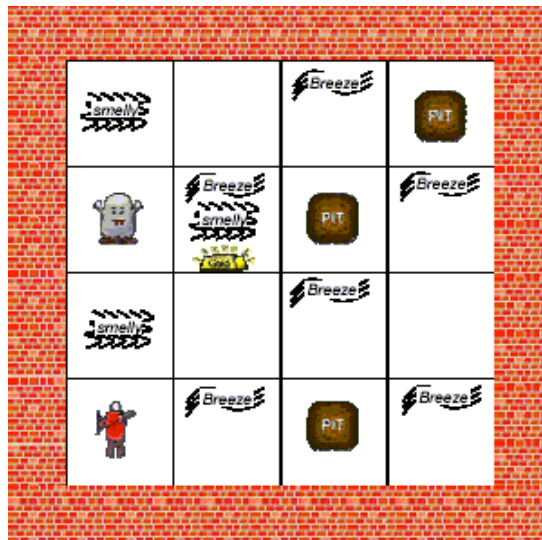
# Knowledge Bases

- *Knowledge base (KB)* is a set of sentences in a formal language
- Declarative approach to building an agent  
TELL it what it needs to know
- Then it can Ask itself what to do
  - Answers follow from KB – *entailment*
- Agents can be viewed at the knowledge level
  - What they know, regardless of implementation
- Or at the implementation level
  - Data structures in KB and algorithms that manipulate them

# A Knowledge-Based Agent

- A knowledge-based agent must be able to
  - Represent states, actions, etc.
  - Incorporate new percepts
  - Update internal representations of the world
  - Deduce hidden properties about the world
  - Deduce appropriate actions
- We will
  - Describe properties of languages to use for logical reasoning
  - Describe techniques for deducing new information from current information
  - Apply search to deduce (or learn) specifically needed information

## The Wumpus World Environment



## Percepts

- Percept = [Stench, Breeze, Glitter, Bump, Scream]
  - Stench in Wumpus square and adjacent (L, R, U, D) squares
  - Breeze in squares adjacent to pit
  - Glitter in gold square
  - Bump when walk into wall or obstacle
  - Everyone hears scream when Wumpus is defeated
  - Agent cannot perceive its own location

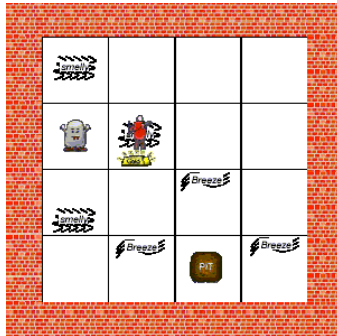
## Goals

- Actions = [goforward, turnleft, turnright, grab, shoot, climb]
- Agent is defeated upon entering room with pit or live Wumpus
- Agent's goal is to find gold, return to [1,1], and climb out of cave
- Score (-1) for each action attempted, (-1000) for being defeated, (+1000) for leaving cave with gold

## Wumpus World Properties

- Is the world deterministic?
  - Yes
- Is the world fully accessible?
  - No, only local perception
- Is the world static?
  - Yes (for now), Wumpus and pits do not move
- Is the world discrete?
  - Yes
- Single-agent?
  - Yes, Wumpus is a domain feature

## Sample Run



- Now we look at
- How to represent facts / beliefs  
"There is a pit in (2,2) or (3,1)"
- How to make inferences  
"No Breeze in (1,2), so pit in (3,1)"

## Formal Representation Through Logic

- **Logic:** a **formal language** for representing information such that conclusions can be drawn
  - **Sentence:** individual piece of knowledge  
English sentence forms one piece of knowledge in English language  
A statement in Java forms one piece of knowledge in Java programming language
  - **Syntax:** form used to represent sentences  
Syntax of Java indicates legal combinations of symbols ( $a = 2+3$ ;)   
Syntax alone does not indicate meaning
  - **Semantics:** mapping from sentences to facts in the world  
They define the truth of a sentence in a world. Add the values of 2 and 3, store the result in the memory location indicated by variable a
  - **Proof System:** a way of manipulating syntactic expressions (enumeration or inference) to get other syntactic expressions (which will tell us something new)

# A few Formal Logics

Language	Ontological Commitment	Epistemological Commitment
Propositional Logic	facts	true/false/unknown
Predicate Logic	facts, objects, relations	true/false/unknown
Temporal Logic	facts, objects, relations, time	true/false/unknown
Probability Theory	facts	degree of belief in [0.0 .. 1.0]
Fuzzy Logic	degree of truth	degree of belief in [0.0 .. 1.0]

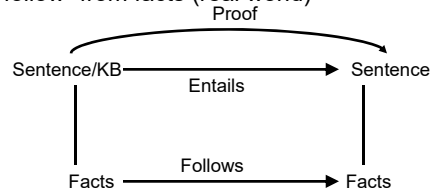
- Which Logic to use?
  - Ontological commitment: What exists? Facts? Objects? Time? Beliefs?
  - Epistemological commitment: What are the states of knowledge?

# Definitions

- The **meaning** of a sentence is a mapping onto the “real world”.
  - This mapping is an **interpretation**.
- A sentence is **valid** (necessarily true, tautology) iff true under all possible interpretations.
  - $A \vee \neg A$
  - A could be:
    - Stench at [1,1]
    - $2+3=5$
  - These statements are not valid.
    - $A \wedge \neg A$
    - $A \vee B$
- The last statement is **satisfiable**, meaning there exists at least one interpretation that makes the statement true. The previous statement is **unsatisfiable**.

# Entailment

- The proof capability relies on relationships between items in the language:
  - Sentences “entail” sentences (representation level)
  - Facts “follow” from facts (real world)



- Entail/Follow means the new item is true if the old items are true
- A collection of sentences, or knowledge base ( $KB$ ), entail a sentence  $\alpha$ 
  - $KB \models \alpha$
- $KB$  entails the sentence  $\alpha$  if and only if  $\alpha$  is true in all worlds where the  $KB$  is true


## Entailment Examples

- KB
  - The Giants won
  - The Reds won
- Entails
  - The Giants won and the Reds won
- KB
  - To get a perfect score your program must be turned in today
  - I always get perfect scores
- Entails
  - I turned in my program today
- Entailment is a relationship between sentences (*syntax*) that is based on *semantics*.



## Entailment in the Wumpus World

- Situation after detecting nothing in (1,1) and moving right with a breeze in (2,1)
- Consider ALL possible models for ?s assuming only pits
  - 3 possibilities  $\Rightarrow$  8 models

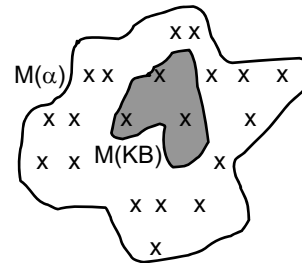


## Proof Methods

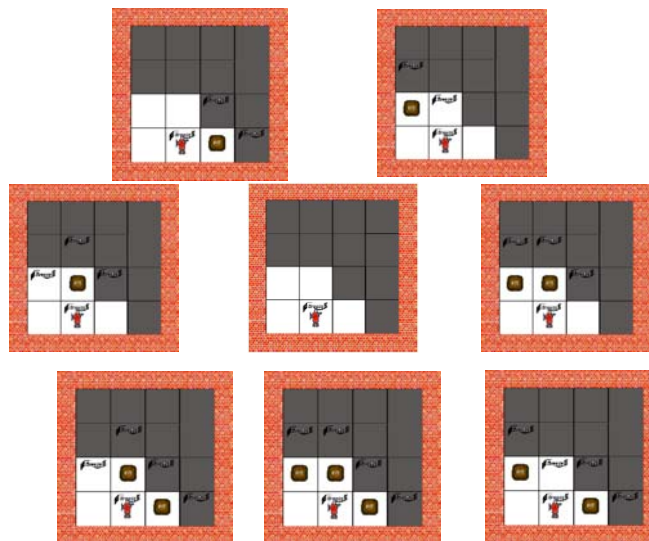
- Proof methods divide into (roughly) two kinds:
  - Model checking
    - Truth table enumeration (sound and complete for propositional logic sentence) (exponential in  $n$ )
    - Improved backtracking
    - Heuristic search in model space
      - Show that all interpretations in which the left hand side of the rule is true, the right hand side is also true (sound but incomplete)
  - Application of inference rules
    - Sound generation of new sentences from old
    - Proof = a sequence of inference rule applications
      - Can use inference rules as operators in a standard search algorithm.
      - Typically requires translation of sentences into a **normal form**

# Proof by Enumeration (Models)

- Logicians think in terms of models, which are formally structured worlds with respect to which truth can be evaluated
- We say  $m$  is a model of a sentence  $\alpha$  if  $\alpha$  is true in  $m$
- $M(\alpha)$  is the set of all models which hold for  $\alpha$
- The set  $M(KB)$  is an enumeration of the entailment of a KB
- Then  $KB \models \alpha$  if and only if  $M(KB) \subseteq M(\alpha)$   
 $KB = \text{Giants won and Reds won}$   
 $\alpha = \text{Giants won}$

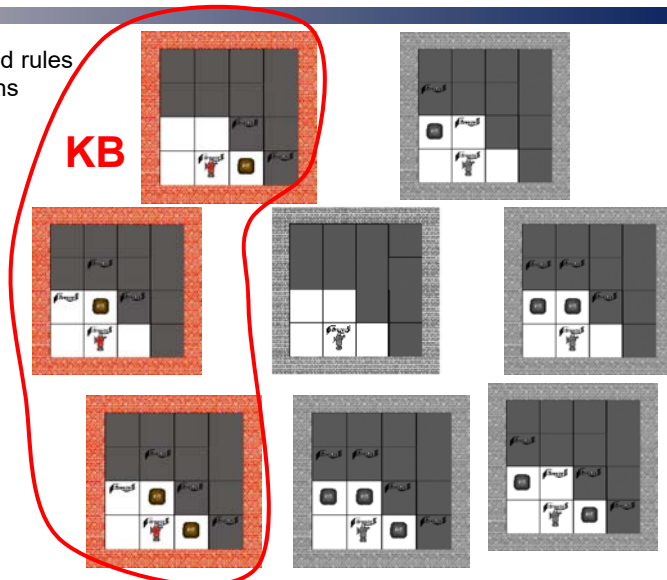


# Wumpus Models



# Wumpus Models

KB = Wumpus-world rules  
+ observations

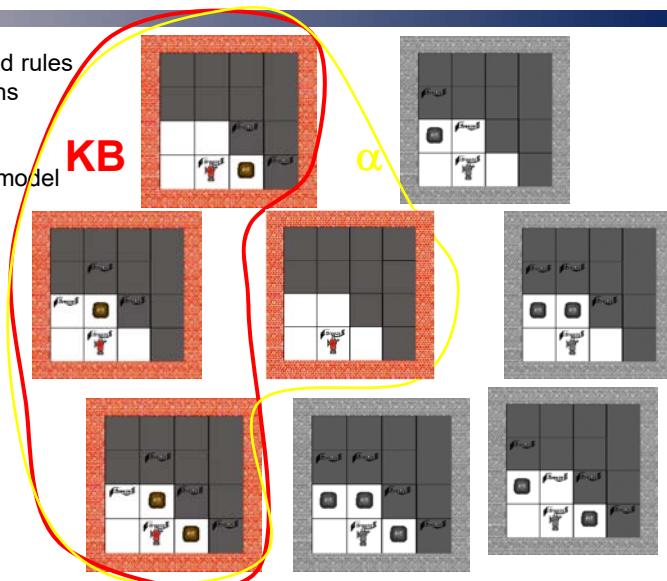


# Wumpus Models

KB = Wumpus-world rules  
+ observations

$\alpha$  = "(1,2) is safe"

KB  $\models$   $\alpha$ , proved by model  
checking

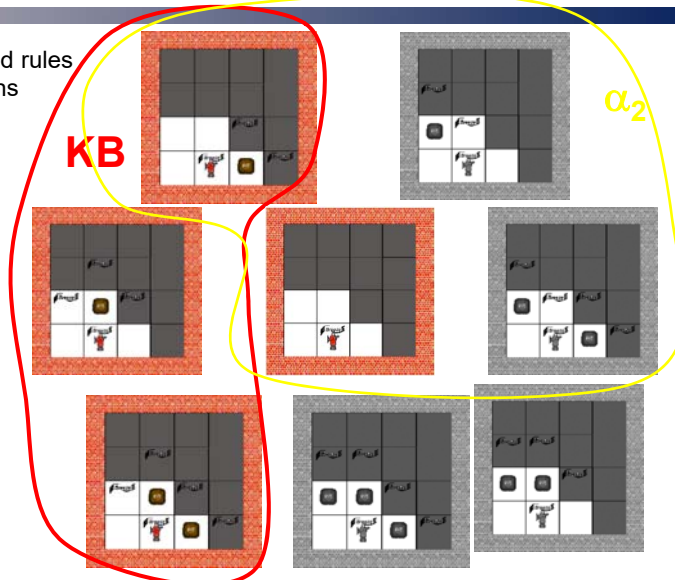


# Wumpus Models

**KB** = Wumpus-world rules  
+ observations

$\alpha_2$  = "(2,2) is safe"

$KB \models \alpha_2$



## Proof by Inference

- $KB \vdash_i \alpha$  : sentence  $\alpha$  can be derived from  $KB$  by procedure  $i$
- Inference can be used two different ways:
  1. Generate new sentences that are entailed by  $KB$
  2. Determine whether or not sentence is entailed by  $KB$
- A **sound** inference procedure generates only entailed sentences.
  - Whenever  $KB \vdash_i \alpha$ , then  $KB \models \alpha$  is true
- A **complete** inference procedure can generate all sentences in the knowledge base.
  - Whenever  $KB \models \alpha$ , then  $KB \vdash_i \alpha$
- Modus ponens yes      Abduction no
 

$$\frac{A, A \Rightarrow B}{B}$$

$$\frac{B, A \Rightarrow B}{A}$$

# Propositional Logic

- Propositional logic is the simplest logic – illustrates basic ideas
- Proposition symbols  $P$ ,  $Q$ , etc., are **unit atoms** and **sentences**
- If  $S_1$  and  $S_2$  are sentences then so are  $\neg S_1$ ,  $S_1 \wedge S_2$ ,  $S_1 \vee S_2$ ,  $S_1 \Rightarrow S_2$ ,  $S_1 \Leftrightarrow S_2$ 
  - Negation, conjunction, disjunction, implication, and biconditional
- An interpretation  $i$  consists of an assignment of truth values to all proposition symbols  $i(s)$   
An interpretation is a “possible world”  
Given 3 proposition symbols  $P$ ,  $Q$ , and  $R$ , there are 8 interpretations  
Given  $n$  proposition symbols, there are  $2^n$  interpretations
- Models** are worlds in which a particular sentence is true under at least one interpretation
- The true/false value of propositions and combinations of propositions can be calculated using a truth table

# Propositional Logic

- For propositional logic, a row in the truth table is one interpretation
- A logic is “monotonic” as long as entailed sentences are preserved as more knowledge is added

$P$	$Q$	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
T	T	F	T	T	T	T
T	F	F	F	T	F	F
F	T	T	F	T	T	F
F	F	T	F	F	T	T

# Rules of Inference for Propositional Logic

Logical Identities and Equivalences		
Idempotency	$A \vee A \equiv A$	$A \wedge A \equiv A$
Commutativity	$A \vee B \equiv B \vee A$	$A \wedge B \equiv B \wedge A$
	$A \leftrightarrow B \equiv B \leftrightarrow A$	
Associativity	$(A \vee B) \vee C \equiv A \vee (B \vee C)$	$(A \wedge B) \wedge C \equiv A \wedge (B \wedge C)$
Absorption	$A \vee (A \wedge B) \equiv A$	$A \wedge (A \vee B) \equiv A$
Distributivity	$A \wedge (B \vee C) \equiv (A \wedge B) \vee (A \wedge C)$	$A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$
Tautology	$A \vee B \equiv A$ if A is a tautology	$A \wedge B \equiv B$ if A is a tautology
Unsatisfiability	$A \vee B \equiv B$ if A is unsatisfiable	$A \wedge B \equiv A$ if A is unsatisfiable
De Morgan's laws	$\neg(A \wedge B) \equiv \neg A \vee \neg B$	$\neg(A \vee B) \equiv \neg A \wedge \neg B$
Contradiction	$\neg A \wedge A \equiv F$	$\neg A \vee A \equiv T$
Identity	$A \vee F \equiv A$	$A \wedge T \equiv A$
Domination	$A \vee T \equiv T$	$A \wedge F \equiv F$
Double-Negation Elimination	$\neg\neg A \equiv A$	

# Rules of Inference for Propositional Logic

Logical Identities and Equivalences		
Implication	$A \Rightarrow B \equiv \neg A \vee B$	
	$A \leftrightarrow B \equiv (A \wedge B) \vee (\neg A \wedge \neg B)$	
	$A \leftrightarrow B \equiv (\neg A \vee B) \wedge (A \vee \neg B)$	
Modus ponens	$A, A \Rightarrow B \models B$	
Modus tollens	$A \Rightarrow B, \neg B \models \neg A$	
And introduction	$A, B \models A \wedge B$	
Or introduction	$A \models A \vee B \vee C \vee D \vee \dots$	
And elimination	$A \wedge B \wedge C \wedge \dots \wedge Z \models A$	
Unit Resolution	$A \vee B, \neg B \models A$	$(A \vee B) \wedge \neg B \Rightarrow A$
Resolution	$A \vee B, \neg B \vee C \models A \vee C$	$\neg A \Rightarrow B, B \Rightarrow C \models \neg A \Rightarrow C$

# Normal Forms

- Other approaches to inference use syntactic operations on sentences, often expressed in standardized forms
- Conjunctive Normal Form (CNF) *conjunction* of disjunction of literals  
 $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$
- Disjunctive Normal Form (DNF) *disjunction* of conjunction of literals  
 $(A \wedge B) \vee (A \wedge \neg C) \vee (A \wedge \neg D) \vee (\neg B \wedge \neg C) \vee (\neg B \wedge \neg D)$
- Horn Form (restricted) *conjunction* of *Horn clauses* (clauses with only 1 positive literal)  
 $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$
- Often written as a set of implications or Implicative Normal Form (INF):  
 $B \Rightarrow A$   
 $(C \wedge D) \Rightarrow B$

## Propositional Conversion to CNF

$$B \Leftrightarrow A \vee C$$

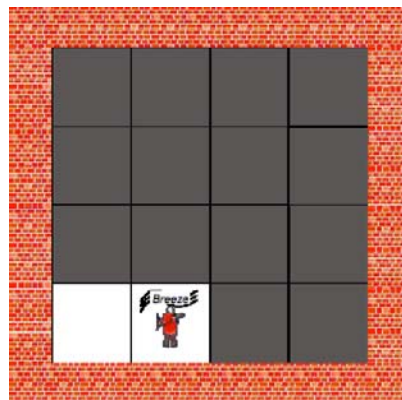
- Eliminate  $\Leftrightarrow$ , replacing  $\alpha \Leftrightarrow \beta$  with  $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$   
 $(B \Rightarrow A \vee C) \wedge (A \vee C \Rightarrow B)$
- Eliminate  $\Rightarrow$ , replacing  $\alpha \Rightarrow \beta$  with  $\neg \alpha \vee \beta$   
 $(\neg B \vee A \vee C) \wedge (\neg(A \vee C) \vee B)$
- Move  $\neg$  inwards using de Morgan's rules and double-negation  
 $(\neg B \vee A \vee C) \wedge ((\neg A \wedge \neg C) \vee B)$
- Apply distributivity law ( $\vee$  over  $\wedge$ ) and flatten  
 $(\neg B \vee A \vee C) \wedge (\neg A \vee B) \wedge (\neg C \vee B)$

# Proof Methods

- Proof methods divide into (roughly) two kinds:
  - Model checking
    - Truth table enumeration (sound and complete for propositional logic sentence) (exponential in  $n$ )
    - Improved backtracking
    - Heuristic search in model space
      - Show that all interpretations in which the left hand side of the rule is true, the right hand side is also true (sound but incomplete)
  - Application of inference rules
    - Sound generation of new sentences from old
    - Proof = a sequence of inference rule applications
      - Can use inference rules as operators in a standard search algorithm.
      - Typically requires translation of sentences into a **normal form**

# Wumpus World Sentences

- Imagine we are at a stage in the game where we have had the following experiences. What is in our knowledge base?
- What can we deduce about the world?
- Given the initial world:  
R1:  $\neg P_{1,1}$
- Also know:  
R2:  $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$   
R3:  $B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$
- And we found out:  
R4:  $\neg B_{1,1}$   
R5:  $B_{2,1}$





## Proof by Truth Table/Enumeration

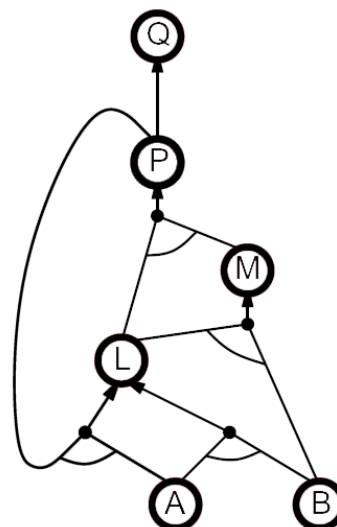
B <sub>1,1</sub>	B <sub>2,1</sub>	P <sub>1,1</sub>	P <sub>1,2</sub>	P <sub>2,1</sub>	P <sub>2,2</sub>	P <sub>3,1</sub>	R1	R2	R3	R4	R5	KB
F	F	F	F	F	F	F	T	T	T	T	F	F
F	F	F	F	F	F	T	T	T	F	T	F	F
:	:	:	:	:	:	:	:	:	:	:	:	:
F	T	F	F	F	F	F	T	T	F	T	T	F
F	T	F	F	F	F	T	T	T	T	T	T	T
F	T	F	F	F	T	F	T	T	T	T	T	T
F	T	F	F	F	T	T	T	T	T	T	T	T
F	T	F	F	T	F	F	T	F	F	T	T	F
:	:	:	:	:	:	:	:	:	:	:	:	:
T	T	T	T	T	T	T	F	T	T	F	T	F

$$KB \Leftrightarrow R1 \wedge R2 \wedge R3 \wedge R4 \wedge R5$$

- Depth-first enumeration of all models is sound and complete
- $O(2^n)$  for  $n$  symbols; problem is co-NP-complete

## Representing Horn Clauses as And-Or Graphs

- $P \Rightarrow Q$
- $L \wedge M \Rightarrow P$
- $B \wedge L \Rightarrow M$
- $A \wedge P \Rightarrow L$
- $A \wedge B \Rightarrow L$
- $A$
- $B$



## Forward and Backward Chaining

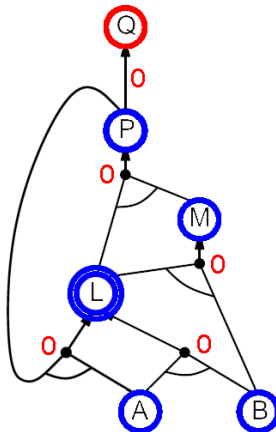
- Modus Ponens: complete for Horn KBs

$$\frac{\alpha_1 \wedge \alpha_2 \wedge \dots \wedge \alpha_n, \quad \alpha_1 \wedge \alpha_2 \wedge \dots \wedge \alpha_n \Rightarrow \beta}{\beta}$$

- Can be used in a forward or backward chaining manner
- Are very natural and run in linear time

## Forward Chaining (FC)

- Fire any rule whose premises are satisfied in the KB, add its conclusion to the KB until query is found



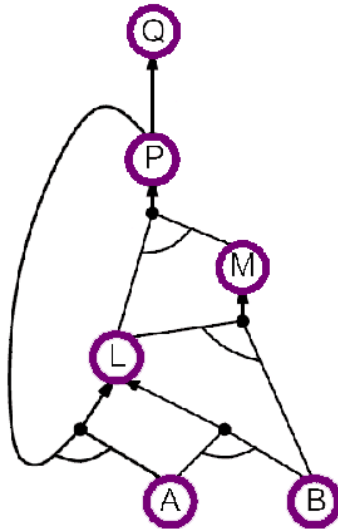
## Proof of Completeness

- FC derives every atomic sentence that is entailed by  $KB$ 
  1. FC reaches a fixed point where no new atomic sentences are derived
  2. Consider this final state as a model  $m$ , assigning true/false to symbols
  3. Every clause in the original  $KB$  is true in  $m$   
Proof: Suppose a clause  $a_1 \wedge \dots \wedge a_k \Rightarrow b$  is false in  $m$ .  
Then  $a_1 \wedge \dots \wedge a_k$  is true in  $m$  and  $b$  is false in  $m$ . Therefore the algorithm has not reached a fixed point
  4. Hence  $m$  is a model of  $KB$
  5. If  $KB \models q$ ,  $q$  is true in every model of  $KB$ , including  $m$

## Backward Chaining (BC)

- Work backwards from the query  $q$ :  
to prove  $q$  by BC  
    check if  $q$  is known already, or  
    prove by BC all premises of some rule  
    concluding  $q$
- Avoid loops: check if new subgoal is already on the goal stack
- Avoid repeated work: check if new subgoal
  - Has already been proved true, or
  - Has already failed

## Backward Chaining Example



## Forward Chaining vs. Backward Chaining

- FC is data driven, automatic, unconscious processing
  - Object recognition, routing decisions
- May do lots of work that is irrelevant to the goal
- BC is goal-driven, appropriate for problem-solving
  - Where are my keys? How do I get into the PhD program?
- Complexity of BC can be much less than linear in size of KB

## Resolution

- Resolution inference rule (for CNF) complete for propositional logic

$$\frac{\alpha_1 \vee \dots \vee \alpha_n \quad \beta_1 \vee \dots \vee \beta_m}{\alpha_1 \vee \dots \vee \alpha_{i-1} \vee \alpha_{i+1} \vee \dots \vee \alpha_n \vee \beta_1 \vee \dots \vee \beta_{j-1} \vee \beta_{j+1} \vee \dots \vee \beta_m}$$

where  $\alpha_i$  and  $\beta_j$  are complementary literals,  
 $\neg \alpha_i \equiv \beta_j$

- $$\frac{A \vee B \quad \neg B}{A}$$

- (We will talk about this in detail with FOPC)

## Propositional Logic Summary

- Logical agents apply inference to a knowledge base to derive new information and make decisions
- Basic concepts of logic:
  - Syntax: formal structure of sentences
  - Semantics: truth of sentences with respect to models
  - Entailment: necessary truth of one sentence given another
  - Inference: deriving sentences from other sentences
  - Soundness: derivations produce only entailed sentences
  - Completeness: derivations can produce all entailed sentences
- Wumpus world requires the ability to represent partial and negated information, reason by cases, etc.
- Forward, backward chaining are linear-time, complete for Horn clauses
- Resolution is complete for propositional logic

# Review

---

- Knowledge Representation
- Logic
- Propositional Logic
- Forward and Backward Chaining Proofs for Horn Clauses
  
- Next class
  - First-Order Predicate Calculus
  - Resolution by Refutation
  - Expert Systems